

Designing Transport-level Encryption For Datacenter Networks

Tianyi Gao, Xinshu Ma, Suhas Narreddy,
Eugenio Luo, Steven Chien, Michio Honda


University of Edinburgh

8th Aug 2025
APNet'25



THE UNIVERSITY of EDINBURGH
informatics


Background: Datacenter transport status quo



[1] Alizadeh et al, SIGCOMM'10 [2] Gao et al, CoNEXT'15 [3] Cai et al, SIGCOMM'22 [4] Hand SIGCOMM'17 [5] Montazeri et al, SIGCOMM'18 [6] Ousterhout et al, ATC'21 [7] Olteanu et al,

Datacenters need end-to-end encryption

- Multi tenancy
- Third-party network/hardware/software on the path



Datacenter transports need message abstraction

Efficient RPC (request-response) support is essential



- Head-of-line blocking avoidance [1]
 - Unordered message delivery
- In-network compute (INC) support [2]
 - e.g., Per-message load balancing
 - Network needs message-level buffering with bounded time
- In-host load balancing [1]
 - Flow-based CPU core affinity creates CPU hotspots

Datacenter transports need message abstraction

Efficient RPC (request-response) support is essential

- Head-of-line blocking avoidance [1]
 - Unordered message delivery
- In-network compute (INC) support [2]
 - e.g., Per-message load balancing
 - Network needs message-level buffering with bounded time
- In-host load balancing [1]
 - Flow-based CPU core affinity creates CPU hotspots

In-order bytestream is unfit




Design space

	Encrypt.	Abstract.	Offload	Protocol	Parallelism
TcpCrypt[4]	TcpCrypt	Stream	N	TCP	Conn.
QUIC[25]	TLS	Stream	N	UDP	Conn.
TCPLS[53]	TLS	Stream	TSO	TCP	Conn.
TLS/TCP[45]	TLS	Stream	Crypto+TSO	TCP	Conn.
SDP	TLS	Msg.	Crypto+TSO	New	Msg.
Homa[40]/NDP[21]	-	Msg.	TSO	New	Msg.
MTP[62]	-	Msg.	TSO	UDP	N/A
Falcon [12]	PSP	Msg.	Full	UDP	Msg. Custom NIC
SRD[58]	-	Dgram.	Full	N/A	Packet Custom NIC
KCM[28]/μTCP[38]	-	Msg.	TSO	TCP	Conn.

- Crypto offload with commodity NICs
 - No compromise from TLS/TCP
- Native transport
 - Flexible protocol design and easy network management

Key question - can we use existing TLS offload


- Autonomous offload [1] (NVIDIA ConnectX-6/7)
 - Mainstream today
 - Likely similar architecture in Fungible (Microsoft) and Netronome NICs
- **It works for non-TCP!**



SDP overview


- One-to-many style socket

```
fd = socket(SOCK_DGRAM, IPPROTO_HOMA) // but reliable  
setsockopt(fd, key) // key handshake already performed, like kTLS  
sendmsg(fd, msg, dst) // or io_uring_prep_sendto(sqe, fd, msg, dst)
```
- Transport-level encryption
 - Transport protocol must be aware of encryption, unlike TLS
- Opportunistic HW offload
- Optional 0-RTT handshake
- ~2800 LoC change in Homa/Linux
- ~300 LoC change in the tx^5 driver
- Support Linux 6.2 and 6.6



Two-level segmentation



- An app message can consist of multiple TSO segments
 - Example below: one app message over two TSO segments
- A TSO segment can consist of multiple packets



Per-message record sequence number space

- Granularity of parallelism
 - TCP (Connection-level) - strict in-order delivery
 - SDP (Message-level) - out-order delivery at both message level and segment level
 - A later message or segment in message can received earlier
 - Global incrementing record sequence number no longer work



Solution: a record sequence number that integrates a message ID with an intra-message record index



Per-message record sequence number space


- Messages can reuse one hardware crypto engine by sharing the record sequence number
- Different messages can be sent and received independently with 48-bits Message ID
- Unique record sequence number for all records across and inside messages to prevent replay attack

Solution: a record sequence number that integrates a message ID with an intra-message record index




Unloaded latency

- SDP outperforms kTLS by 21-32% with hw offload and 16-35% without it
 - Homa is faster than TCP by 5-35 %




Redis throughput

- SDP outperforms kTLS by 5-13 % with TLS offload and 8-17 % without it



Implications

- Opportunity for (proper) evolution of transport in datacenters
 - Internet: TCP-as-a-substrate philosophy for middleboxes
 - Datacenter transports can still evolve
- Is Homa/Linux a right basis?
 - Generic and documented enough for abstraction, packet format, and reasonable performance to build other receiver-driven protocols



Trammell, B. et al., (2014). Evolving transport in the Internet.
IEEE Internet Computing, 18(5), 60-64.

Thank you!
Any questions?