

Building Better Network Simulators with Machine Learning

Mohammad Alizadeh



Massachusetts
Institute of
Technology



How to evaluate new ideas?

- Randomized Control Trial (RCT): The gold standard – measures reality
 - Measures reality
 - Time consuming and risks disruption
 - Not repeatable → poorly suited to iterative design & experimentation
 - Not always possible (e.g., new hardware)
- Simulation: A model of reality
 - Great flexibility, control, repeatability
 - Widely used in the design, planning and operation of networks

Simulation in emerging AI-driven systems

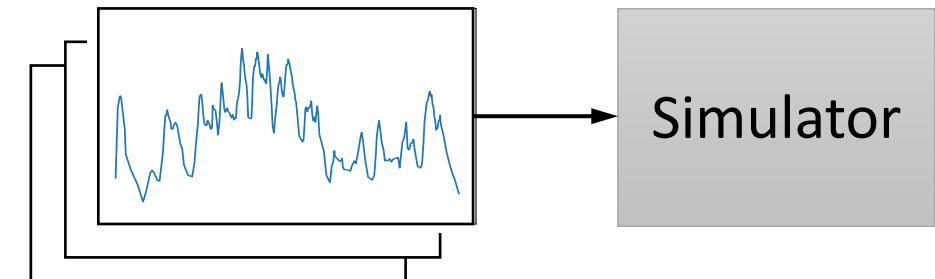
Simulation landscape

Full system

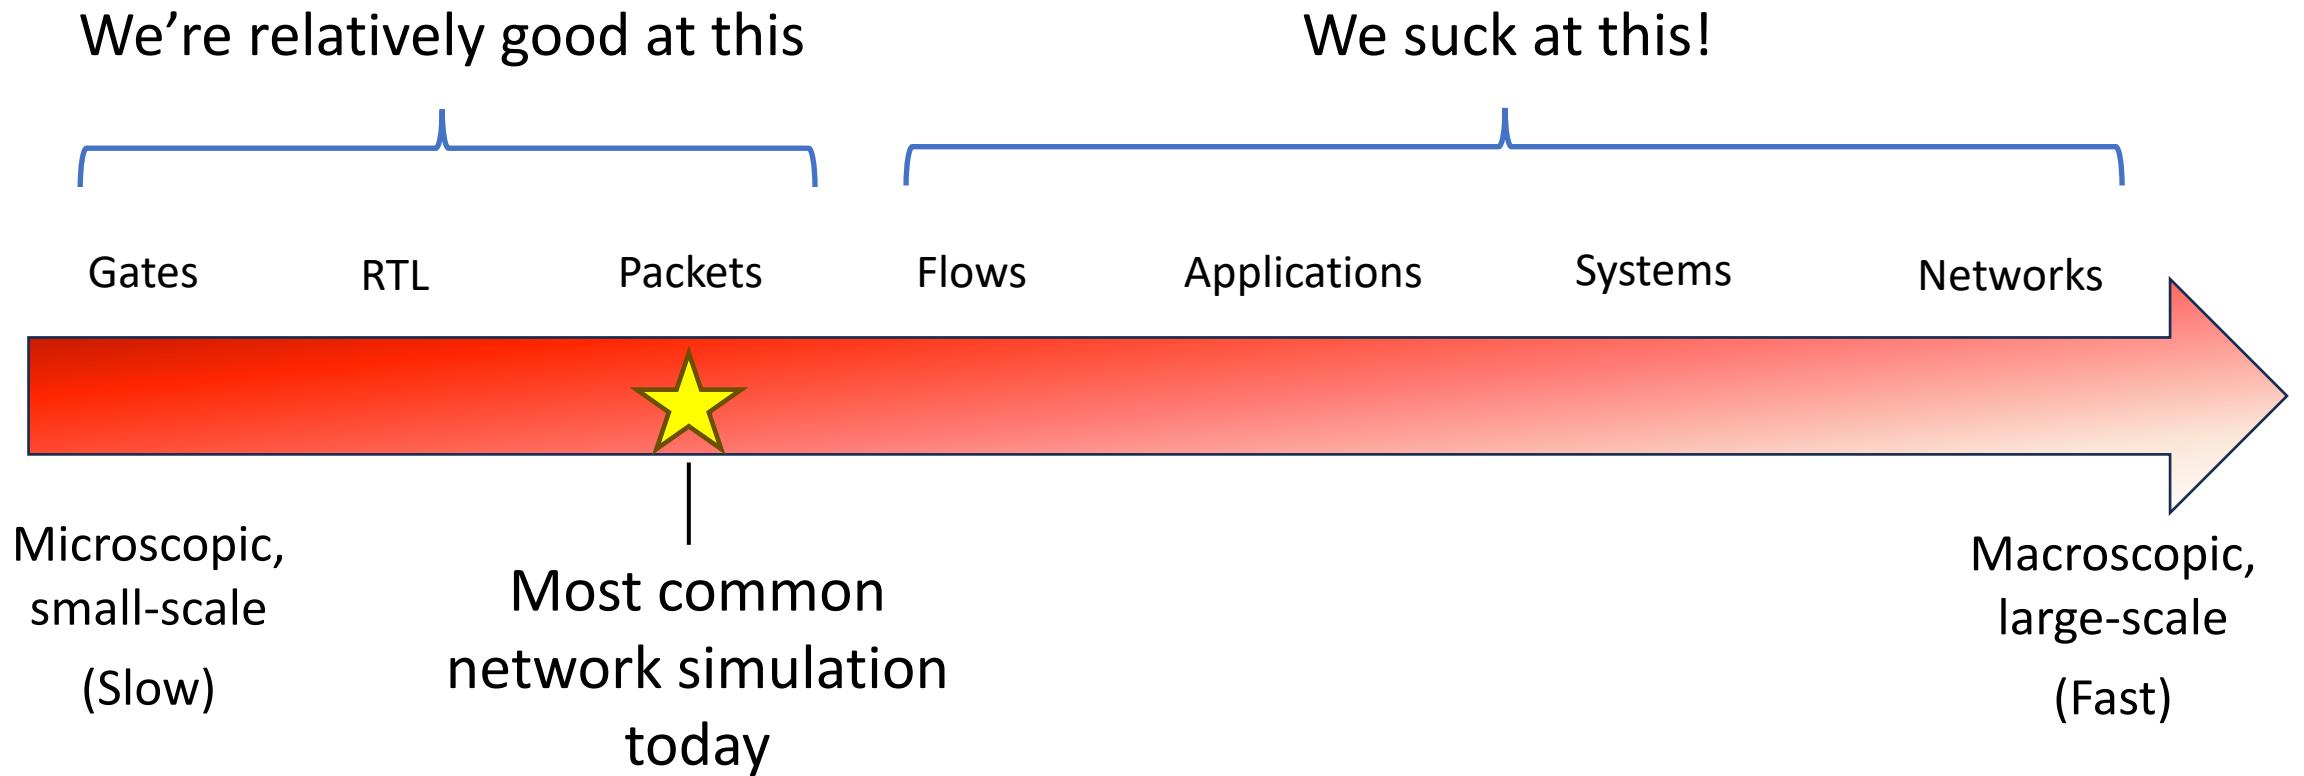
- Model all system components
Needs detailed knowledge
Tedious to build

Trace-driven

- Use traces to capture parts of the system
- Lower complexity & effort
Often less accurate than full-system sim.



From microscopic to macroscopic



Today's simulators are too low-level for many questions

We lack accurate high-level models of real networks,
so we resort to low-level simulation

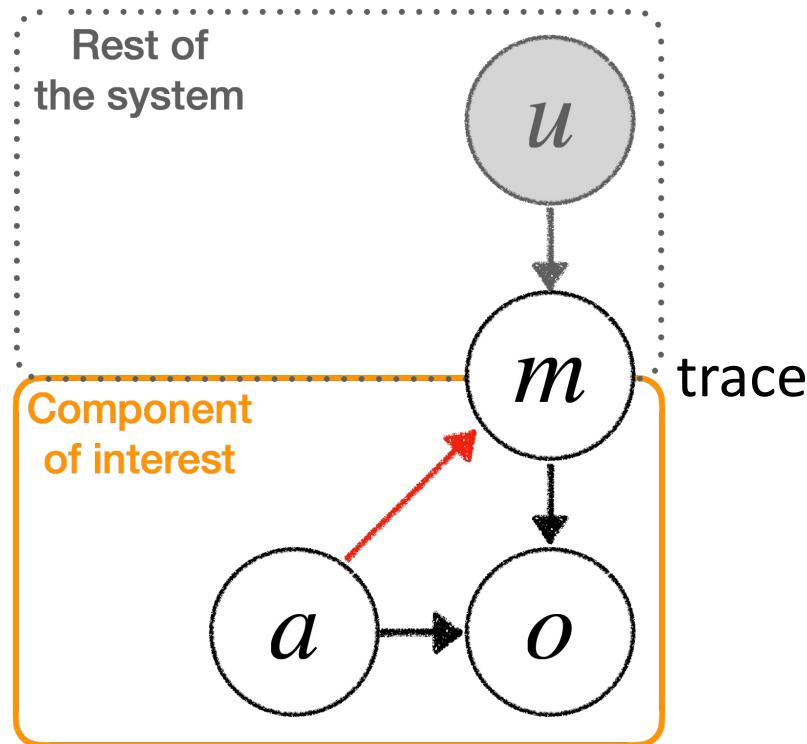
A magnifying glass is positioned over a complex network graph. The network consists of numerous nodes (represented by circles) connected by lines (edges). Some nodes are highlighted in a darker shade of blue or grey, indicating they are being analyzed or are central to the current view. The background is a light beige color.

Vision: Learning high-level network dynamics from data

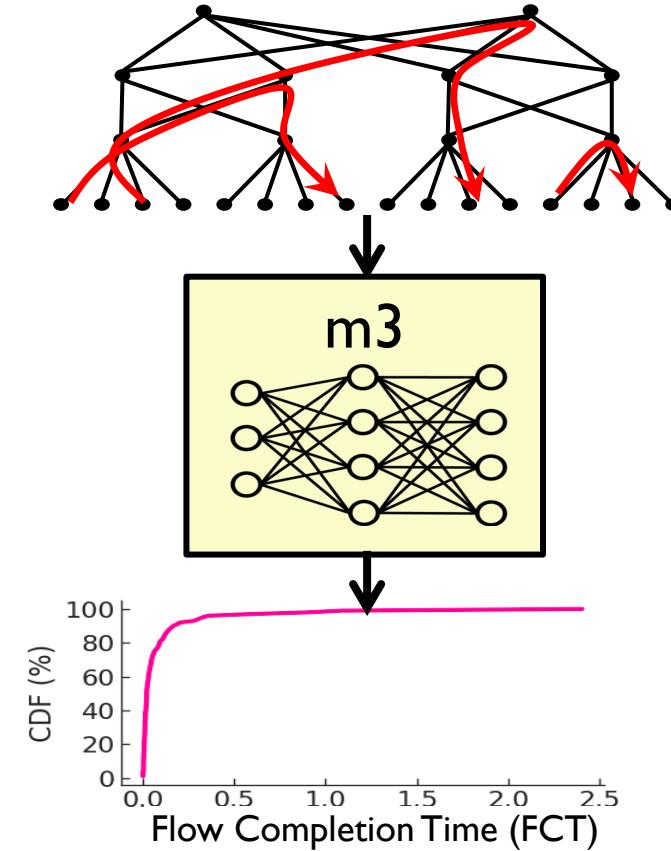
- Networks produce massive amounts of data
- Relatively easy to add instrumentation, run active experiments, ...

Can we use large-scale data to learn accurate high-level network models?

Rest of this Talk...



CausalSim [NSDI'23]:
Unbiased trace-driven simulation



m3 [SIGCOMM'24]:
Fast flow-level performance estimation

CausalSim: A Causal Framework for Unbiased Trace-Driven Simulation



Abdullah
Alomar*

Pouya
Hamadanian*

Arash
Nasr-Esfahany*

Anish
Agarwal

Mohammad
Alizadeh

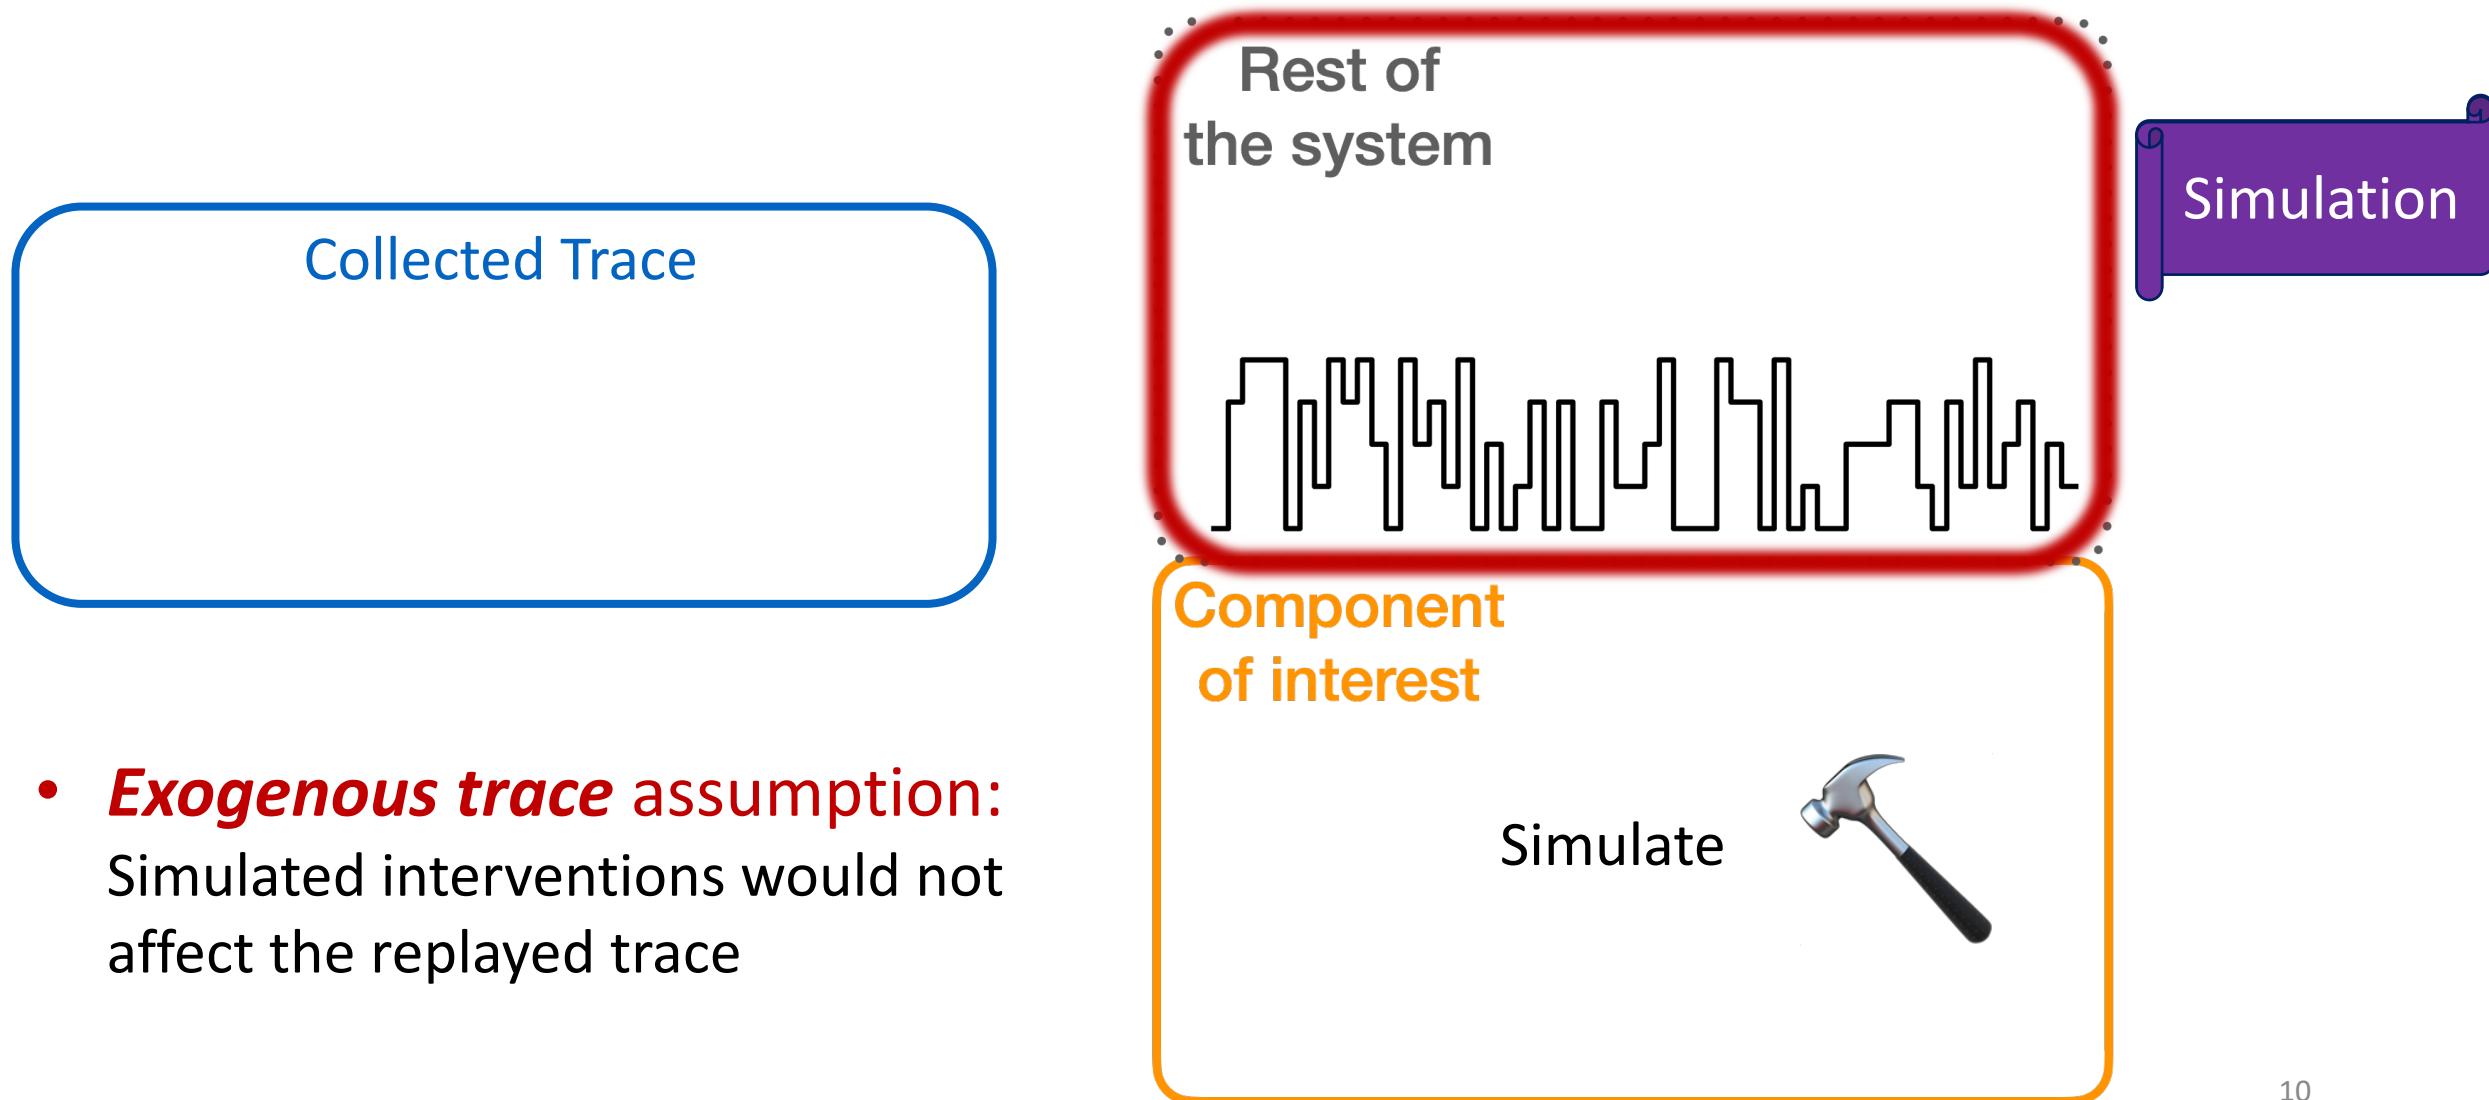
Devavrat
Shah



Massachusetts
Institute of
Technology



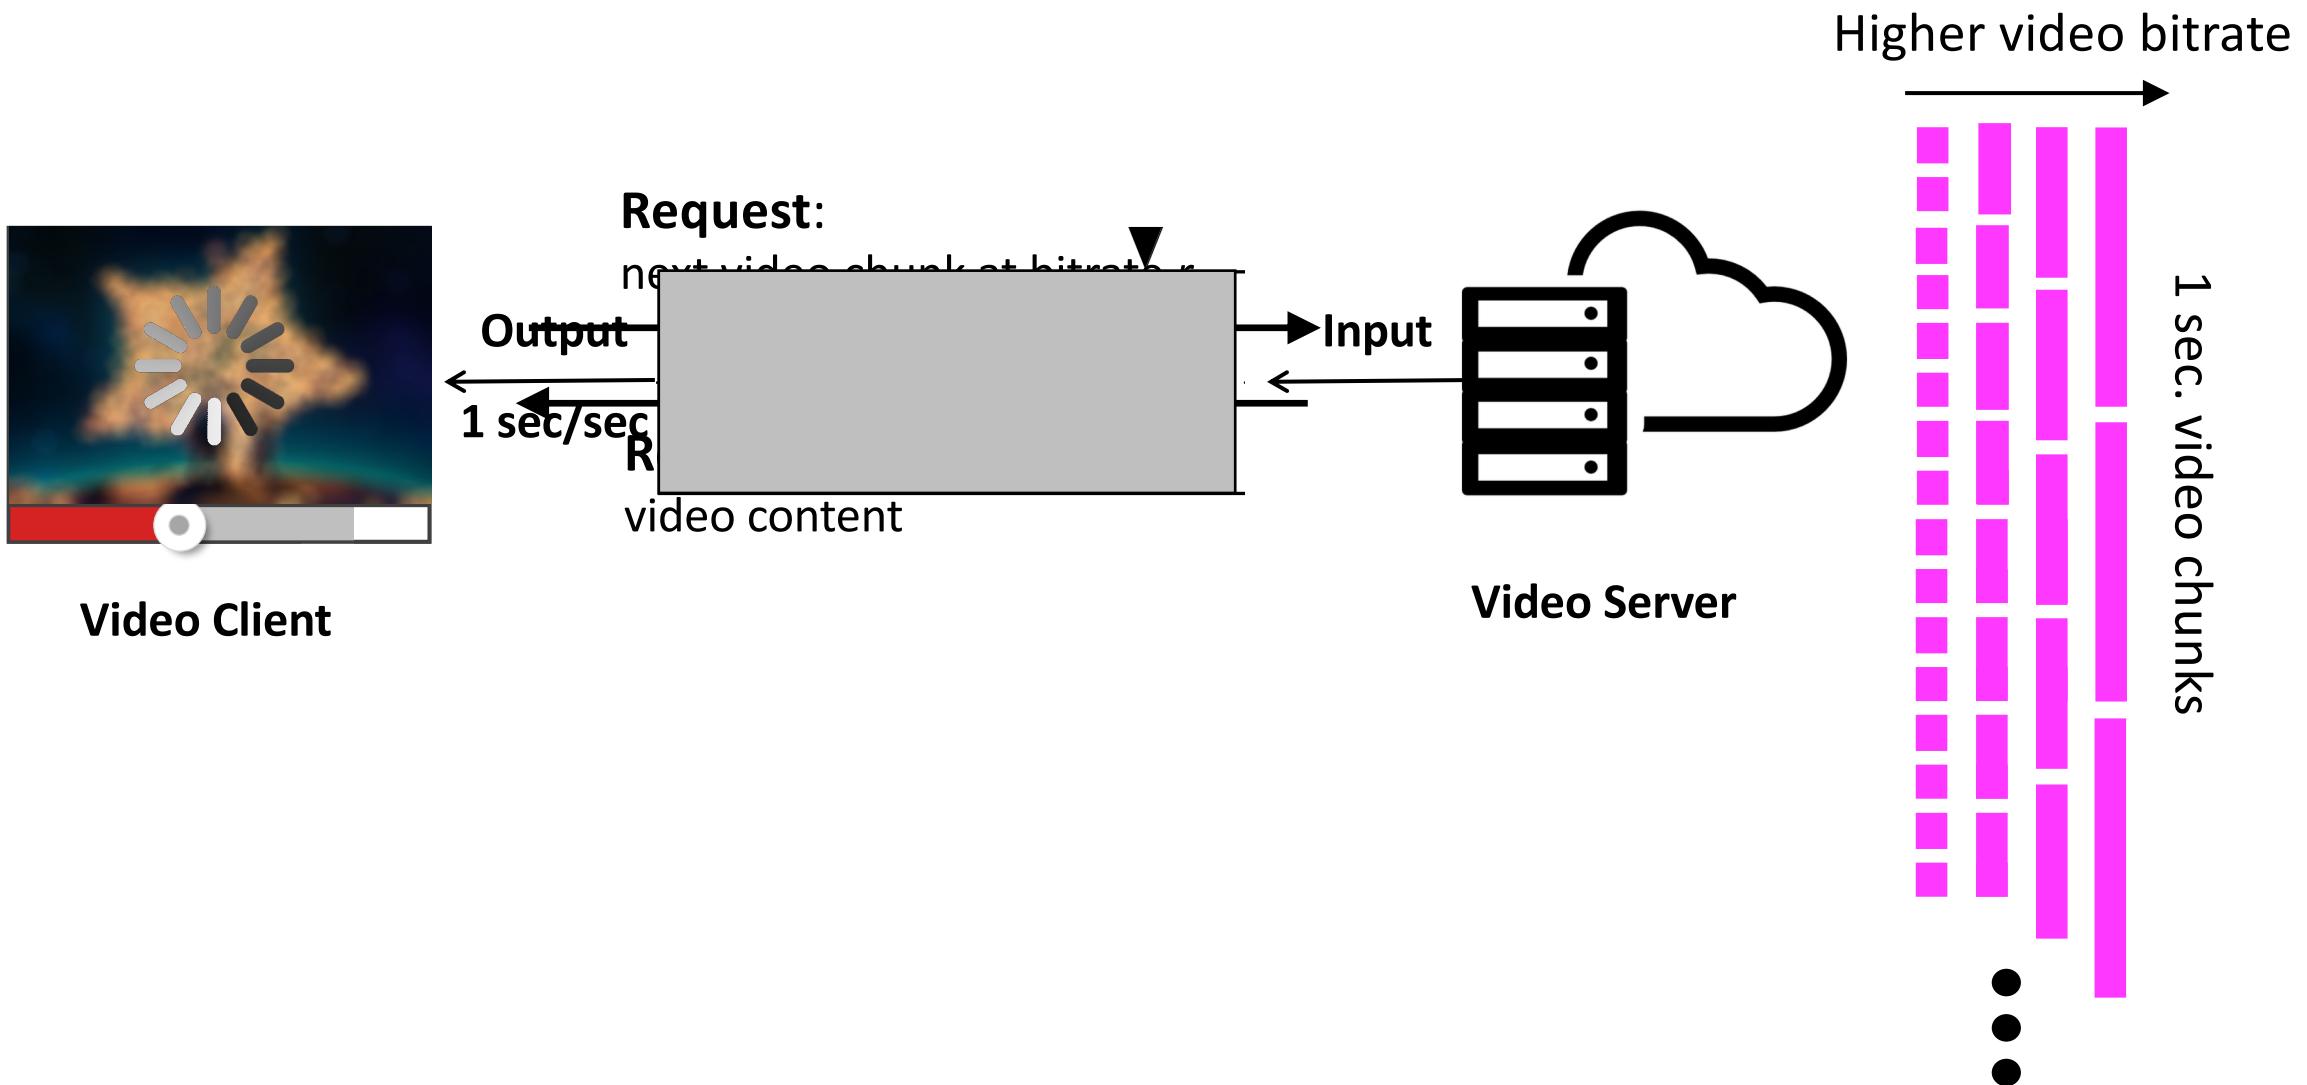
Trace-Driven Simulation



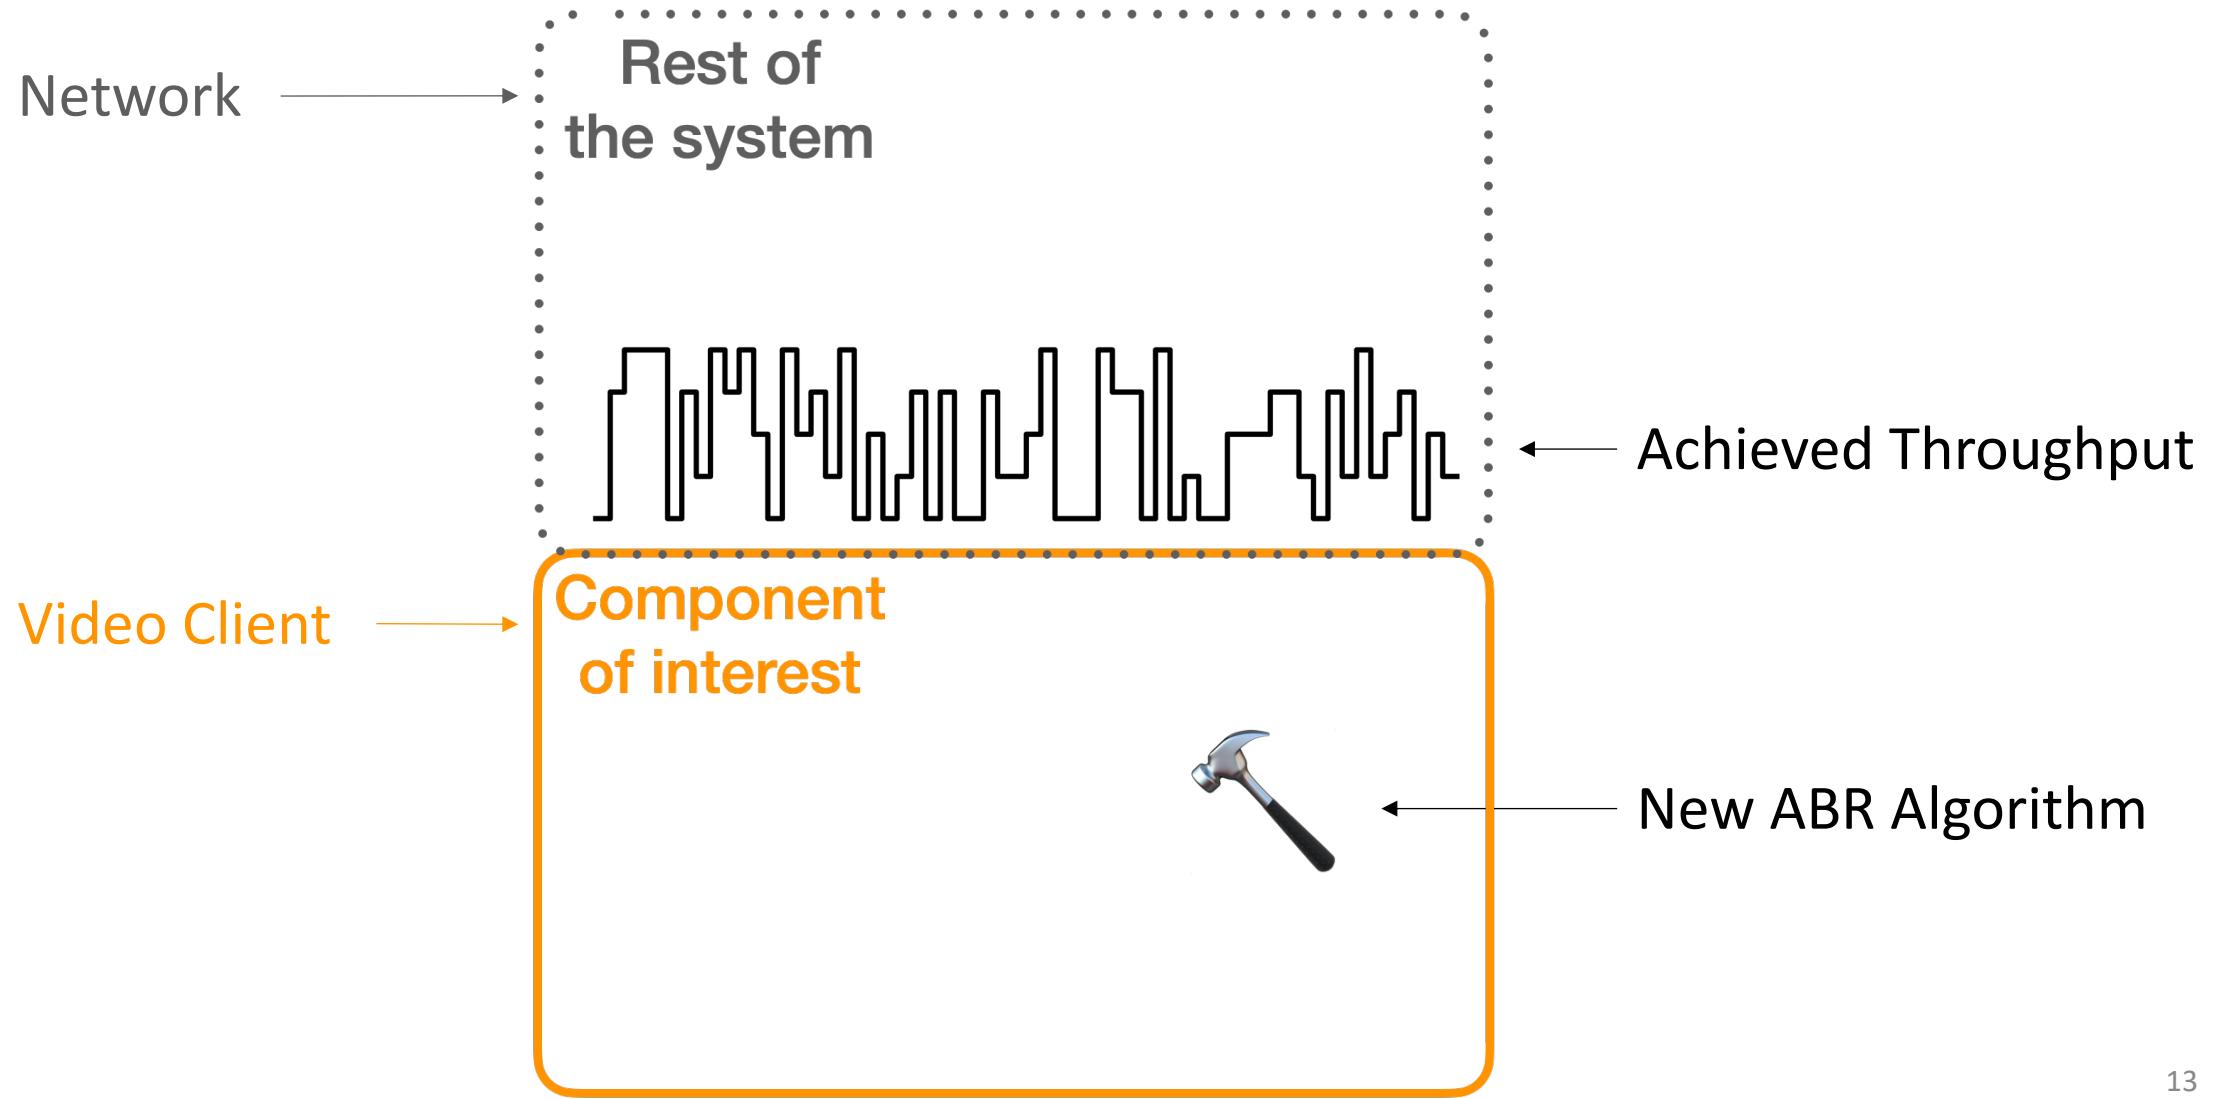
Exogenous trace assumption does not hold for many real-world traces

... hurts accuracy, can lead to completely wrong conclusions

Example: Adaptive Bitrate Video Streaming (ABR)



Trace-driven simulation for ABR



Trace-driven simulation for ABR

- “ExpertSim” [e.g., Yin et al. SIGCOMM’15, Mao et al. SIGCOMM’17]

$$b_{t+1} = \max \left(b_t - d_t , 0 \right) + T$$

playback buffer (in sec) →

chunk size (the new algorithm's choice)

download time

achieved throughput (from the replayed trace)

The Puffer dataset

- The Puffer Randomized Control Trial [Yan et al. NSDI 2020]
 - July 2020 – June 2021
 - 5 ABR algorithms
 - 56M downloaded chunks over 230K streaming sessions
 - 3.5 years worth of streamed video

Task: Given the traces for all except one ABR algorithm, simulate the held out algorithm on the same paths

A typical trace

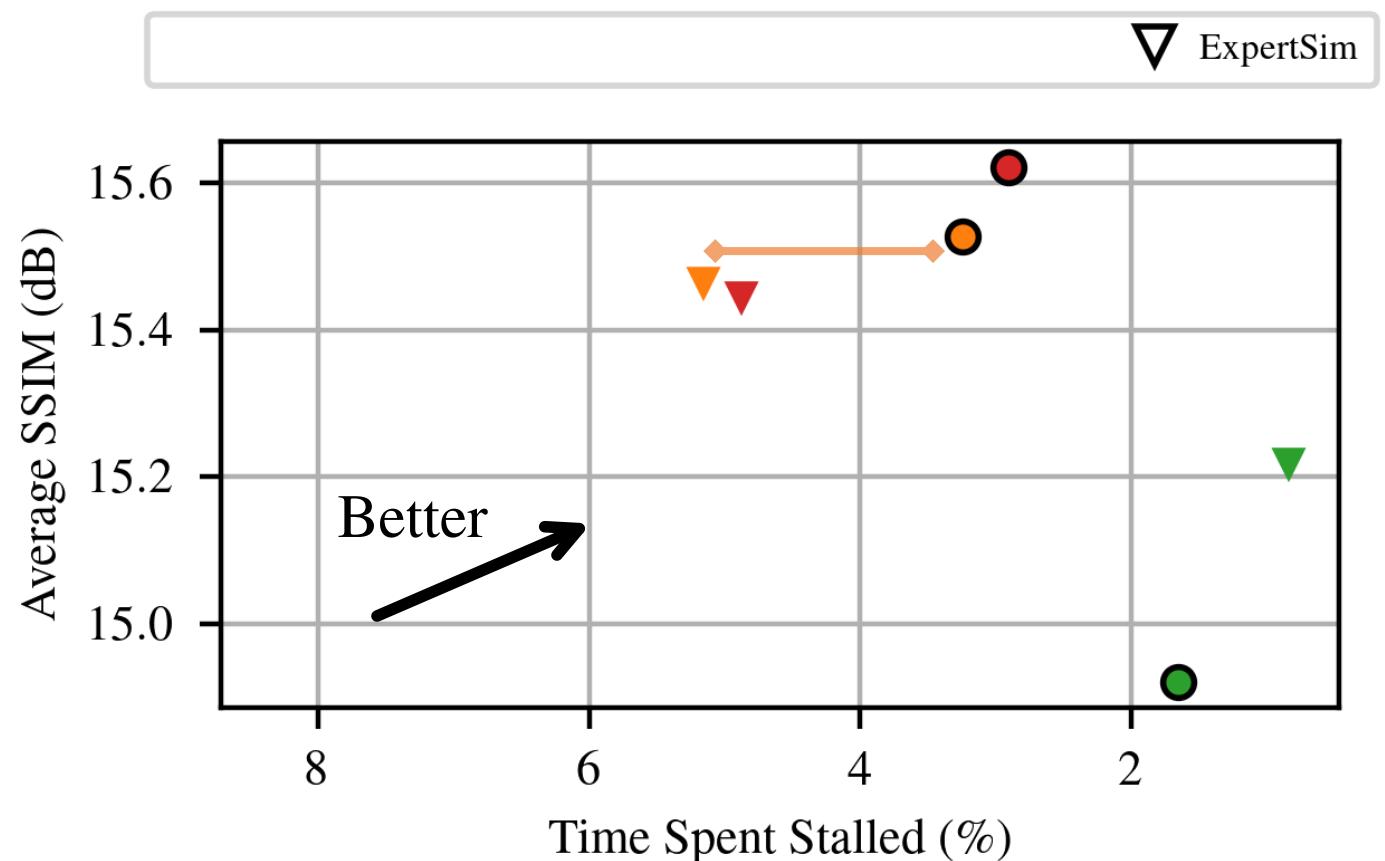
Chunk #	1	2	3	Algorithm
Bitrate	360p	480p	480p	BBA
Achieved Throughput	1Mbps	0.8Mbps	1.2Mbps	(“source”)
Playback Buffer	5s	3s	7s	

A simulated trajectory	Chunk #	1	2	3	Algorithm
	Bitrate	720p	?	?	BOLA1
	Achieved Throughput	?	?	?	(“target”)
	Playback Buffer	?	?	?	

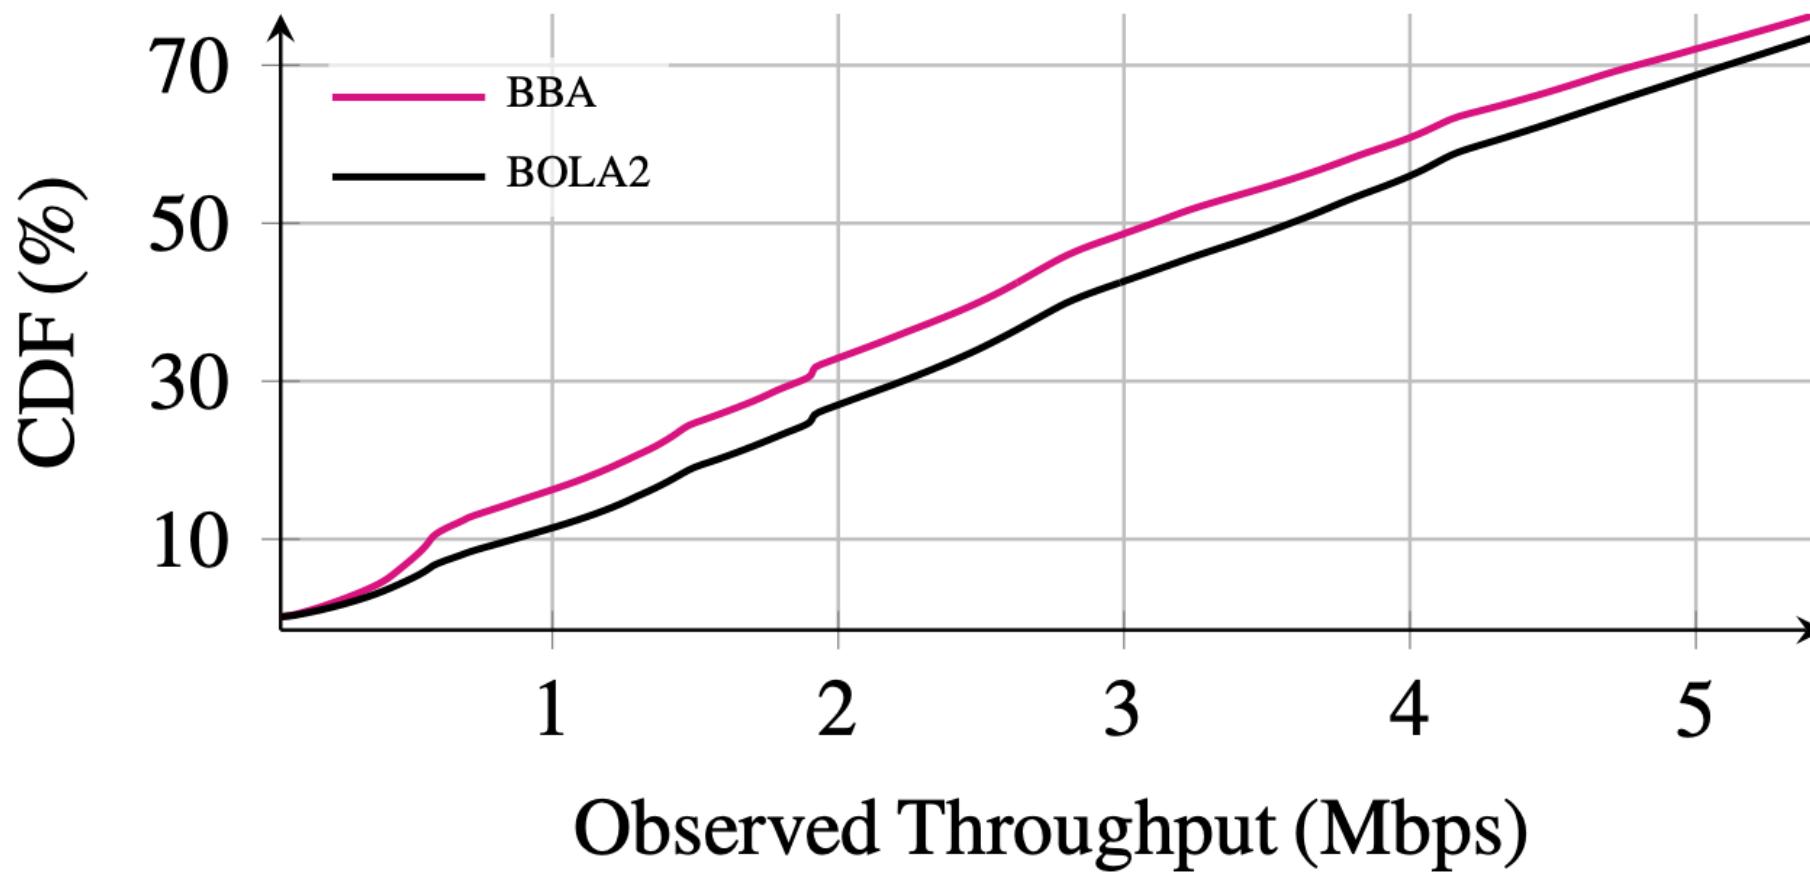
How accurate is trace-driven simulation?

- Source data: The Puffer Dataset with 5 algorithms
- Target algorithm (unseen): BBA, BOLA1, BOLA2

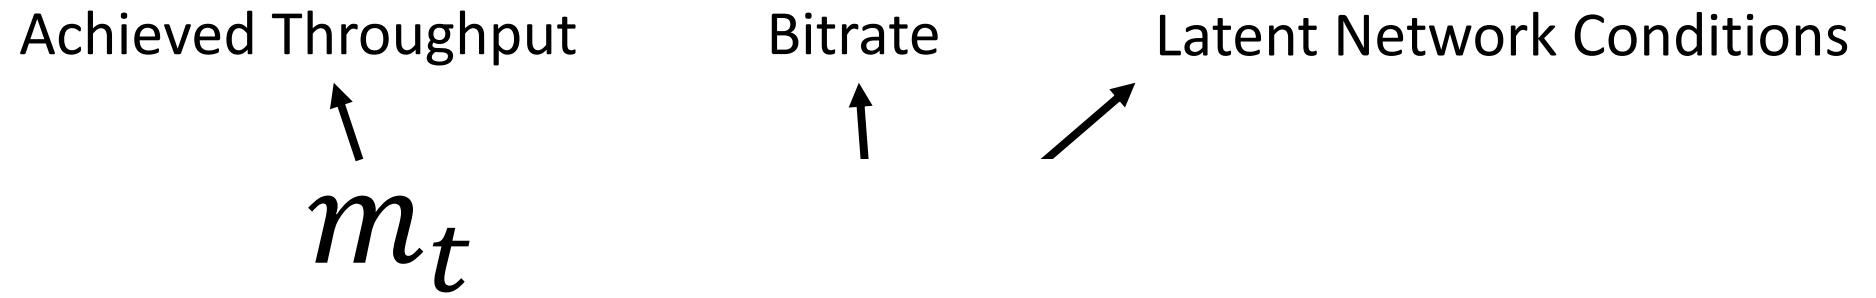
Exogenous trace assumption
Bitrates chosen by the ABR
algorithm affect the achieved
throughput



ABR algorithms affect throughput



Can we relax the exogenous trace assumption?



Both u_t and $f(\cdot)$ are unknown

Towards a solution

- If u and $f(\cdot)$ were known...

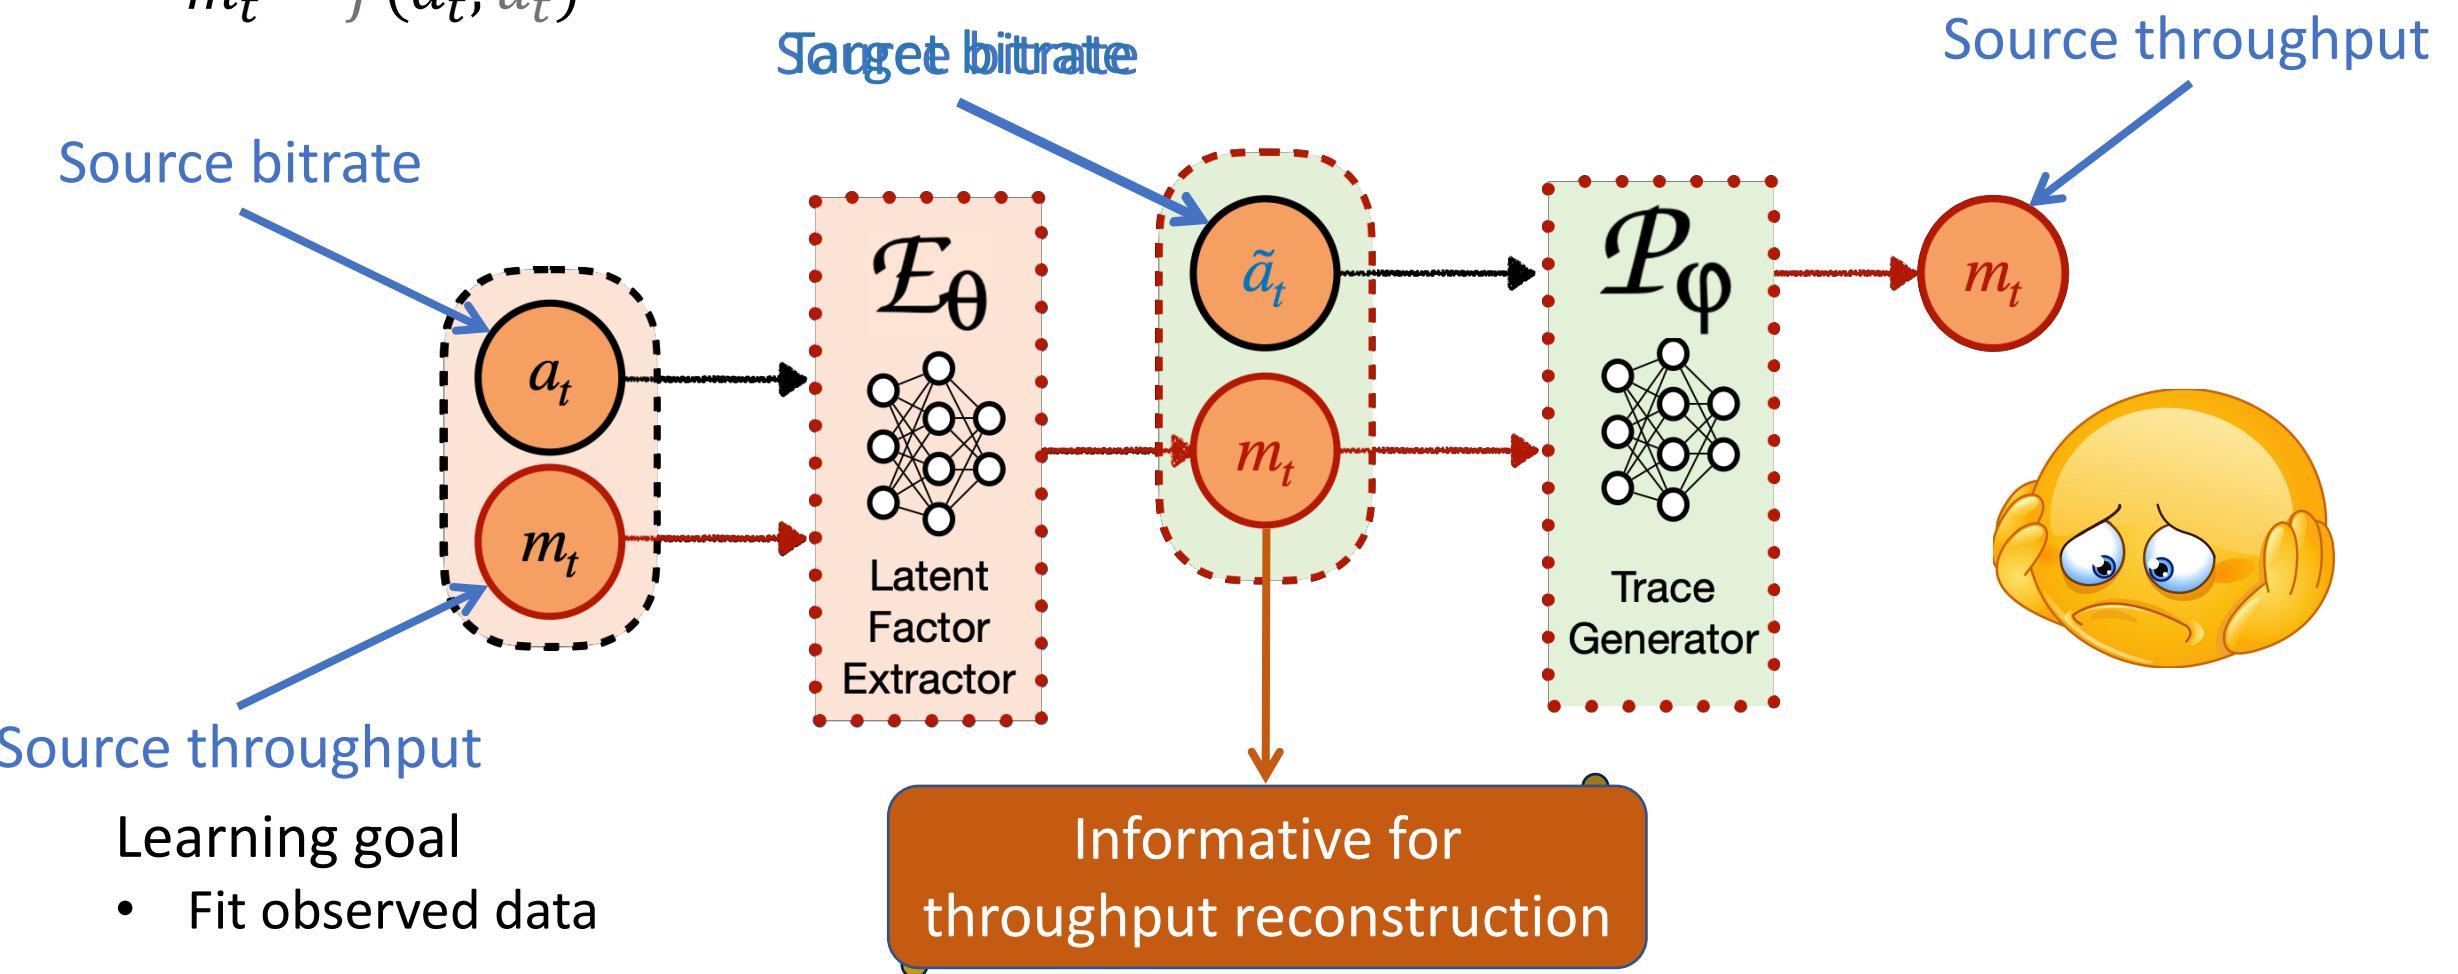
$$\tilde{m}_t = f(\tilde{a}_t, u_t)$$



simulated bitrate
(counterfactual)

A learning approach

$$m_t = f(a_t, u_t)$$



Learning goal

- Fit observed data

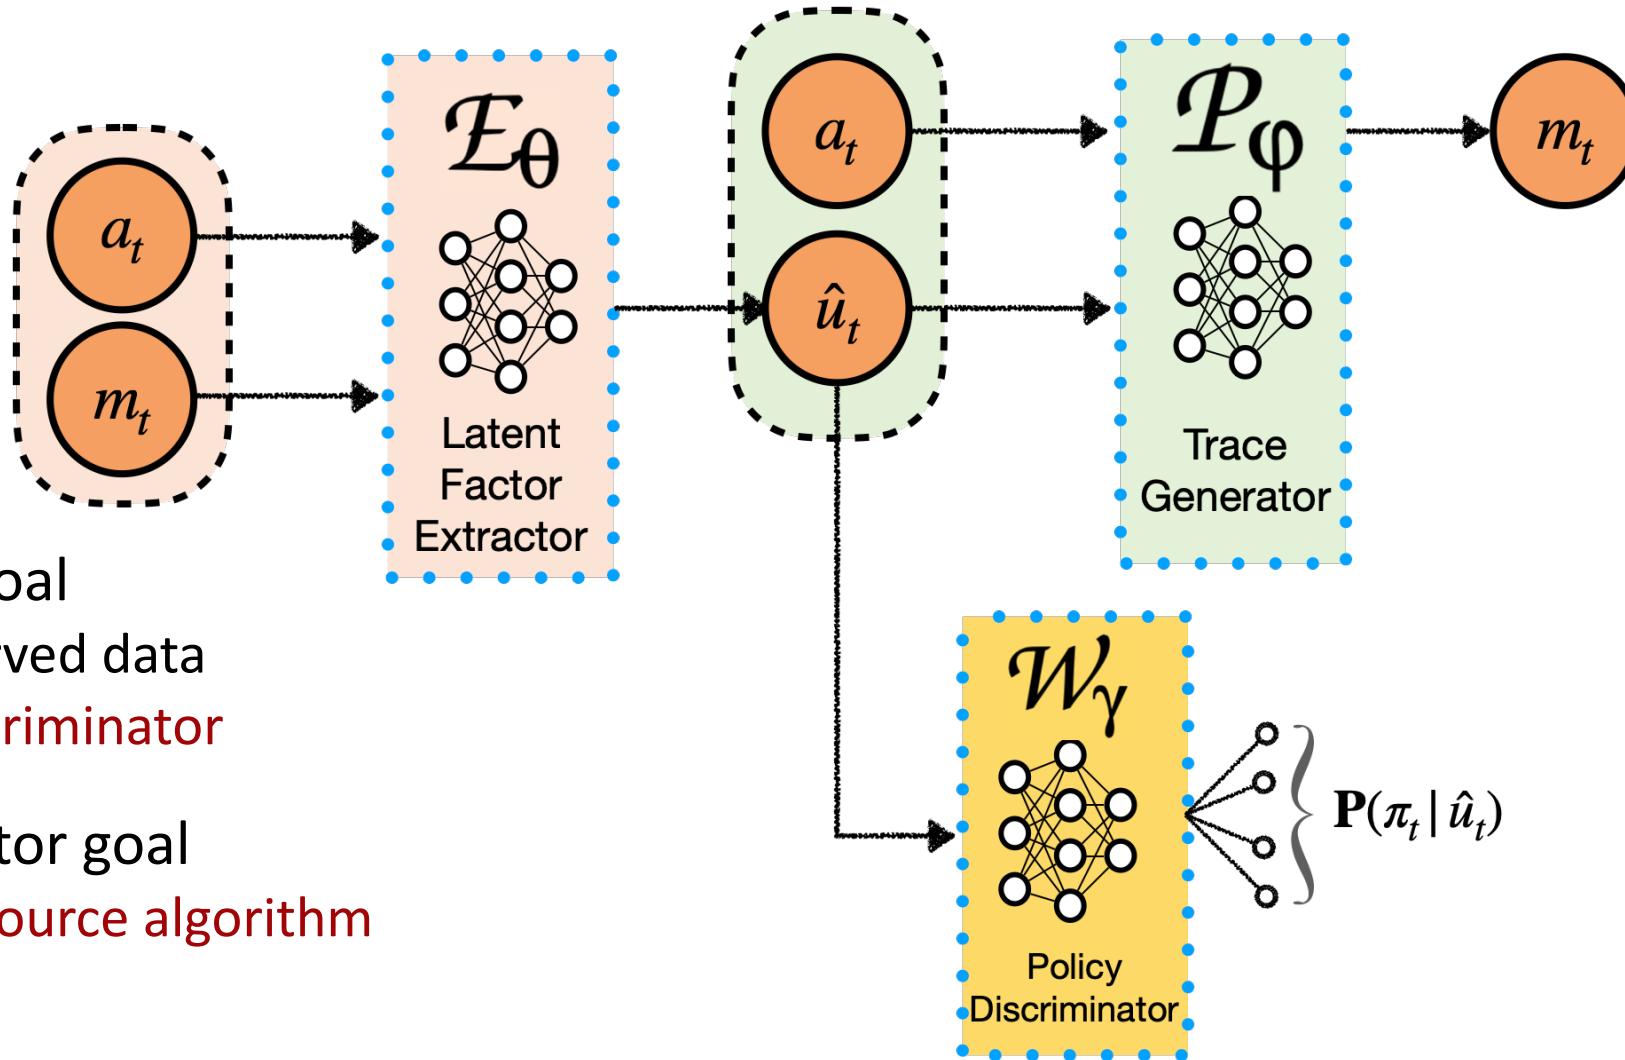
RCT to the rescue!



RCT property: Distribution of latent network conditions is the same in trajectories assigned to different algorithms.

- ↳ Latent network condition is independent of the source algorithm (used for tracing).
 - ↳ Latent network condition information about the source algorithm.
- Discriminator goal
- Predict source algorithm
-
- $\{ P(\pi | u_t)$
not give any
algorithm.

Exploiting the RCT property



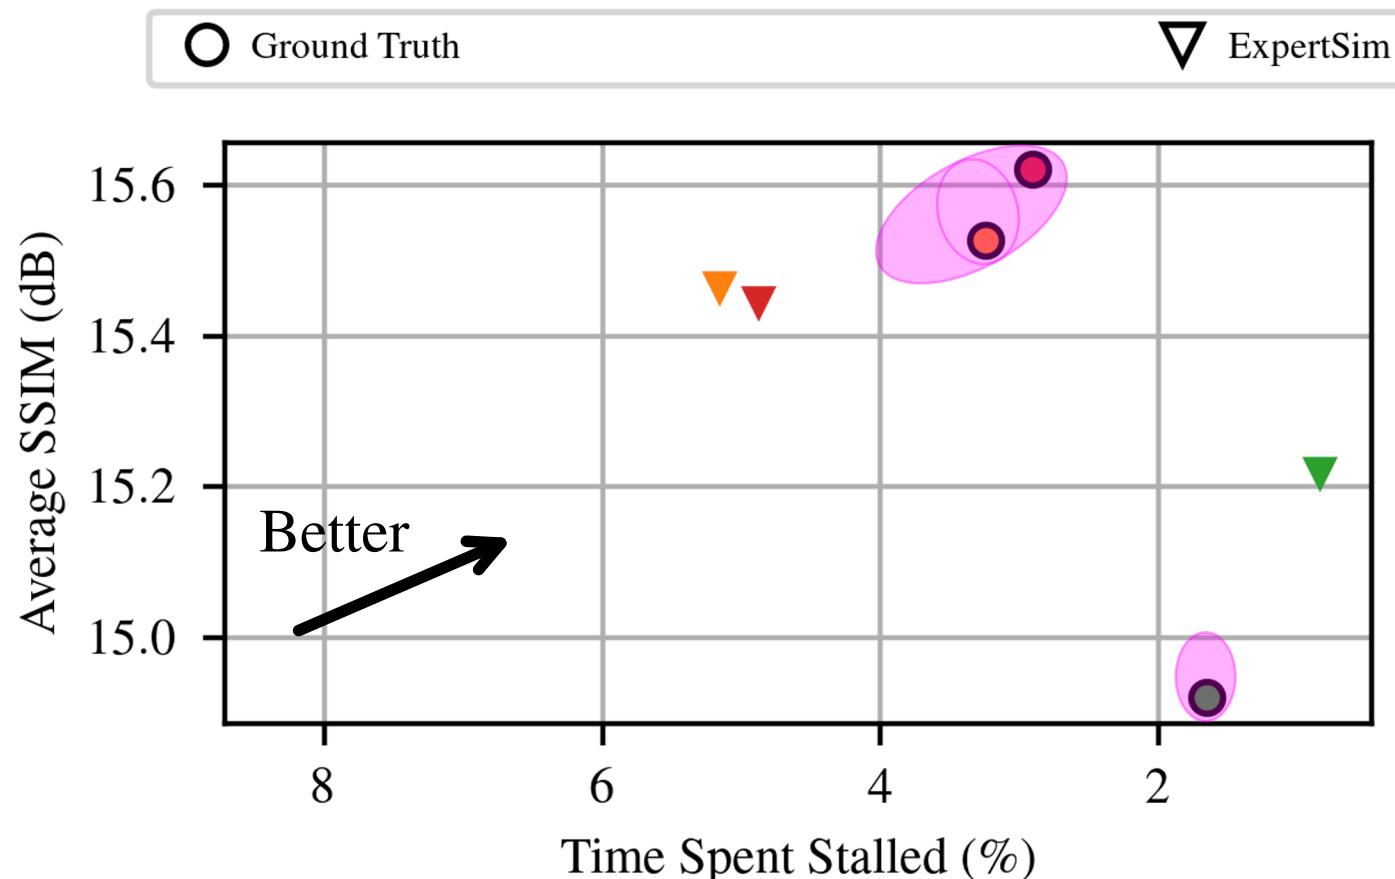
RCT property is sufficient for unbiased simulation

Theorem: The RCT property (independence of u and the algorithm) is sufficient for estimating the counterfactual trace if

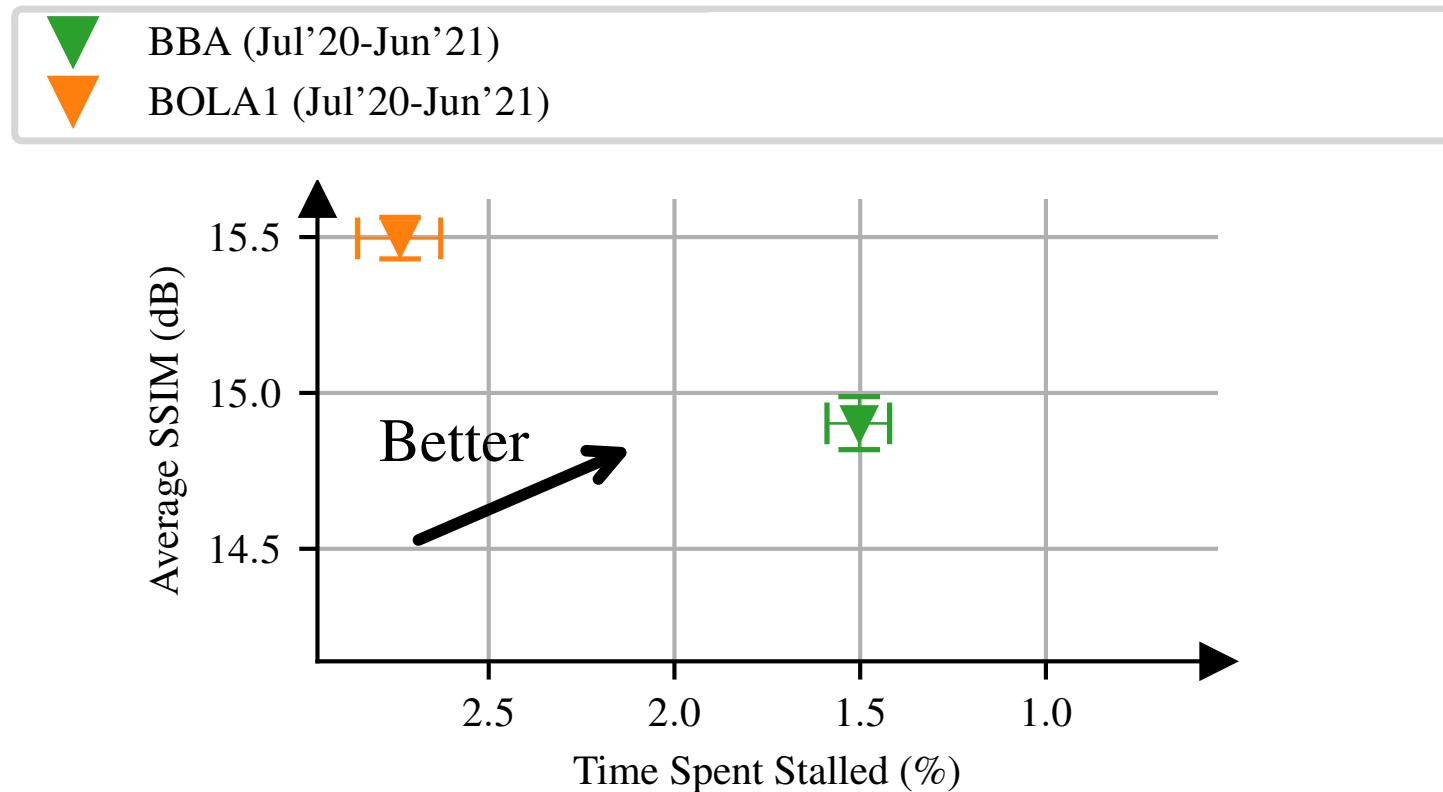
1. **(Invertibility)** $\forall a: m = f(a, u)$ is invertible.
2. **(Low-rank factorization)** Matrix representation of f has rank r , and $r \leq \text{dim}(\text{trace})$.
3. Traces are collected using sufficient number of **diverse** algorithms (See the paper for the precise statement).

Simulation accuracy

- Source data: The Puffer Dataset with 5 algorithms
- Target algorithm (unseen in training): **BBA**, **BOLA1**, **BOLA2**



Case study: CausalSim in the wild

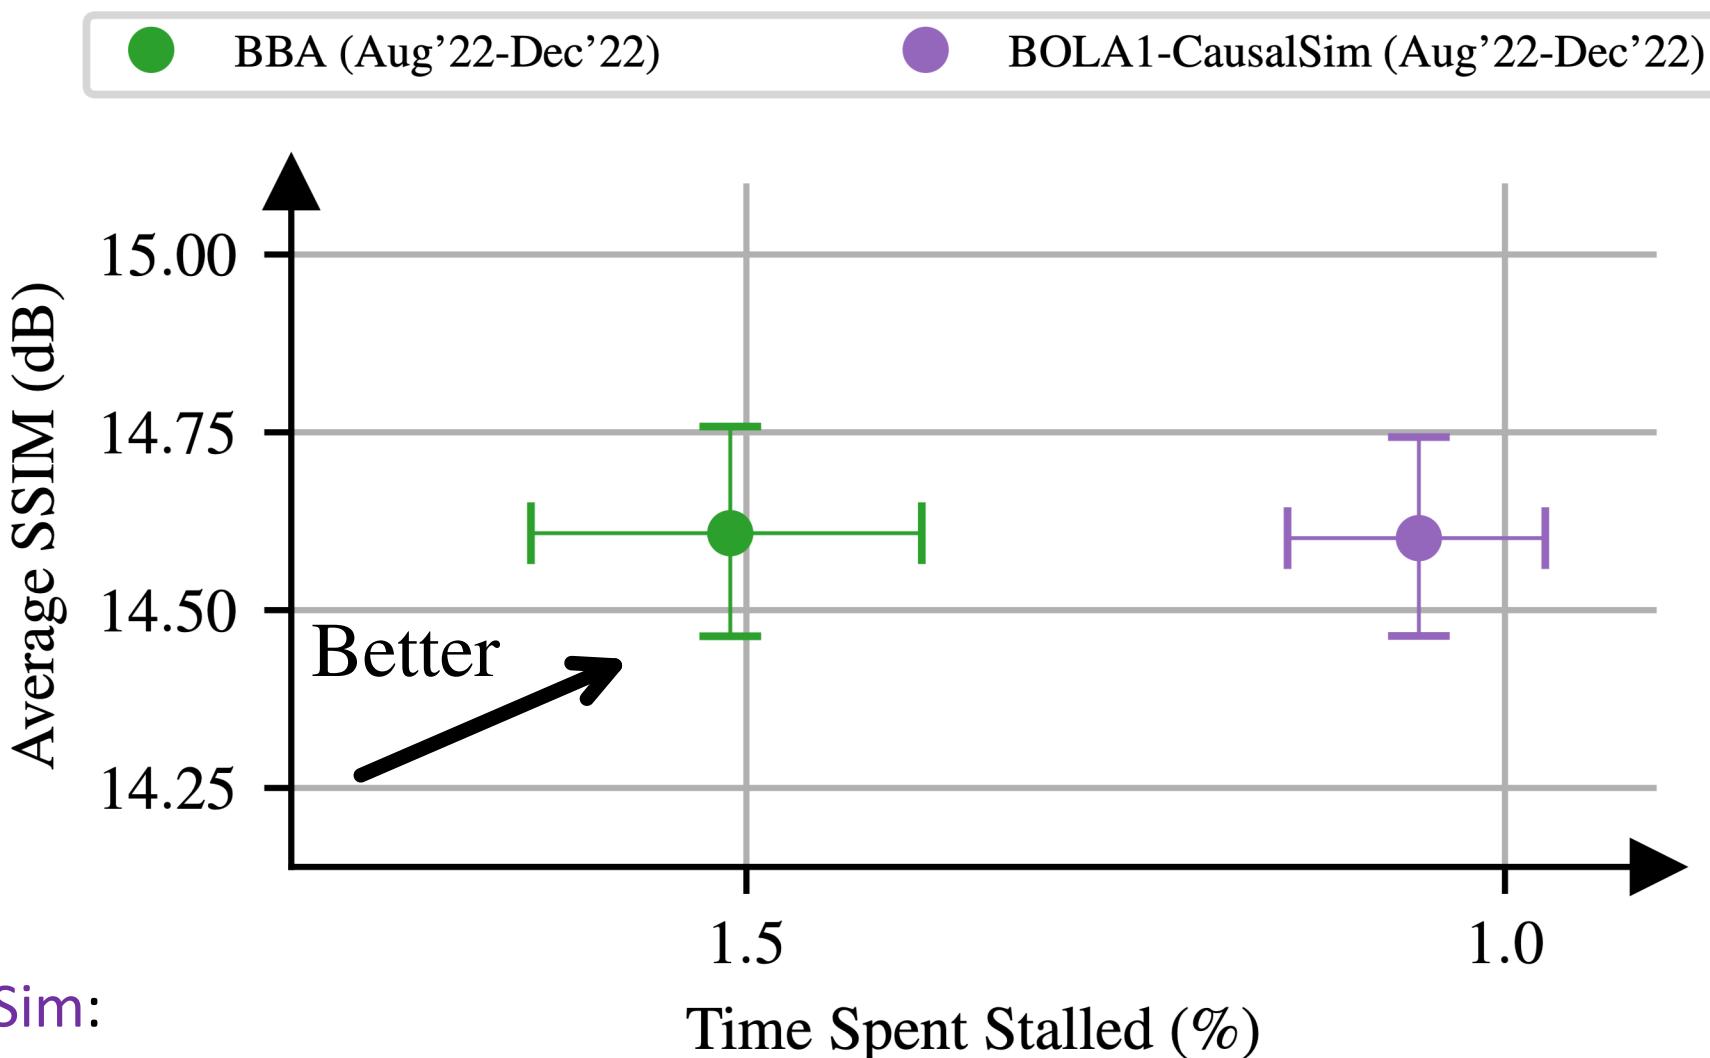


- Can we use CausalSim to improve algorithms?

Case study: CausalSim in the wild



Case study: CausalSim in the wild



BOLA1-CausalSim:

- Deployment: 1.4x less stalling than BBA

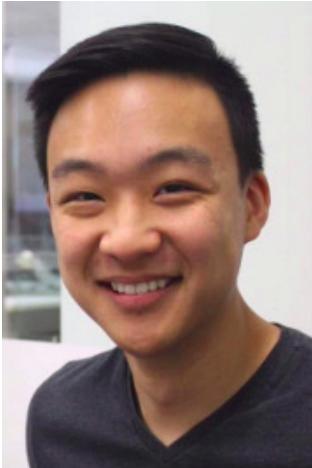
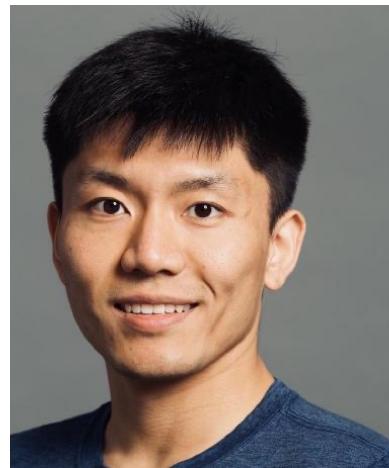
Summary

- ❑ The *Exogenous Trace Assumption* causes bias in trace-driven simulation
- ❑ **CausalSim** eliminates bias by modeling the effect of interventions on the trace
- ❑ We can use **CausalSim** to optimize and deploy real-world ABR algorithms



causalsim.csail.mit.edu

m3: Accurate Flow-Level Performance Estimation using Machine Learning



Chenning
Li*

Arash
Nasr-Esfahany*

Kevin
Zhao

Kimia
Noorbakhsh

Prateesh
Goyal

Mohammad
Alizadeh

Thomas
Anderson

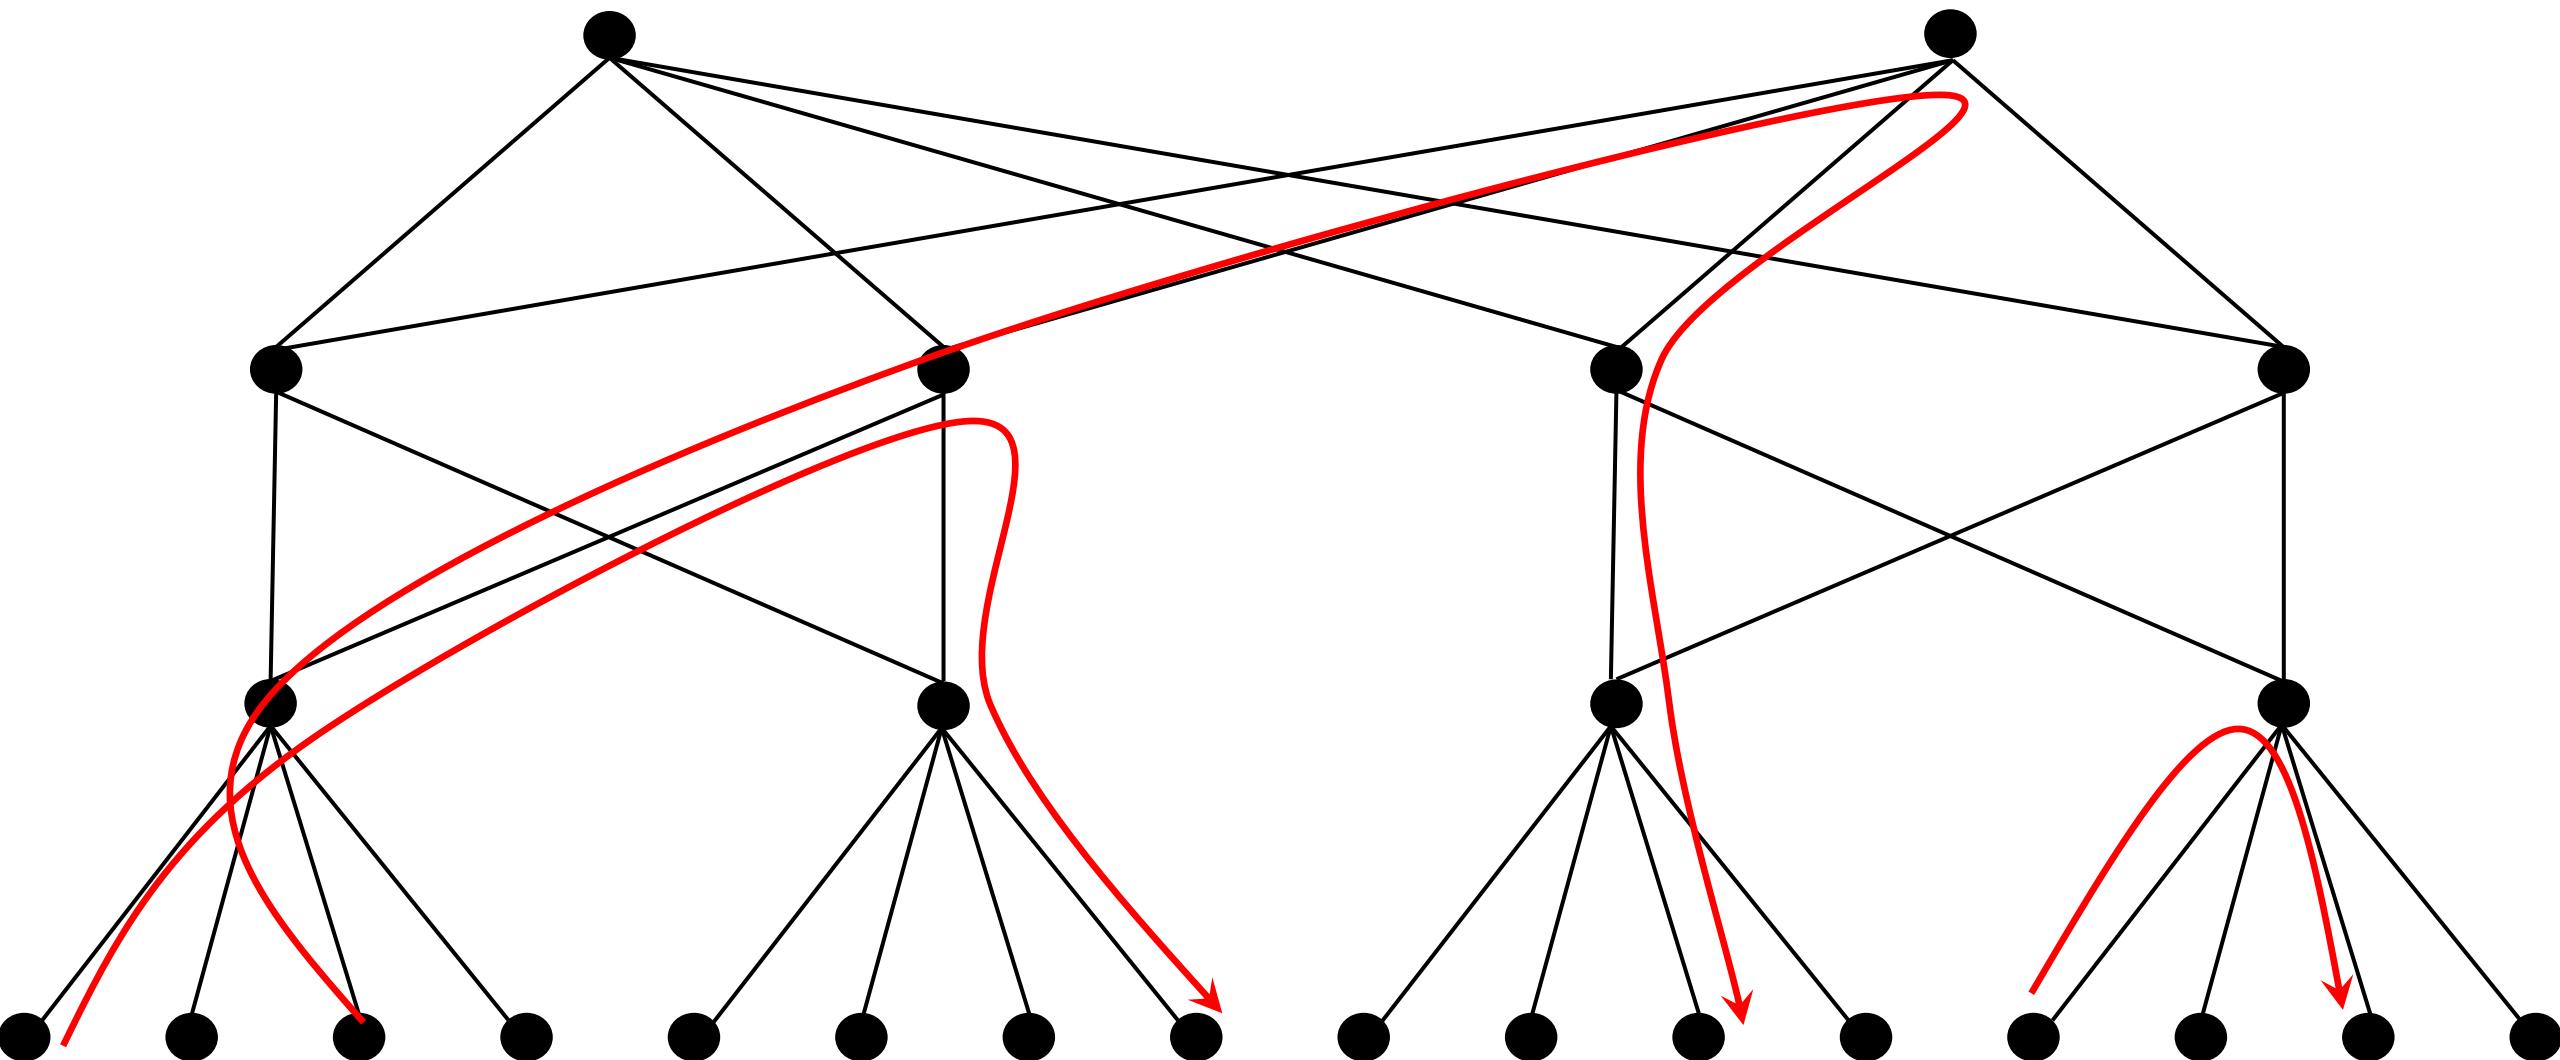


Massachusetts
Institute of
Technology

W UNIVERSITY of WASHINGTON

Microsoft®
Research

Predicting datacenter network performance

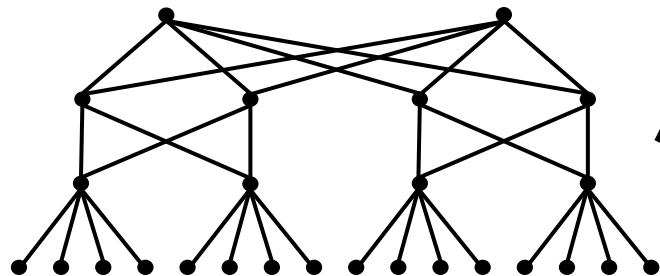


Network topology/configs

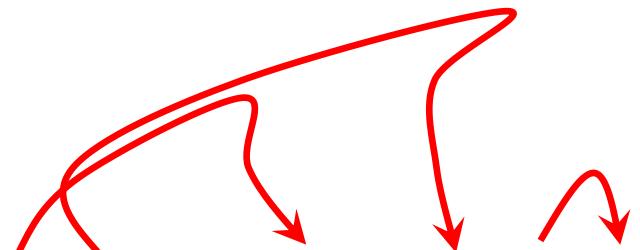
Workload

Existing tools for network perf. prediction

Network topology/configs



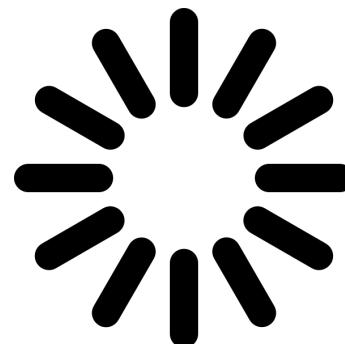
packet loss ratio
buffer occupancies
for specific flows at
specific times
flow jitter



Workload

But ns-3 is very slow...

~6,000 hosts
~ 10 millions of flows



~10 hours for 5s
simulated time!

Prior work on faster simulation

- Exact parallel simulation
 - ❑ ns-3
 - ❑ OMNET++
 - ❑ DONS, SIGCOMM'23
 - Approximative simulation
 - Parsimon, NSDI'23
 - DeepQueueNet, SIGCOMM'22
 - MimicNet, SIGCOMM'21
-  All are packet-level simulators → slow for large-scale networks, especially as the network becomes faster

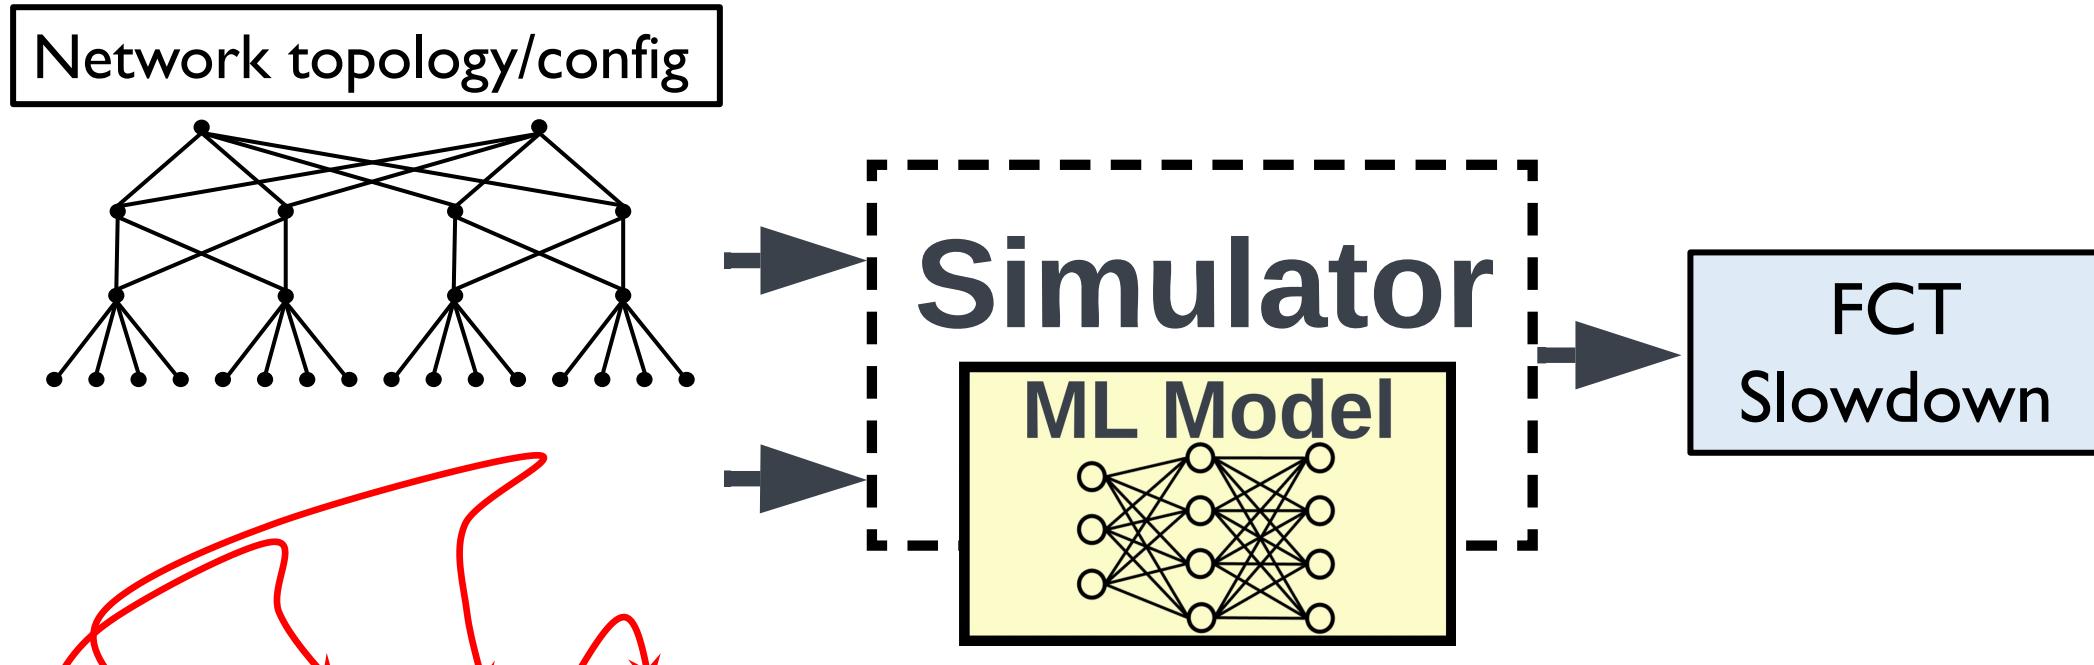
Do we always need packet-level simulation?

- Many questions are about **aggregate performance statistics**
 - Tail latency
 - Throughput
 - Loss rates
 - Short & long flow completion time
- How would workload/design changes impact these statistics?



Can we elevate the abstraction level of answering such questions without compromising fidelity?

Abstract the network simulator as a function

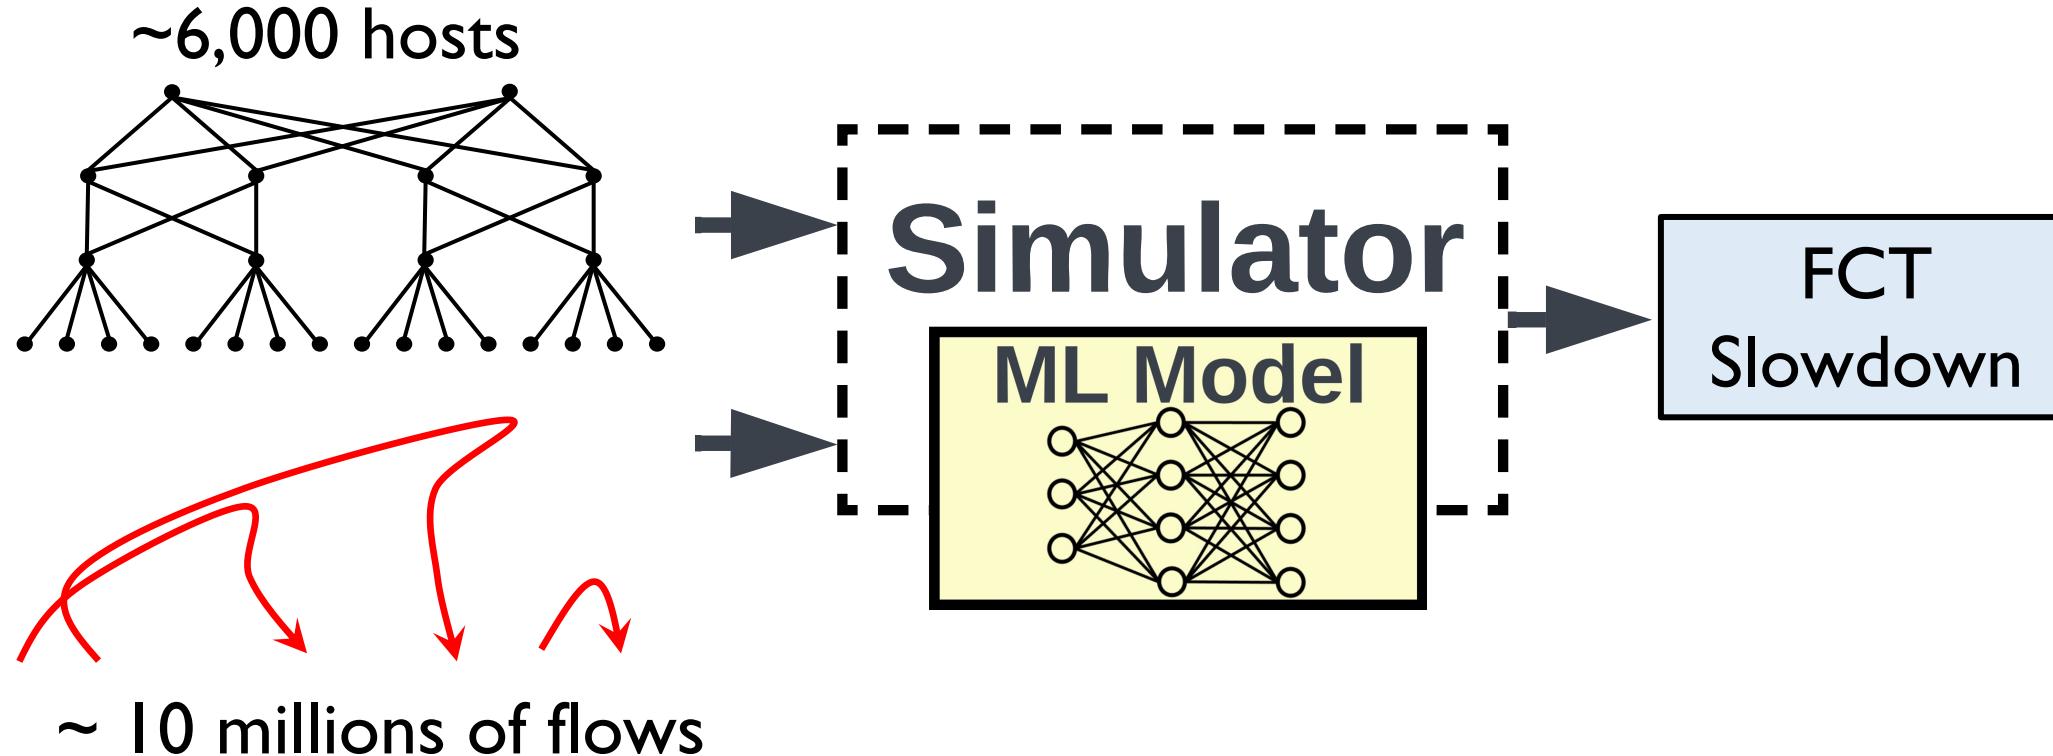


FCT slowdown: The ratio of the flow completion time (FCT) to the ideal completion time for a flow in an unloaded network



Use supervised learning to learn the end-to-end function?

The simulator is a **complex** function

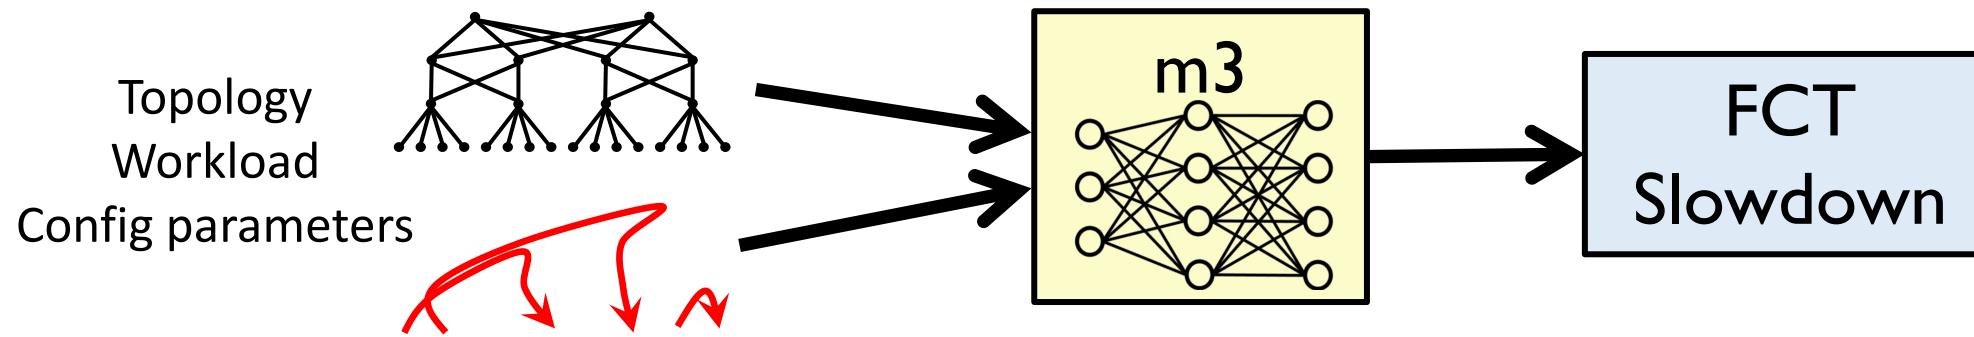


Hard to represent the function in a compact way



Slow to generate the training dataset for the ML model

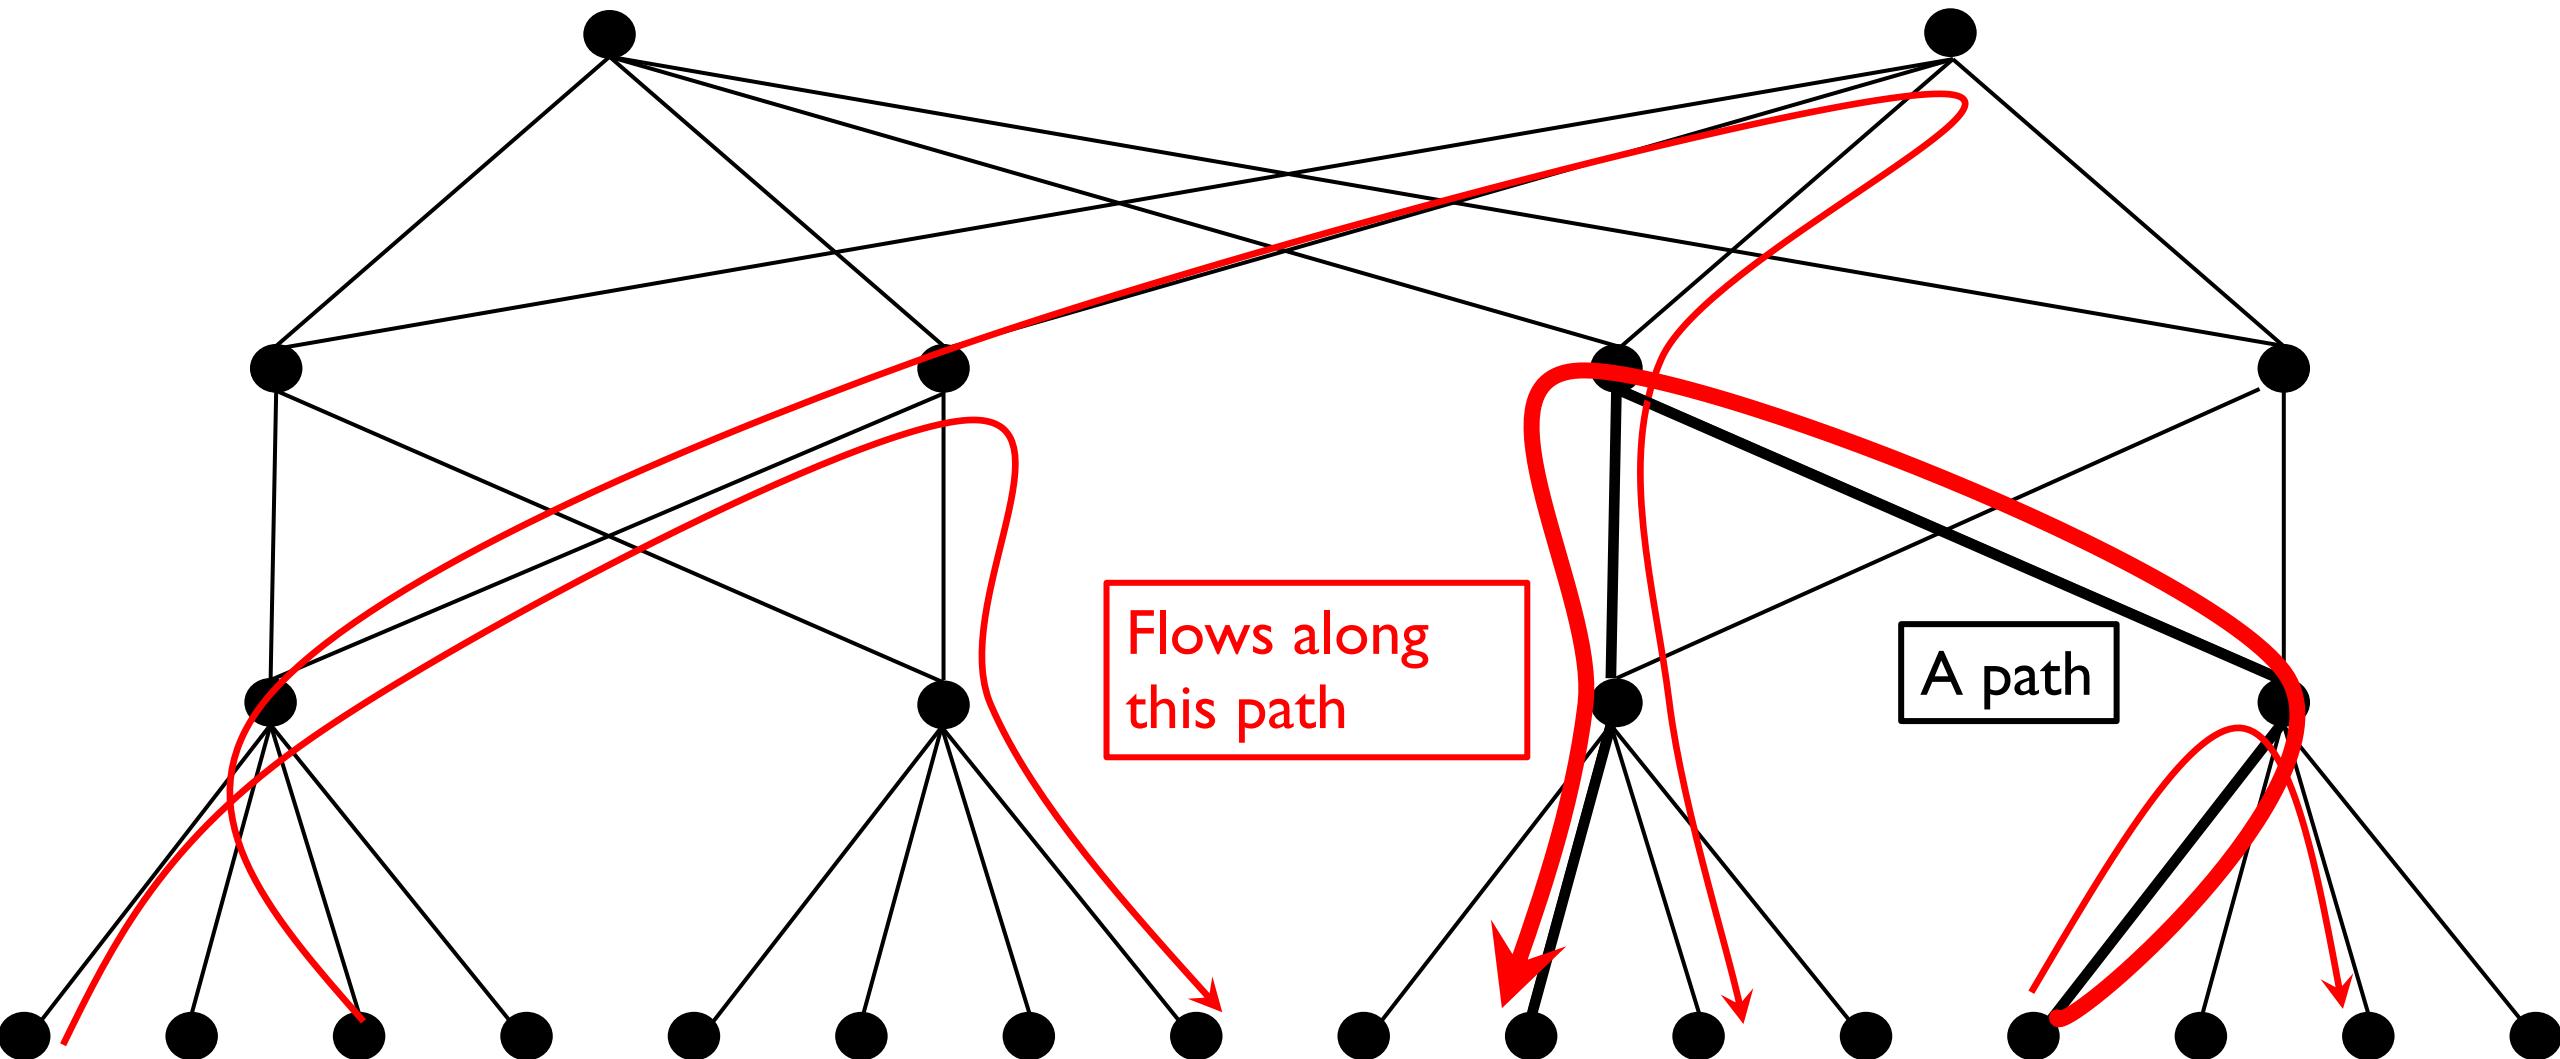
m3: Accurate flow-level performance estimation using machine learning



m3: ~1200X speedup with only ~10% estimation error in a 384-rack, 6144-host topology (vs. ns-3)

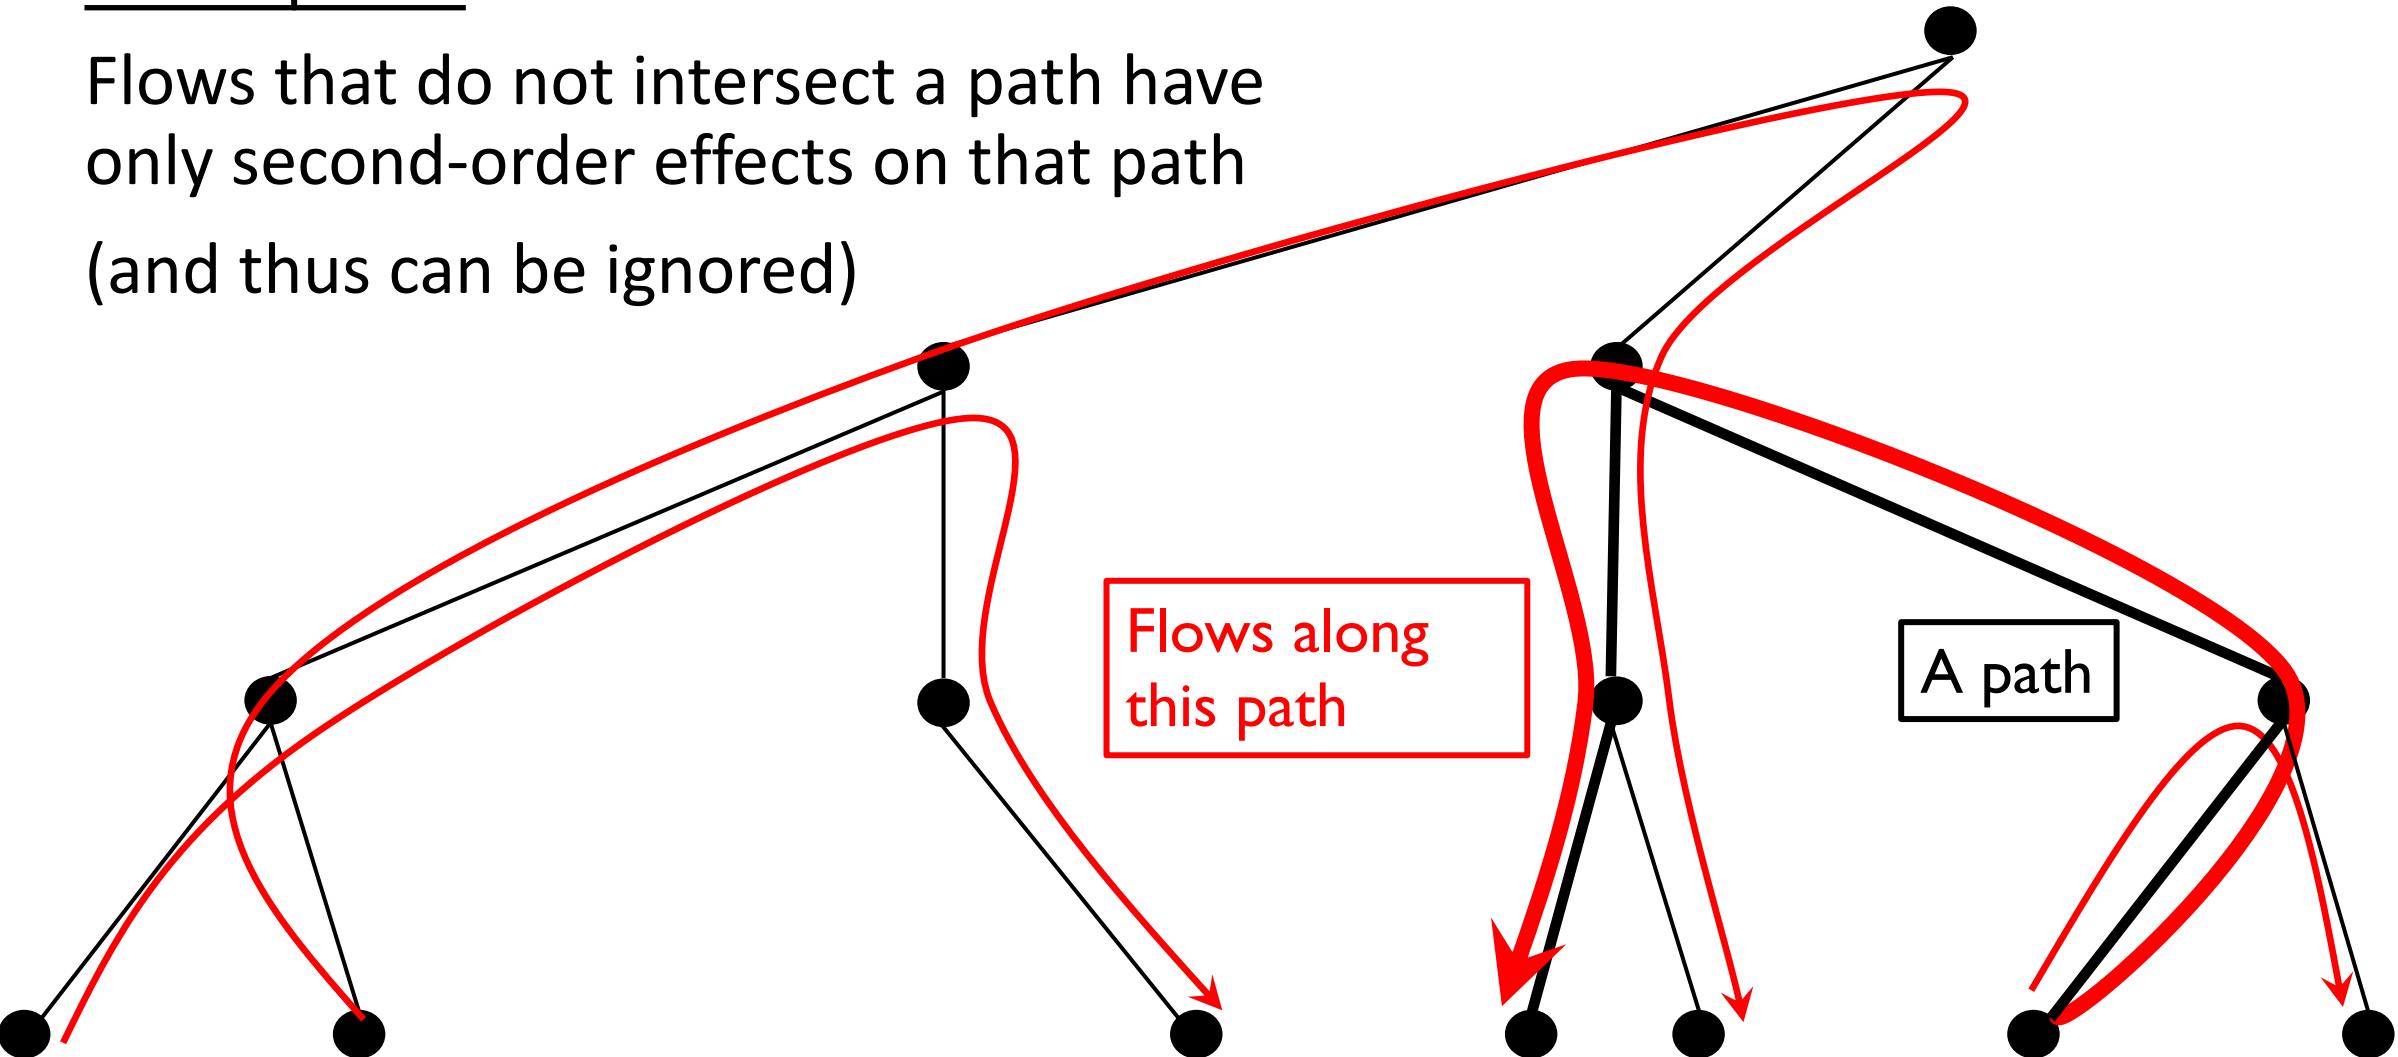
From ~10 hours (ns-3) to less than 1 min (m3)

Building block: A path simulation



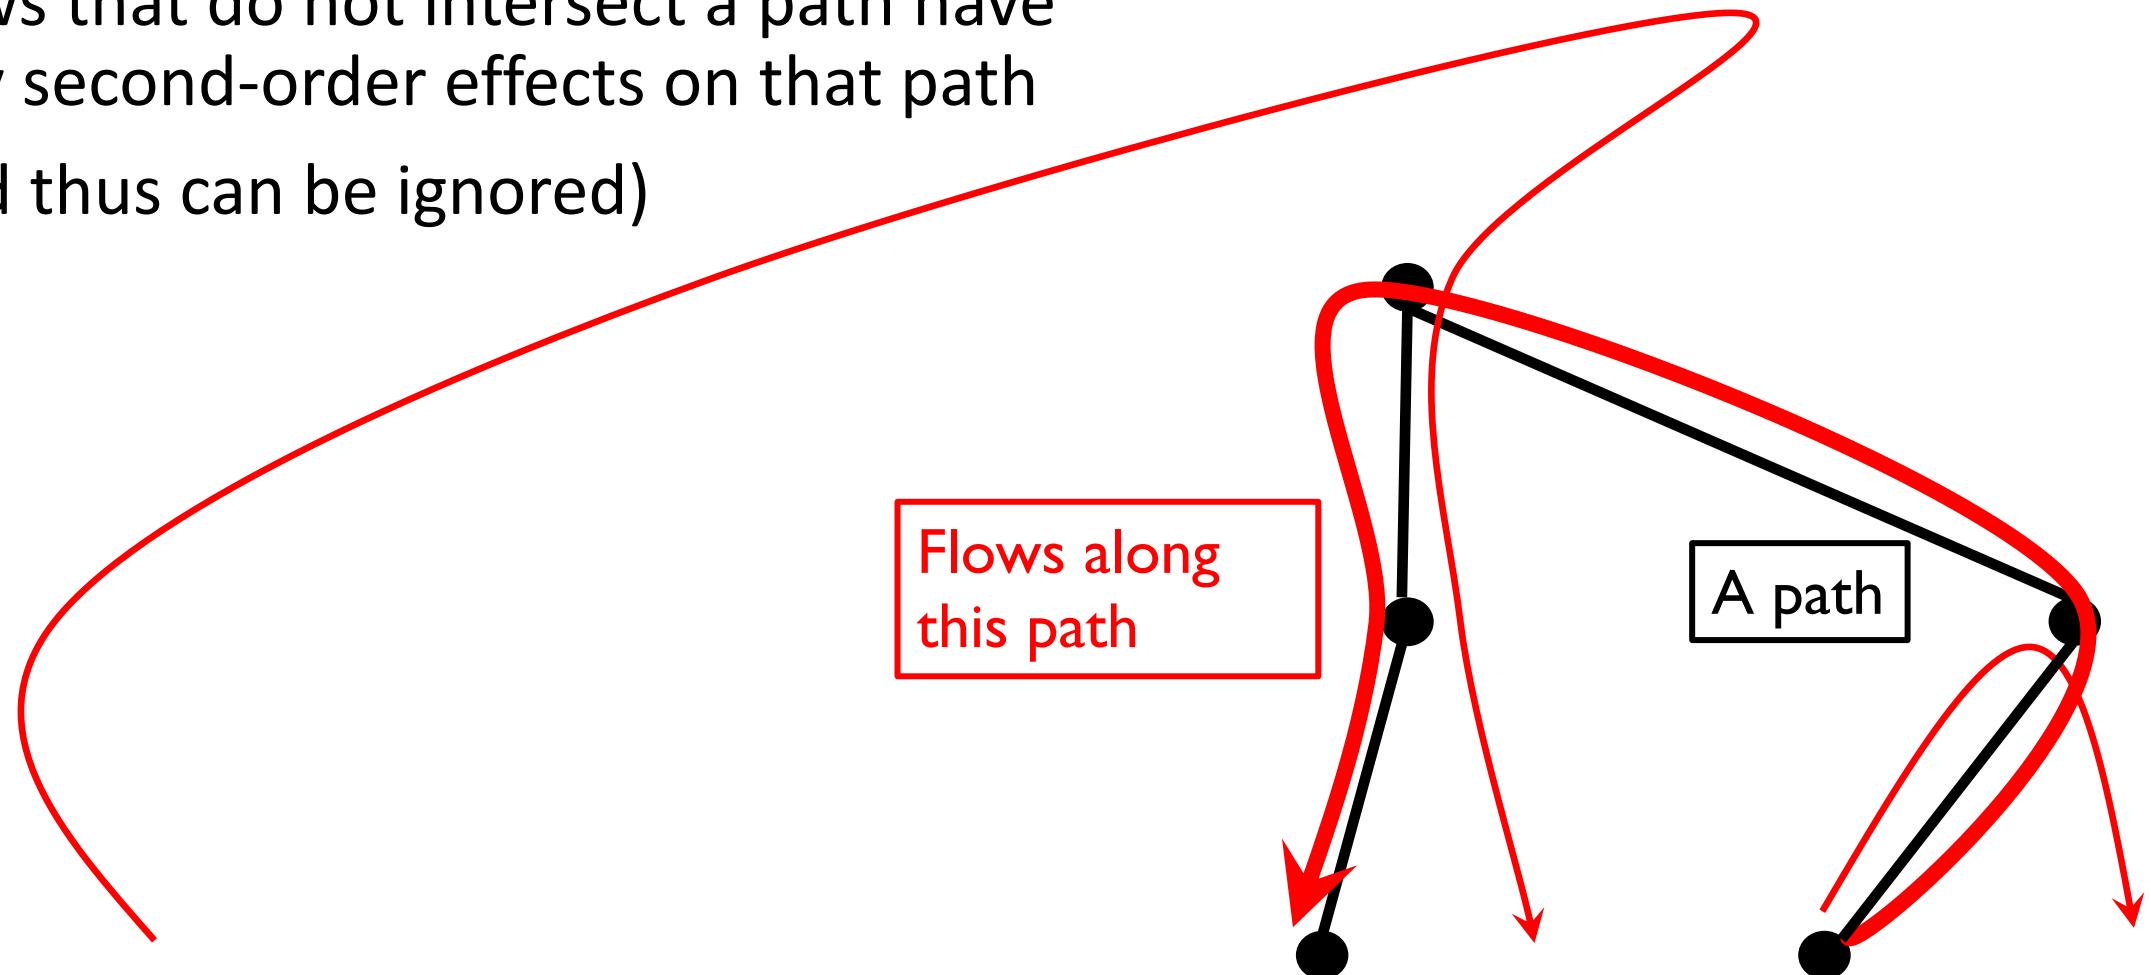
Assumption:

Flows that do not intersect a path have
only second-order effects on that path
(and thus can be ignored)



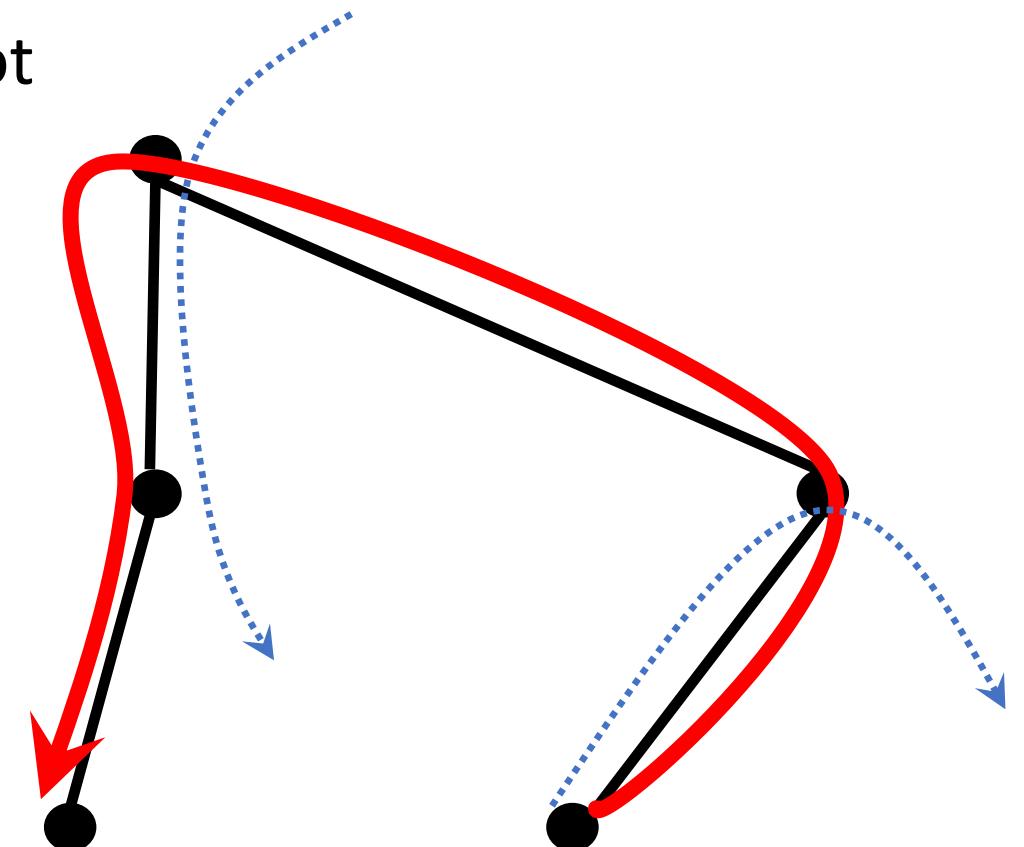
Assumption:

Flows that do not intersect a path have
only second-order effects on that path
(and thus can be ignored)

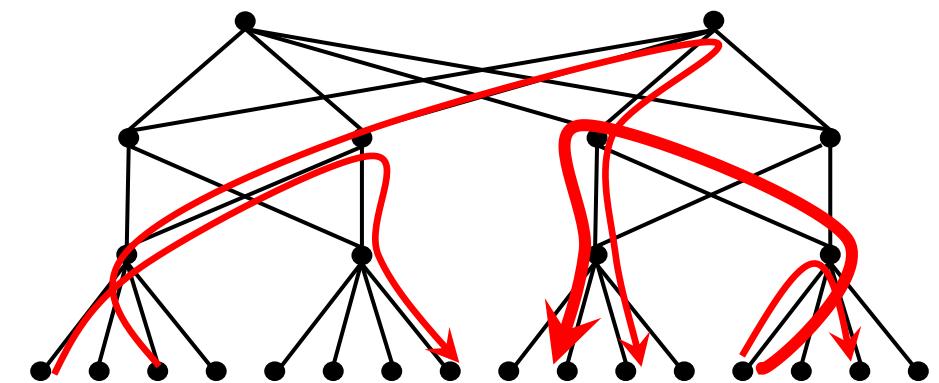


Path-level simulation

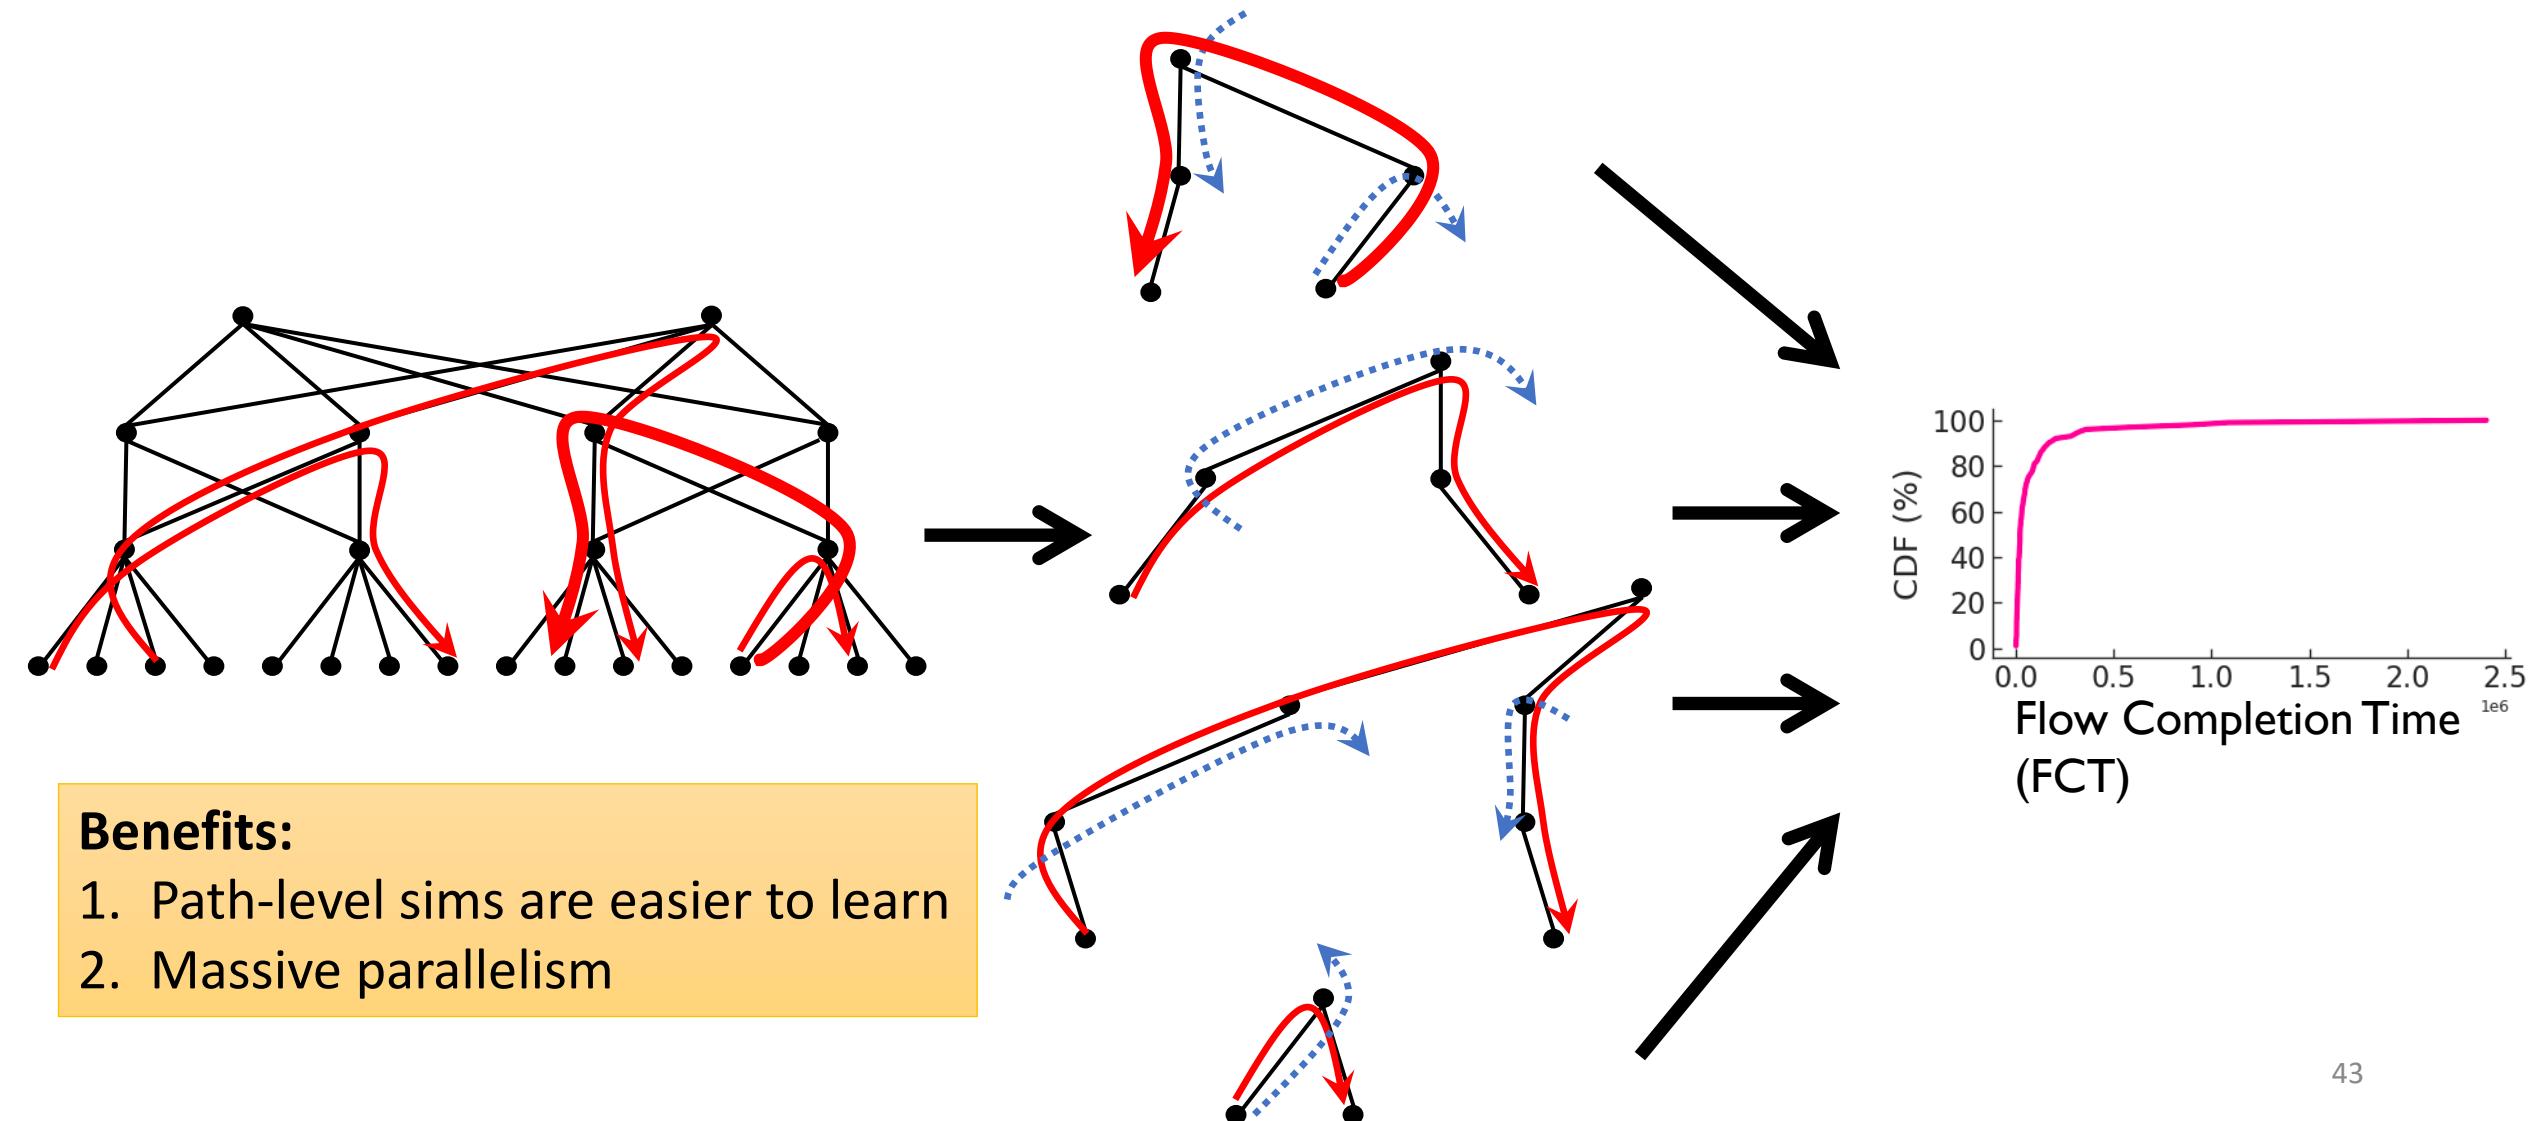
- Select one path and prune flows that do not intersect it
 - ❑ **Foreground flows** follow the entire path
 - ❑ **Background flows** intersect only certain links along this path, not all of them
- Goal: estimate the FCT slowdown distribution of the **foreground flows**



Path decomposition & aggregation



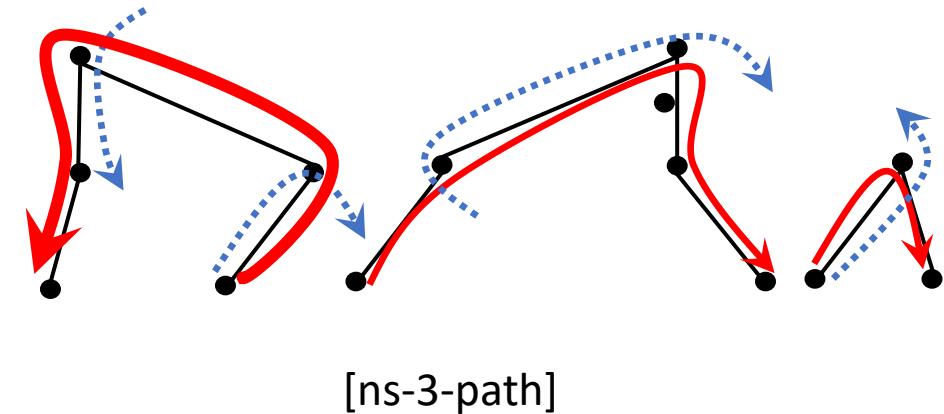
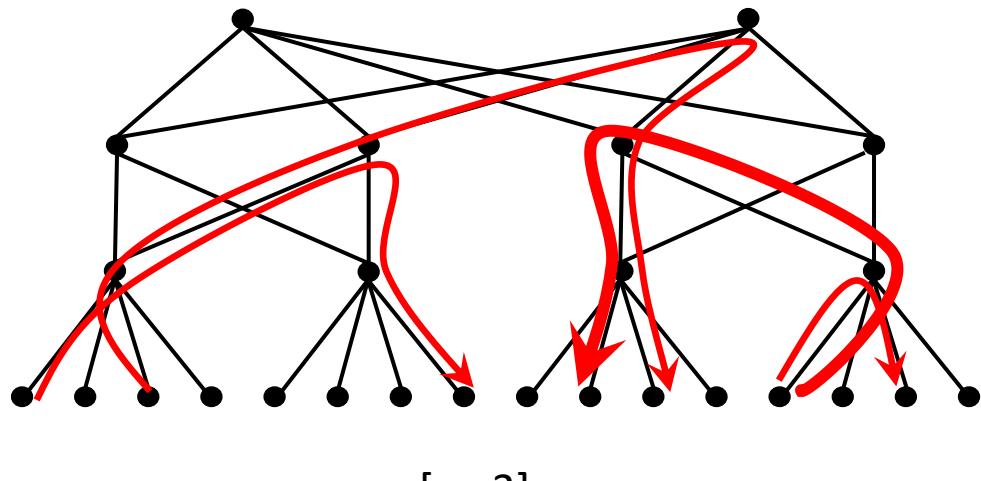
Path decomposition & aggregation



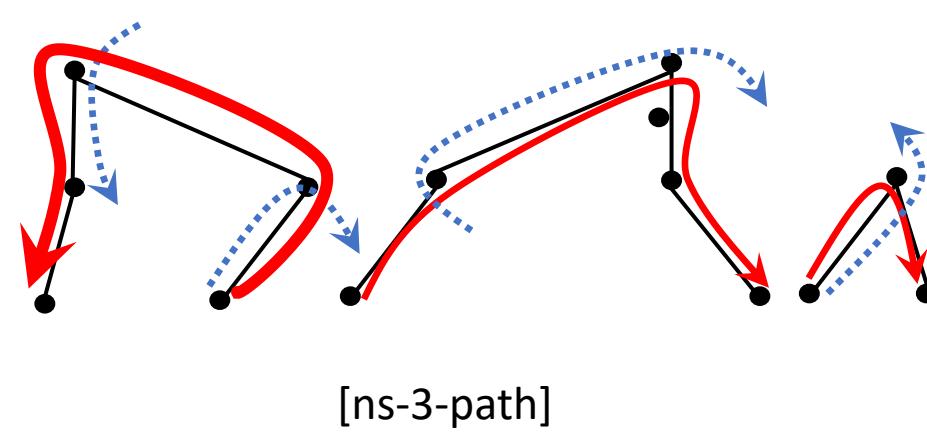
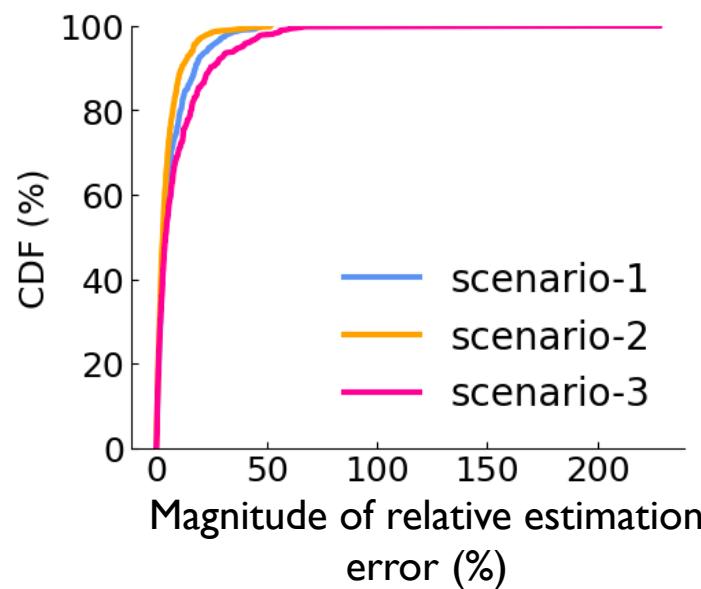
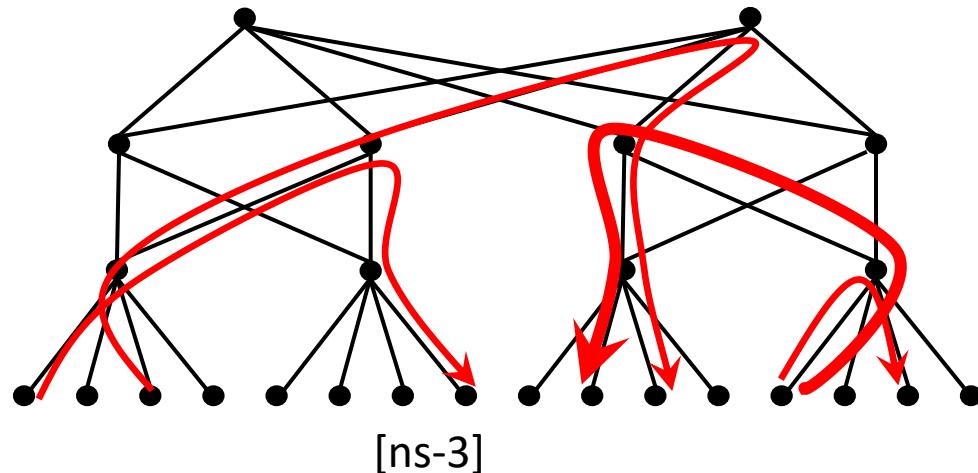
Is path decomposition accurate?

Experiment

- Three scenarios on a 32-rack 256-host topology with different workloads
- Randomly select 500 paths for each scenario
- Compare the **tail latency** (P99 FCT slowdown) of foreground flows for each path

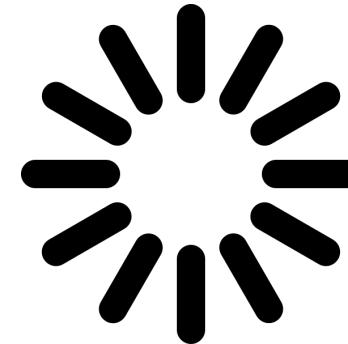
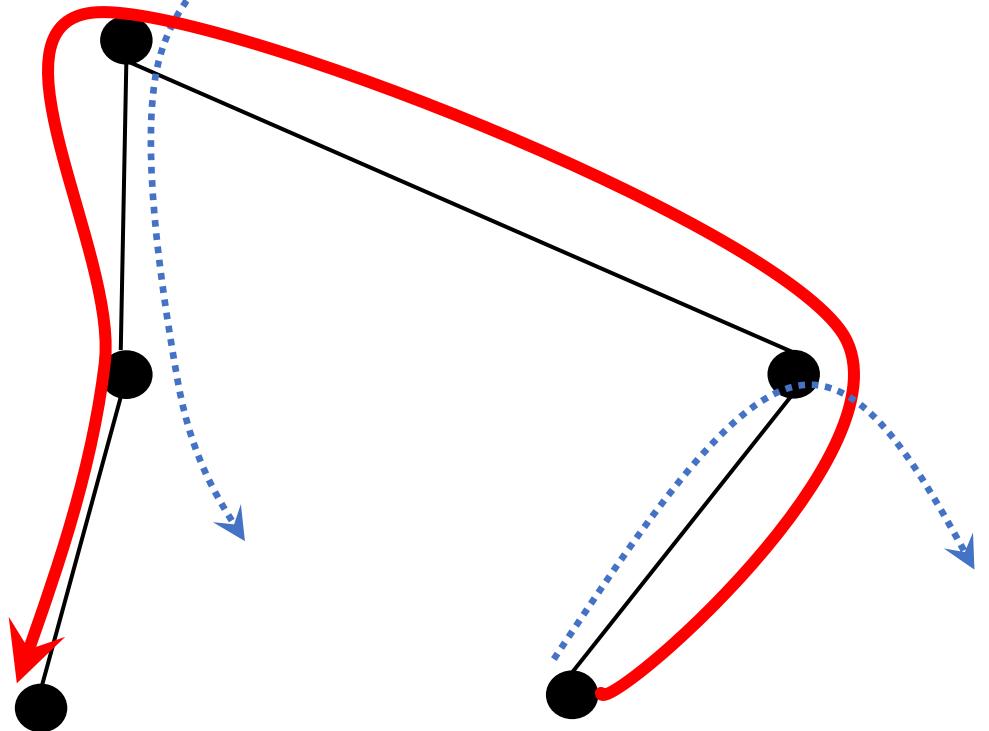


Path-level sim is accurate but still slow



Scenario	Median relative estimation error	Overall speedup
1	-0.83%	3.6X
2	-2.13%	5.4X
3	-5.91%	14.2X

Estimate path-level perf. with ML



~ 10 K foreground flows
~ 1 million background flows



m3 uses a **flow-level simulator** to extract compact features for its ML model

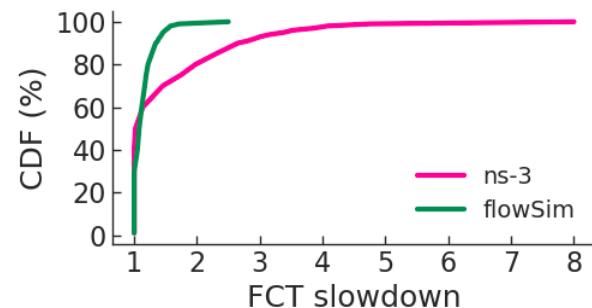
Flow-level simulator (flowSim)

- Max-min flow-level simulation
 - ❑ No packets – A “fluid” flow i sending at time-varying rates $r_i(t)$
 - ❑ Flows transmit at the max-min fair rates^[1] (approximates congestion control)
 - ❑ Events: flow arrivals and departures
 - ❑ No queues!

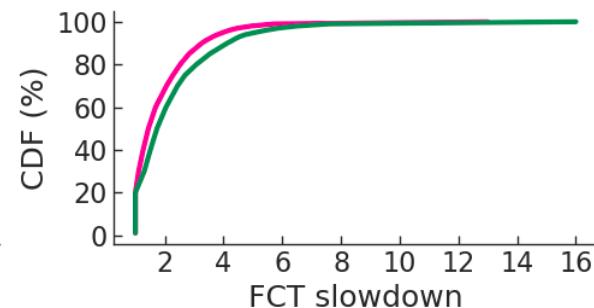


Fast: <1 sec for a path-level sim with
1 million flows

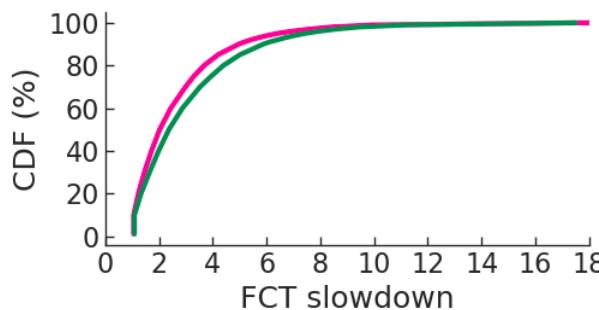
flowSim is fast but inaccurate



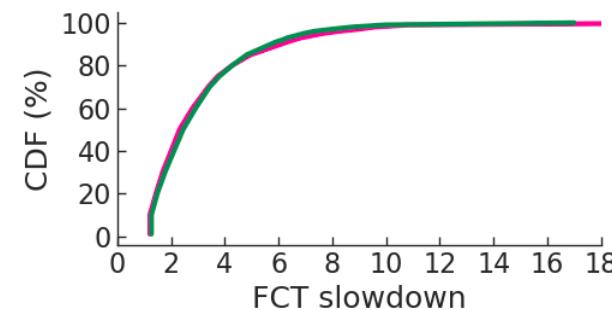
(a) flow size $\in (0, 1\text{KB}]$



(b) flow size $\in (1\text{KB}, 10\text{KB}]$



(c) flow size $\in (10\text{KB}, 50\text{KB}]$



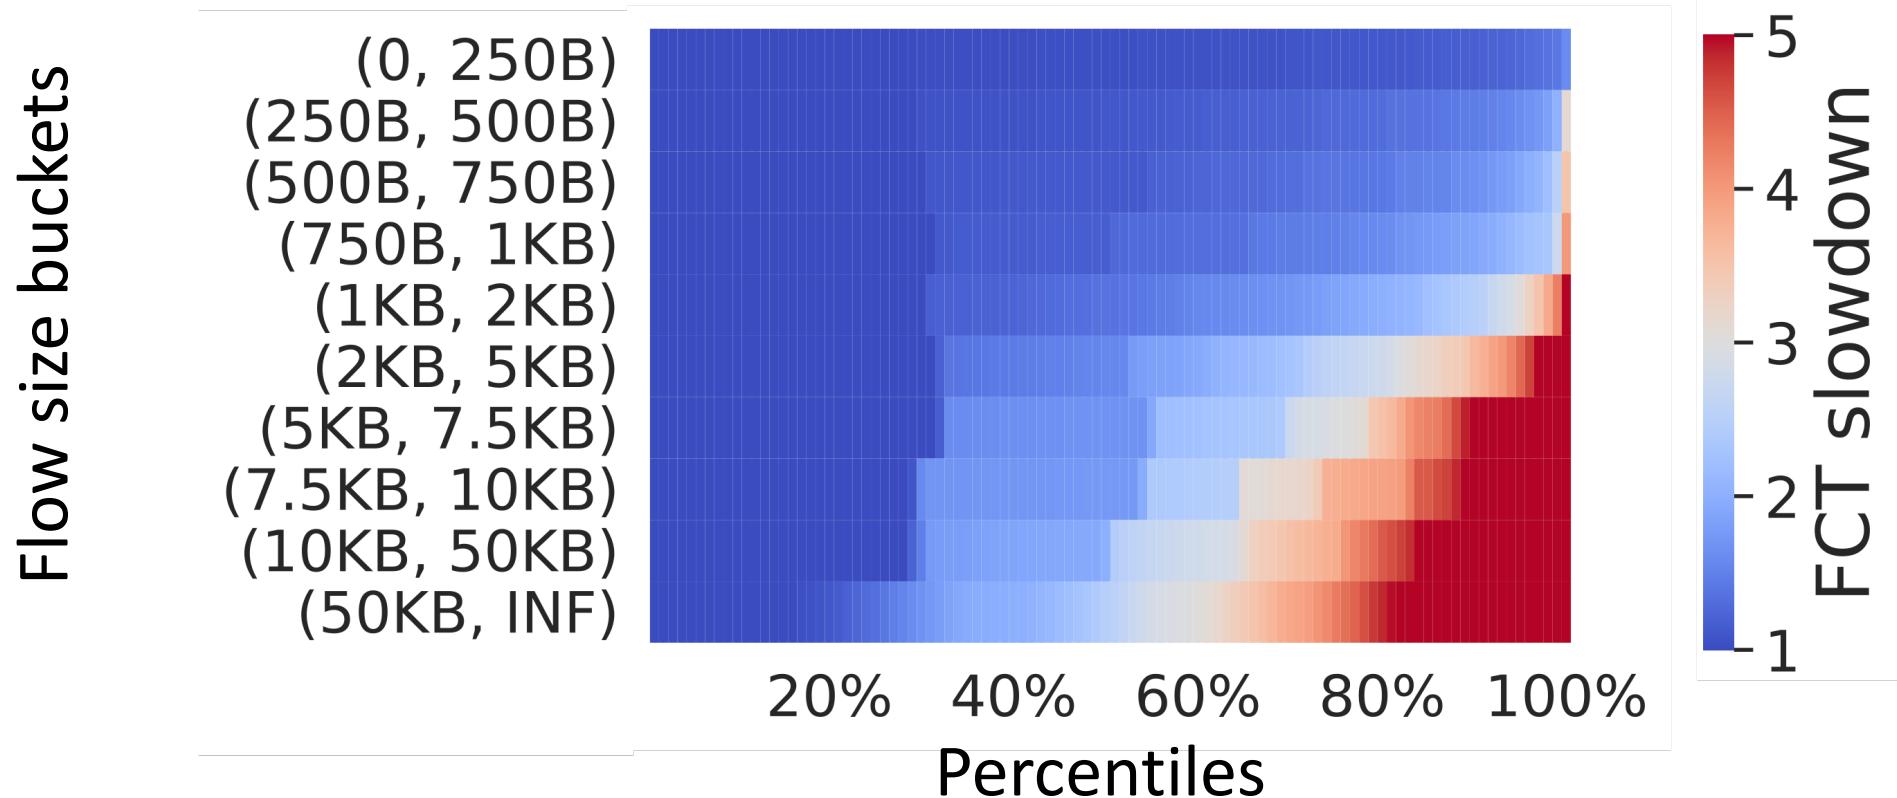
(d) flow size $\in (50\text{KB}, \infty)$



Not accurate, esp. for short flows

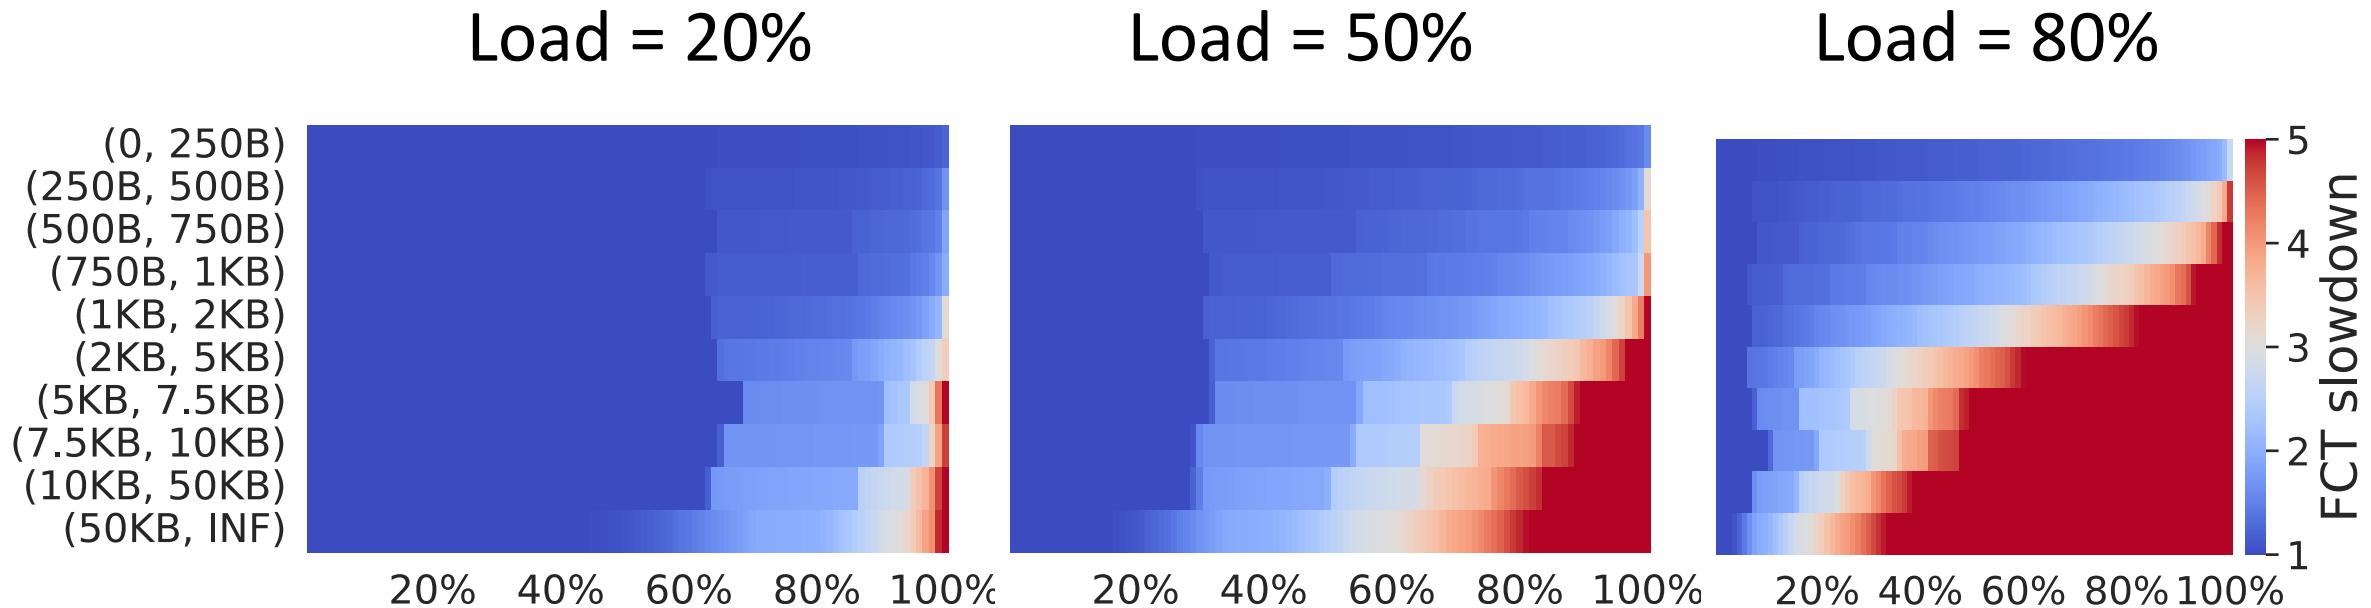
Use flowSim to extract a feature map

- Run path-level simulation on flowSim
- Summarize per-flow FCT slowdown into a feature map



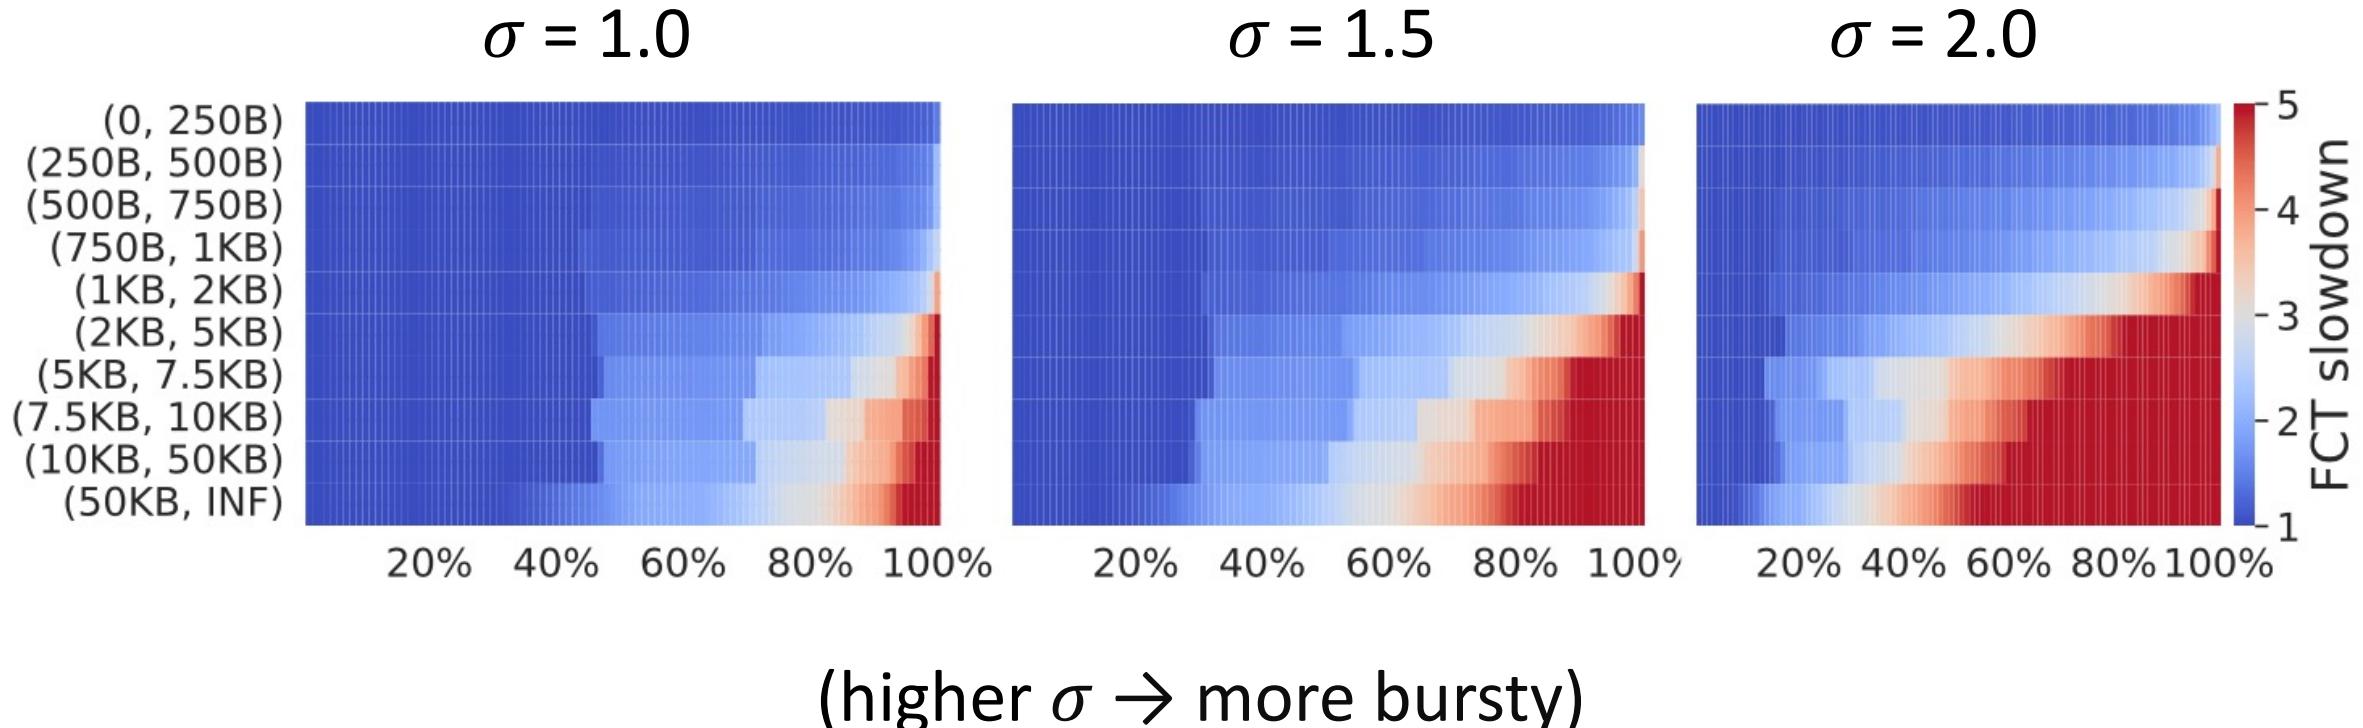
m3's feature map captures workload characteristics

Example: Different levels of load for the simulated network

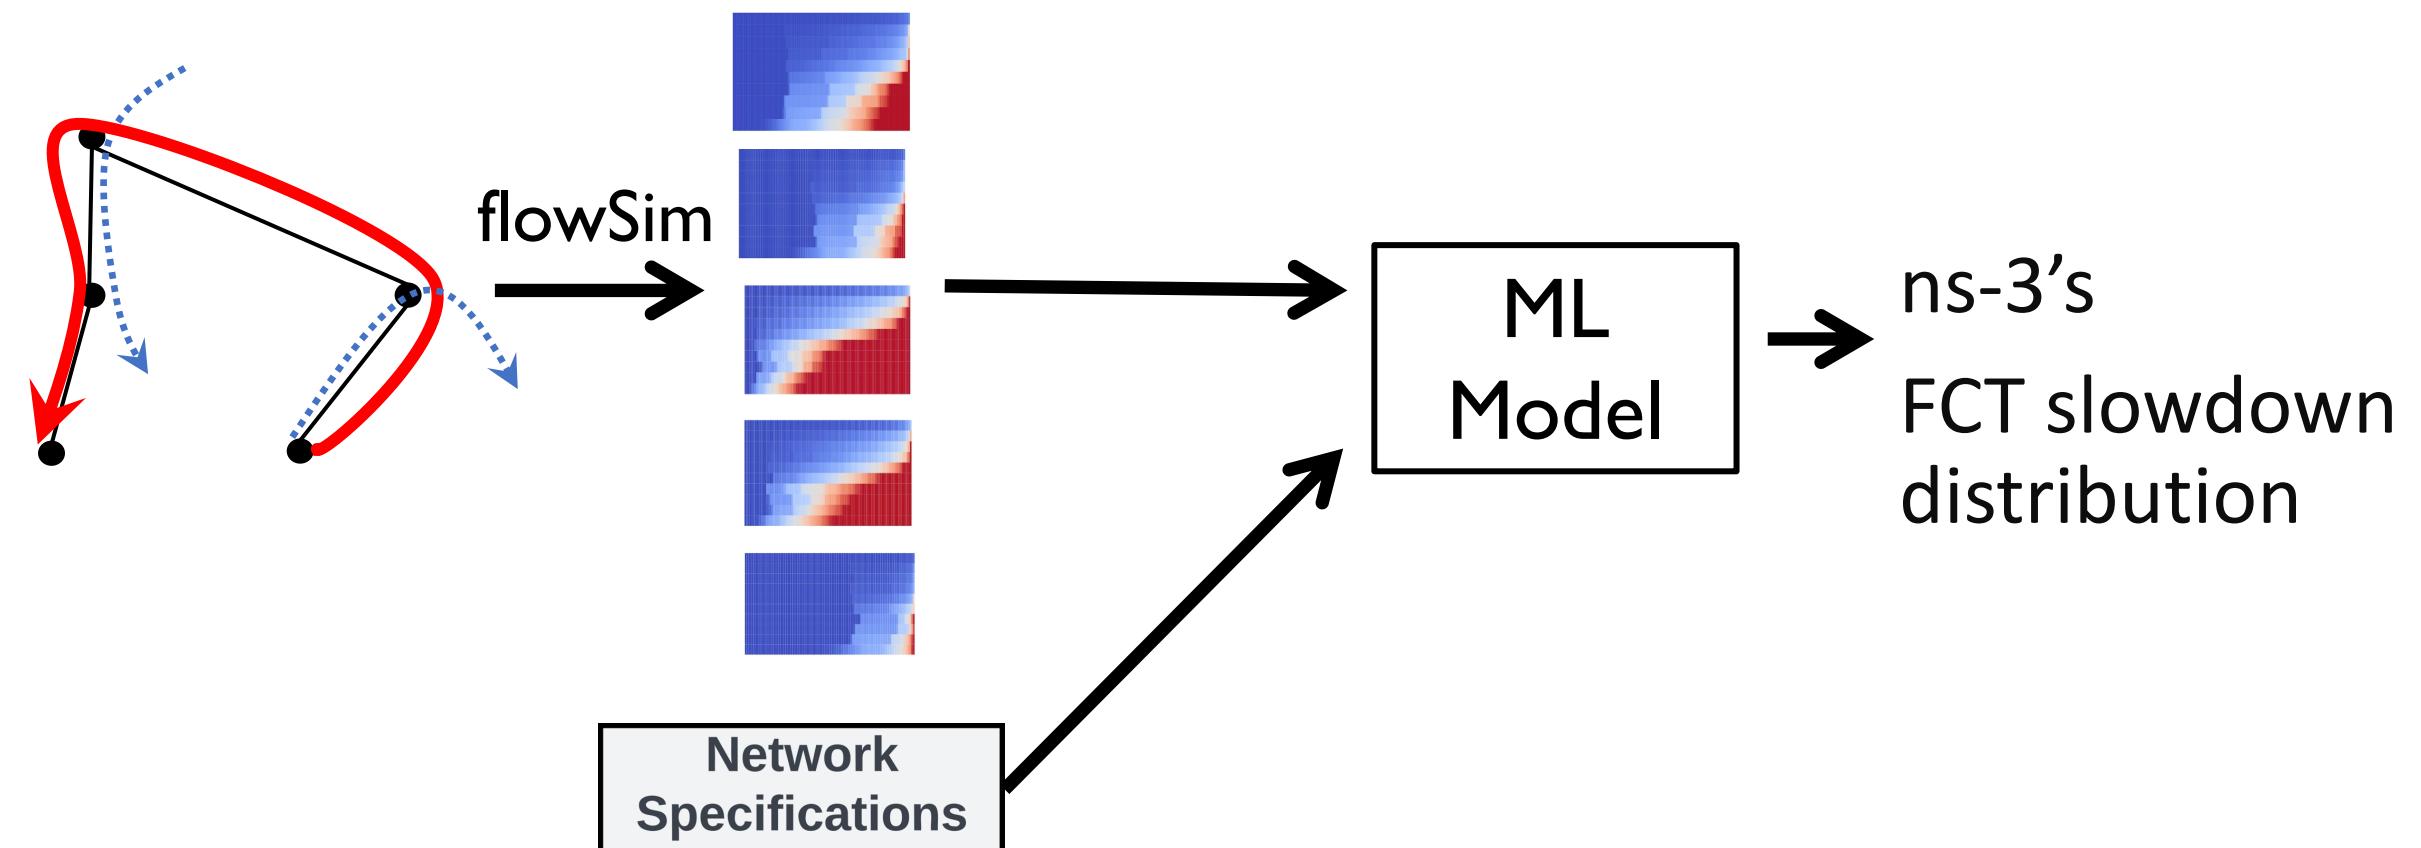


m3's feature map captures workload characteristics

Example: Different levels of burstiness of the workload



m3 uses ML to correct flowSim



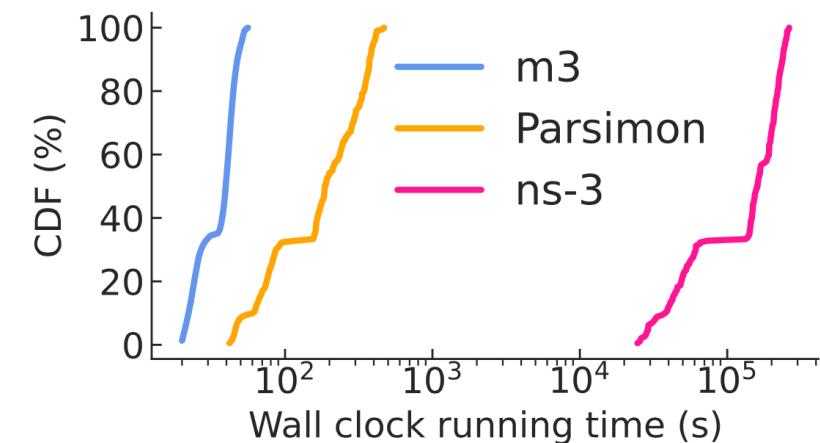
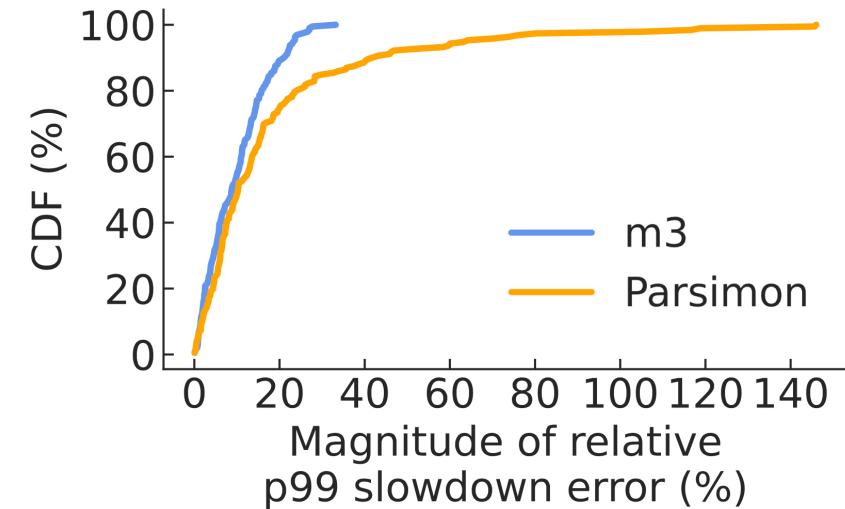
Buffer size, initial window size, different CCAs
as a one-hot vector, and CCA parameters

Training

- Train on diverse synthetic scenarios
 - Path topology: 2-hop, 4-hop, 6-hop parking-lot topology
 - Flow size distribution: Pareto, Lognormal, Exponential, Gaussian
 - Flow inter-arrival time distribution: burstiness from low ($\sigma = 1.0$) to high ($\sigma = 2.0$)
 - Randomly chosen network config: Buffer size, initial window size, CC from DCTCP, TIMELY, DCQCN, HPCC, etc.
- We train the model once and use it directly for inference
 - m3 generalizes to unseen workloads, topologies, and network configurations

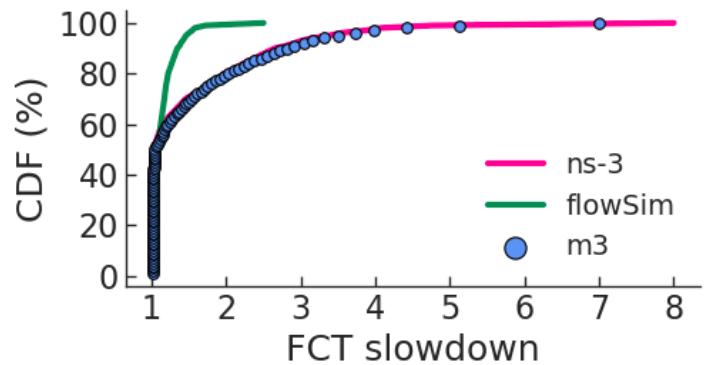
Accuracy and run-time evaluation

- Various production workloads on a sampled Meta's data center fabric
- Baseline: Parsimon^[1]
- m3 reduce Parsimon's mean error from **18.3%** to **9.9%**
- m3 is fast
 - **5.7×** speed-up over Parsimon
 - **3 orders** faster than ns-3

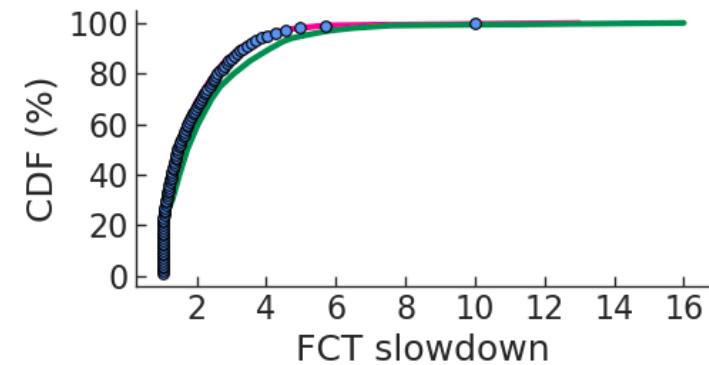


[1] Zhao, Kevin, et al. "Scalable tail latency estimation for data center networks." In proceedings of USENIX Symposium on Networked Systems Design and Implementation (NSDI). 2023.

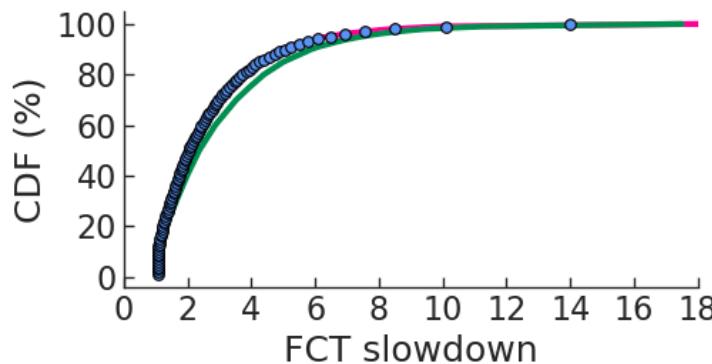
Example Scenario



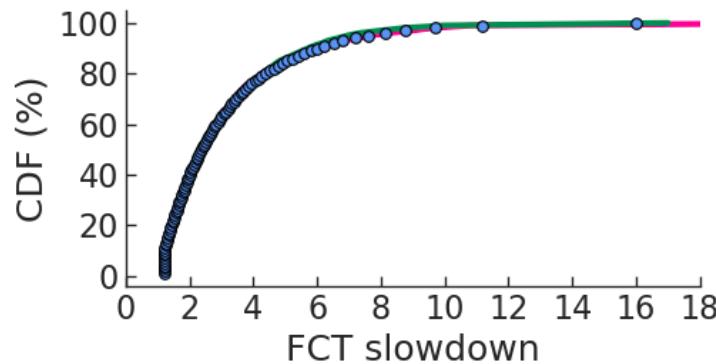
(a) flow size $\in (0, 1\text{KB}]$



(b) flow size $\in (1\text{KB}, 10\text{KB}]$



(c) flow size $\in (10\text{KB}, 50\text{KB}]$



(d) flow size $\in (50\text{KB}, \infty)$

Summary

- ❑ **m3** is fast and accurate for estimating flow-level performance
- ❑ Two key ideas:
 - ❑ Path decomposition
 - ❑ Workload featurization via flow-level simulation
- ❑ **General principle:** Learn to map from simple reference system to real system



<https://github.com/netiken/m3>

Gates

RTL

Packets

Flows

Applications

Systems

Networks

Microscopic,
small-scale

Macroscopic,
large-scale

ML can significantly change how we model and simulate network systems
(esp. at the large-scale macroscopic level)

- Many challenges:
 - Abstractions for learning system dynamics
 - Handling partial observability & building systems for observability
 - Robustness
 - Interpretability

What m3 does and does not do

- ❖ m3 assumes static routes known in advance.
 - ❖ This assumption breaks in the case of dynamic routing strategies like packet-spraying or flowlets
- ❖ m3 predicts the distribution of FCT slowdown for various size buckets
 - ❖ Not per-flow or per-packet performance
- ❖ m3 learns the impact of network parameters from training examples
 - ❖ cannot predict unseen protocols

Large-scale simulation

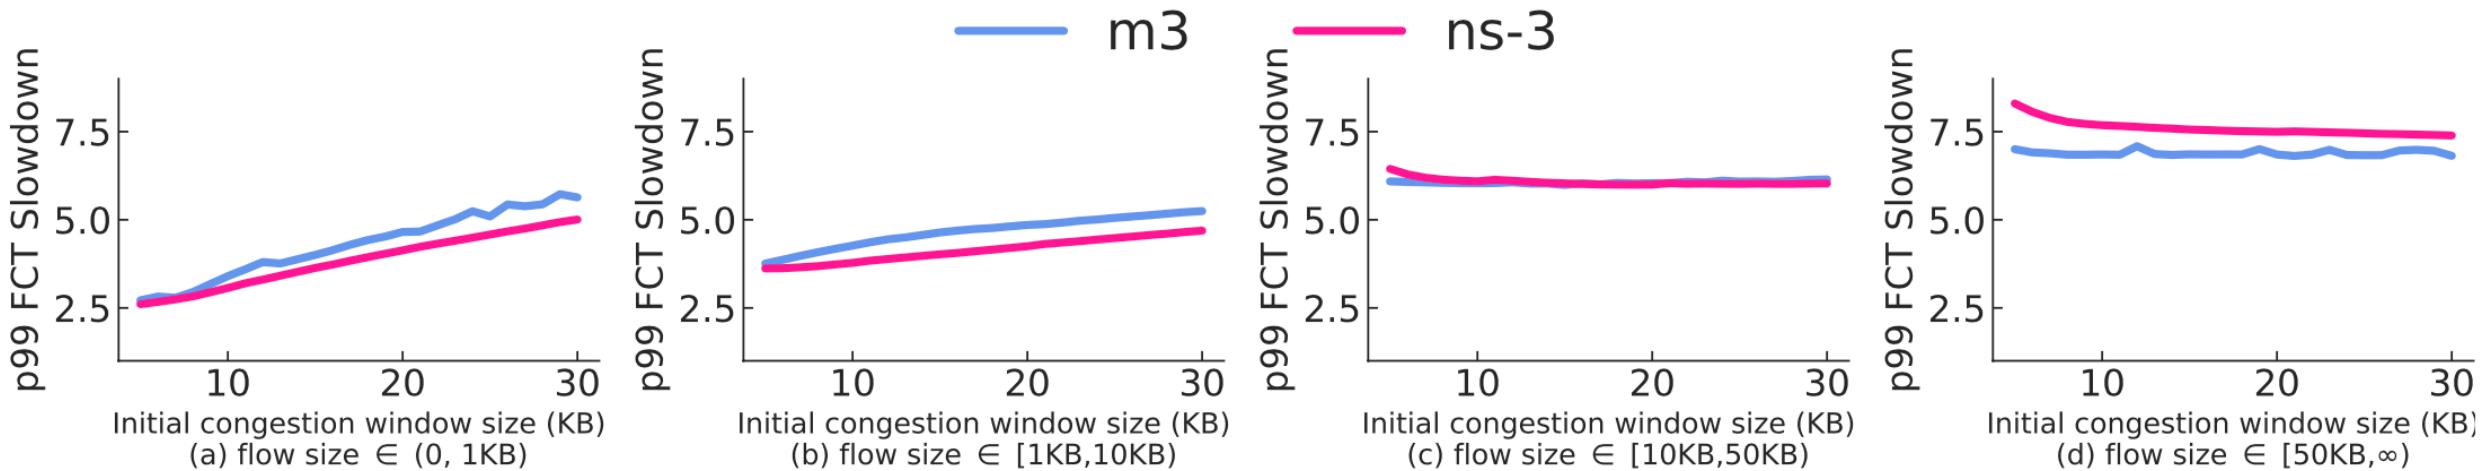
- Meta Production Workload
- A 384-racks, 6,144-host Meta's data center fabric
- Baseline: Parsimon

Init. Window	Methods	p99	Error	Time	Speedup
10KB	ns-3	2.05	-	13.5h	-
	Parsimon	4.29	+109%	1m29s	546×
	m3	2.10	+2.44%	37s	1314×
18KB	ns-3	2.44	-	11.9h	-
	Parsimon	2.73	+11.9%	1m24s	510×
	m3	2.30	-5.74%	40s	1071×

➤ m3 is faster, more accurate, and robust

m3 enables rapid design-space exploration

Example: Effect of initial window size on p99 FCT slowdown for different classes of flows



The mean runtimes of each configuration in the plot

- m3: 25.2 seconds
- Ns-3: 8 hours

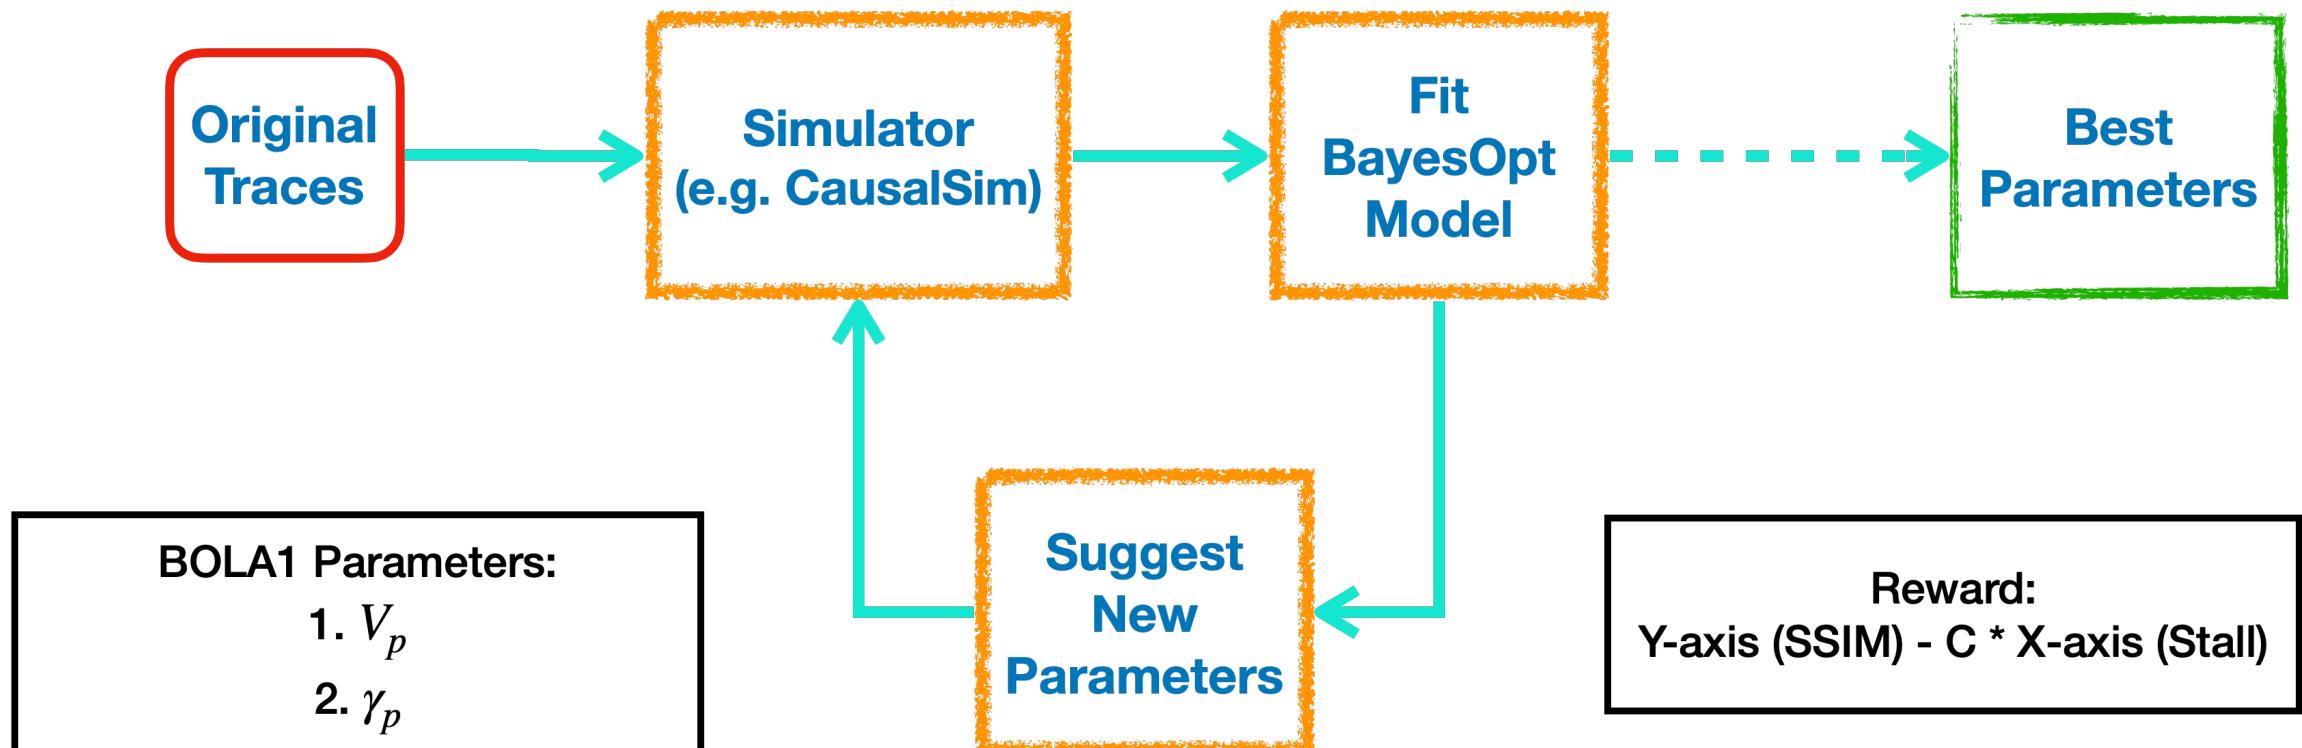
Vision



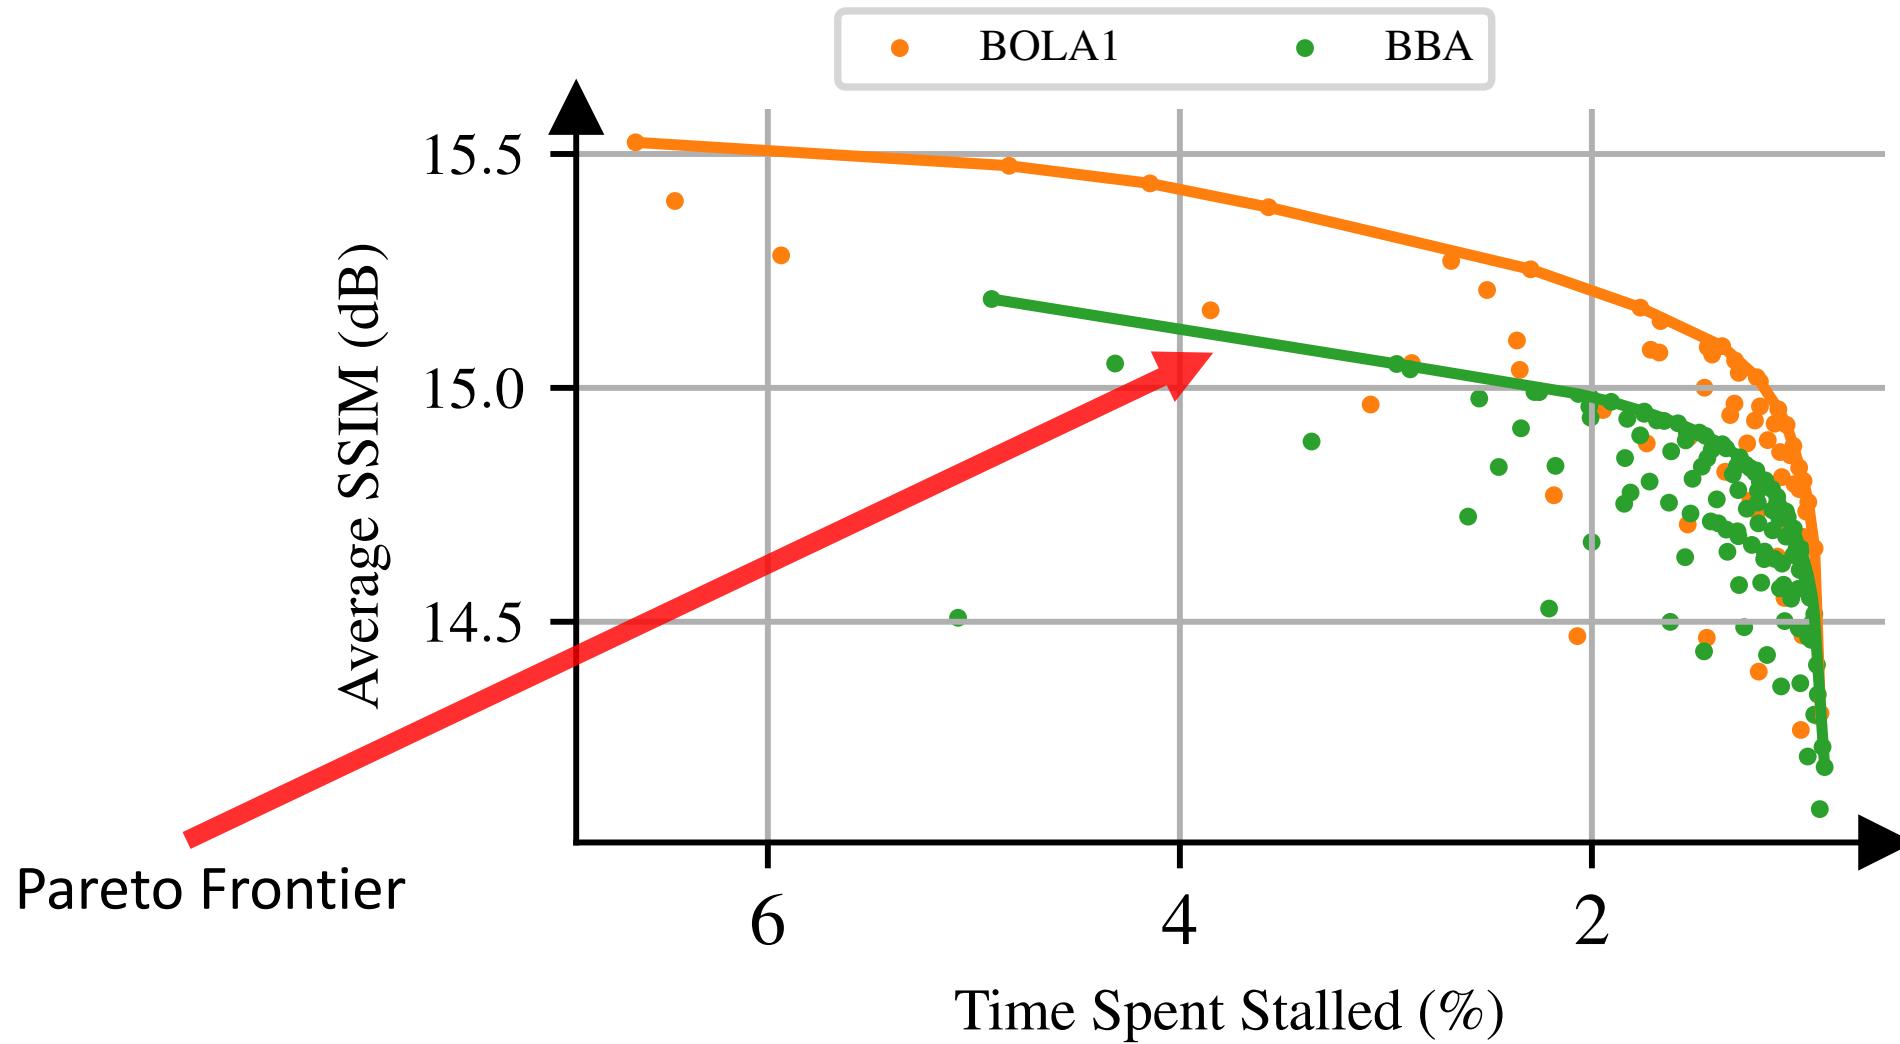
- Networks produce massive amounts of data
- Relatively easy to add instrumentation, run active experiments, ...
- Goal: Learn high-level network dynamics from observational data

Case Study: CausalSim in the Wild

We evaluated 300+ different variants in less than 1 day

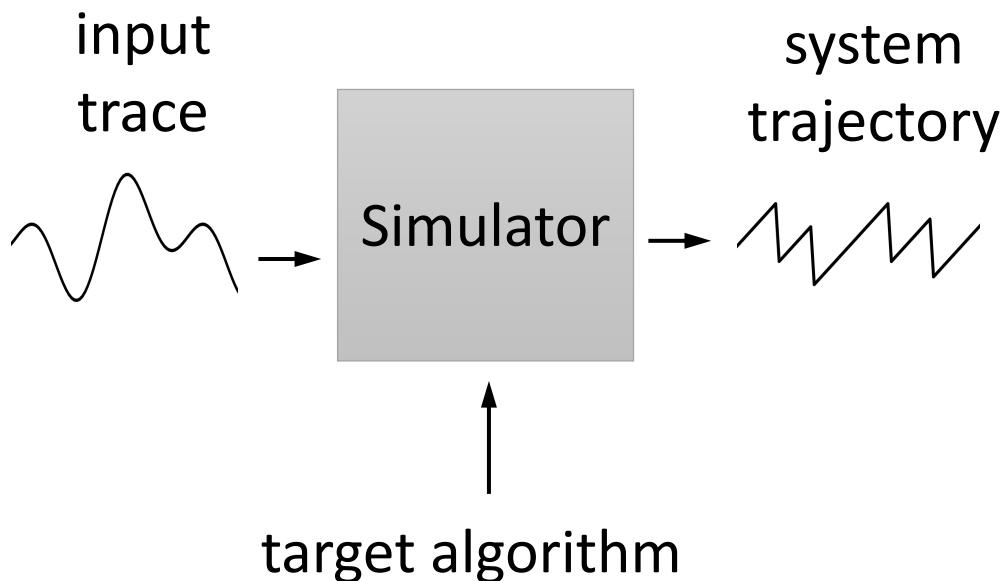


Case Study: CausalSim in the Wild

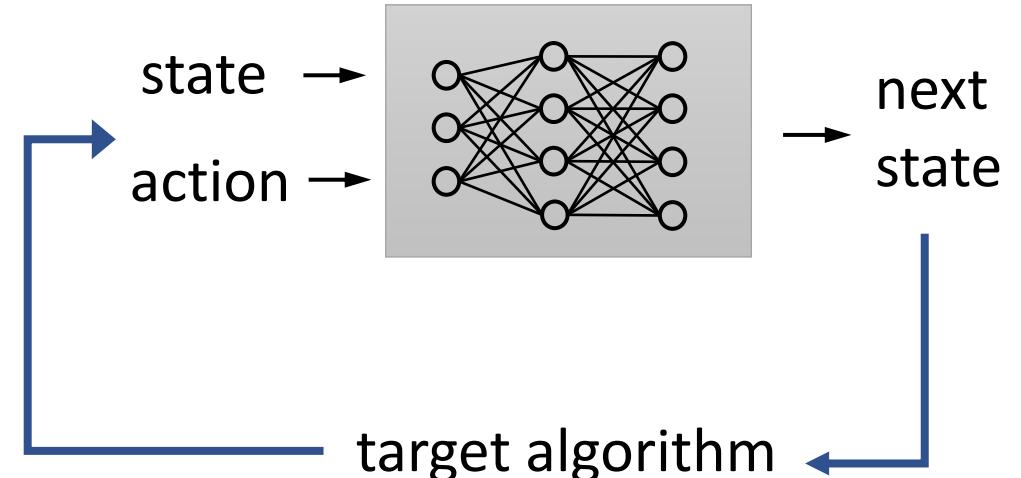


Standard Approaches

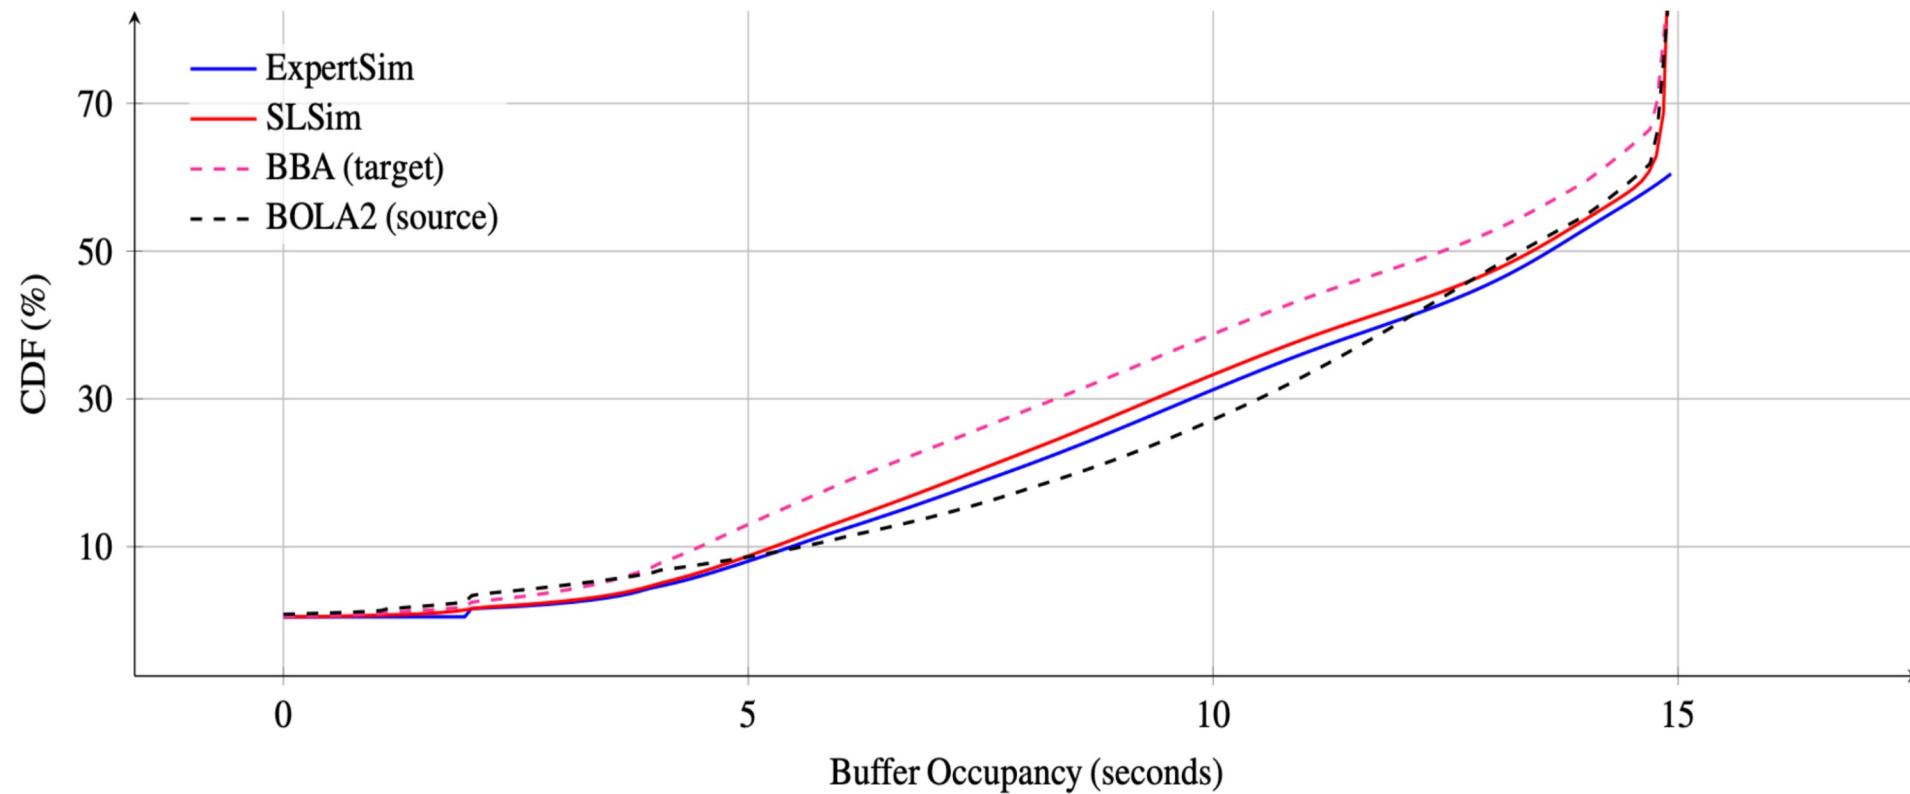
Expert-designed Simulators



Supervised Learning



How accurate is trace-driven simulation?



CausalSim: Unbiased Trace-Driven Simulation

