

From Prediction to Proof: Rethinking AI for Systems and Networks

Matthew Caesar
University of Illinois at Urbana-Champaign
caesar@illinois.edu

What is Artificial Intelligence?

artificial intelligence

/ ,ɑ:tɪfiʃl ɪn'telɪdʒ(ə)ns/ noun

- 1: the capability of computer systems or algorithms to imitate intelligent human behavior
- 2: a branch of computer science dealing with the simulation of intelligent human behavior by computers

“Human thinking” has achieved success



AI Art



AI Music



AI Education

perplexity

Ask anything...

Search Research



AI Search



AI Management



AI Games

“Human thinking” is prone to problems




Semantic incoherence

Hallucination
Mode collapse
Semantic drift
Interface misalignment
Context drift




Model Opacity

Non-transparency
Causal confusion
Overconfidence
Catastrophic forgetting
Prompt sensitivity



Security

Reward hacking
Specification gaming
Role breakage
Deployment drift
Data poisoning




Data Errors

Bias
Model inversion
OOD Errors
Spurious Correlations
Over/underfitting

Our field: Systems and Networking



Financial and
trading networks



Medical devices



Clouds and
online services




Societal
infrastructures

- We work on infrastructures behind modern society
- These infrastructures are complex, huge, dynamic
- Yet critical to get right

Humans aren't good at getting things right

- Human error is #1 cause of problems in systems/networks
 - 80% of data center outages, 95% of data breaches, 95% of data loss incidents as a direct and immediate cause
 - Some categories of faults are 100% human-caused (misconfigurations, vulnerabilities, social engineering attacks, software bugs, insider attacks, ...)
 - Consistently listed by CISOs as top risk; #1 category on IBM Threat Index
- Daily news is filled with vulnerabilities, misconfigurations, errors
 - Human error is the largest (increasing) contributor to failure
- “Human thinking” isn’t such a great approach to designing and operating systems



Should AI be the goal of the systems community?

- Should we be striving to build systems that “think like humans”?
 - AI may be useful, but it might also not be exactly what we want
- We may want to think more deeply about what we actually want from AI
 - We may be able to come up with something better

An alternative proposal


Artificial Reasoning

-- The capability of computer systems to derive optimal, fully-correct, understandable and analyzable solutions, through a formal sequence of logical steps


Focus on things we care about: safety, resilience, correctness, efficiency, explainability

A complement, not a replacement for AI techniques

What would an AR for Systems look like?



What would an AR for Systems look like?




Towards a solution

- Seems difficult, but recent breakthroughs make progress towards this vision
 - *Formal methods* enable perfect and precise modeling, proving, and synthesis of diverse platforms
 - *Virtualization technologies* allow manipulation of time and inputs, log capture, and deterministic analysis and execution of real software
 - *Optimization techniques* enable exact and rigorous derivation of optimal behaviors in complex environments

How can we collectively leverage these components to build generalized Automated Reasoning for networked systems?

System Architecture of an AR



Key Challenges and Solution Approaches

1. How can we formally reason about dynamic, real systems where vagaries of execution behavior really matter?

→ *Solution approach:* [System-Guided Formal Modeling](#)
("guide" formal modeling with running of real implementation code)


2. How can we safely integrate less-trusted AI inputs into the reasoning process?

→ *Solution approach:* [Cognitive Input Autocorrection](#)
(automatically repair inputs from AIs to match correctness specification)


How to create a model for reasoning?

One option: run the system's software in an emulated environment (e.g., VMs)

Limitations: replication overheads, execution overheads, limited coverage



Software-Based Model




State Space Exploration

A more scalable approach to modeling

Idea: create a model with **formal methods**

- Can leverage rigorous techniques to efficiently, exhaustively search
- Lower bandwidth/compute requirements to traverse model states





Benefits of formal modeling

- Unveil corner cases that are hard to argue about without models
- Better predictability compared to pure emulation
 - Automatic state transition without synchronization from the deployment
 - The model-based representation is less likely to deviate from the deployment state
- More optimization opportunities compared to working with real hw/sw
 - E.g., partial-order reduction (Only the orders of packets entering the stateful components are relevant) for simulating multiple connections through a stateful network


Challenges in applying formal modeling for reasoning

- Models must be truthful and precise
 - Real-world systems face complexities, scale, non-determinism, environmental interactions (e.g., time/event triggers), distributed protocols, energy/bw constraints
- Models with unnecessary details have lower scalability




Formal Emulation: Combine models with emulation

- For each emulated component, we communicate with them from formal models by injecting and interpreting concretized packets and events
- State space is trimmed by synchronizing with deployed system





Formal Emulation: Combine models with emulation

- We employ virtual interfaces and network namespaces for lightweight emulation
- The packets to/from emulation instances are interpreted for model state transition




Formal Emulation: Tracking emulation states

- Emulation state := initial state + history of events
 - a. Events: packet arrivals, sensor updates, etc.
- Emulation instances need to be in the right state before injecting packets
 - a. Reset to the initial state
 - b. Replay the history of events
- Hashing histories of events to reduce memory overhead





Formal Emulation: Dynamic multi-connection coordination

- How to model the non-deterministic nature of multiple connections in a networked component?
 - Our model implemented the non-deterministic choices for the model checker explore.
 - Apply partial-order reduction (POR) to reduce unnecessary search space.
 - POR heuristic:
 - Pick an arbitrary connection until every connection is about to enter an emulation.
 - Explore all orderings of the connections entering the emulations. (And repeat.)



Formal Emulation: Dynamic multi-connection coordination

- How to handle new connections initiated by emulations?
 - Parse received packets, and add new connections to the model state.
- What about L3 vs L7 proxies?
 - How to tell if a packet has gone through a L3 proxy or belongs to a new connection?
 - We treat all proxied packets as new connections. “Pause” connections when necessary.





Formal Emulation: Association interpretation

- Challenge
 - Once we inject a packet, how do we know if the packet is dropped or not?
- Method 1 (drop timeout estimation)
 - Adjust the drop timeout based on injection RTT estimate, similar to TCP retransmit timeout.
 - Cumulative estimate of time := $\mu_{latency} + \sigma_{latency} \times \max\left(4, \text{ceil}\left(\frac{2 \times (\text{Number of jobs})^{1.5}}{\text{Number of total cores}}\right)\right)$
- Method 2 (Linux per-packet drop_monitor)
 - Since 5.4, we can request for per-packet drop alerts from kernel.

	Drop timeout	Kernel drop_monitor
Advantages	<ul style="list-style-type: none">• Available for all types of emulations• Easy to implement	<ul style="list-style-type: none">• No false violation when there is no tail drops
Limitations	<ul style="list-style-type: none">• Potential false violations under high load• Longer wait time for dropped packets	<ul style="list-style-type: none">• Only appropriate to hypervisors/kernels supporting this method

Performance Results




- CPU time grows linearly with network size, memory usage almost constant
- Timeout method generally faster than drop-monitor; drop monitor has additional overhead from registering and checking drops in kernel

Can we still use AI?

- AI has proven very useful to systems
 - But it can also be wrong
 - But we don't want to just not use AI
- Idea: can we automatically correct AI inputs?
 - React quickly without need for operator in the loop
 - Transparent integration with existing workflows/APIs
 - Operator can view fixes to get insights on understanding their errors

Cognitive Input Autocorrection



- Autocorrection layer synthesizes repairs to inputs in real time
- Fixes/patches derived from formal methods, guaranteeing compliance with provided specifications
- Optimization-based framework to place observation and correction programs

Two ideas for future work in Automated Reasoning

- Formally Verified AI Systems
 - There have been great strides in formal verification of AI techniques, and networked systems separately
 - Can we apply these techniques to build AI-based systems with formal guarantees on correctness, QoS, etc.?
- Using Reason to Design Systems
 - Much of research is getting automated
 - E.g., we rely on AI more and more for algorithm design
 - Has been harder to do for certain things, like architecture, which require robustness
 - Formal logics are good at deriving rigorous designs in other disciplines
 - Can we automatically derive system architectures with AR?

Conclusions

- The future of AI in systems isn't about mimicking minds, it's about mastering the goals important to our community
 - Artificial Reasoning may open the door to a new kind of AI
 - Centered around critical properties such as correctness, trust, and understanding
- We presented some abstractions that can help bring this vision closer to reality
 - Leveraged recent advances in formal methods, modeling, virtualization, and optimization to achieve scale, completeness, rigor
 - Early results demonstrate benefits and practicality of approaches