



中南大學  
CENTRAL SOUTH UNIVERSITY



湖南大學  
HUNAN UNIVERSITY

中南大學  
CENTRAL SOUTH UNIVERSITY

# LEFT: LightwEight and FasT packet Reordering for RDMA

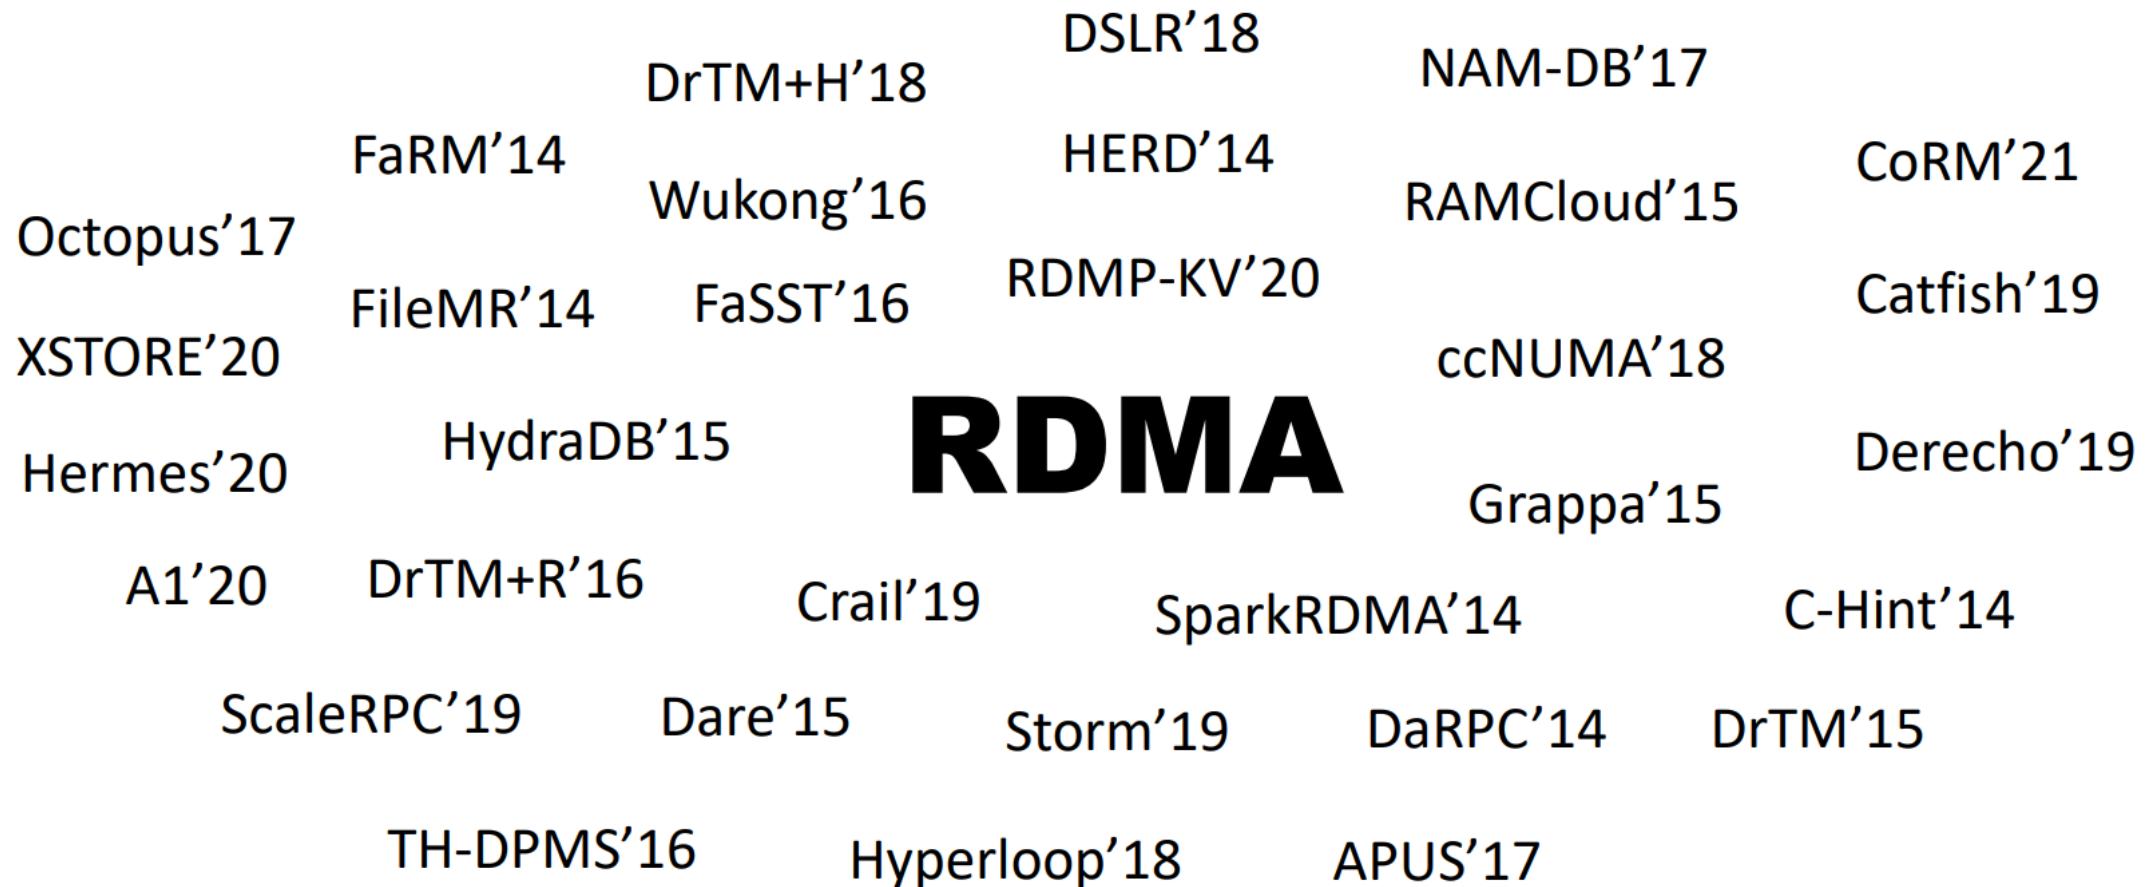
Peihao Huang, Xin Zhang, **Zhigang Chen\***, Can Liu, **Guo Chen\***

[204701003@csu.edu.cn](mailto:204701003@csu.edu.cn)

August 3, 2024

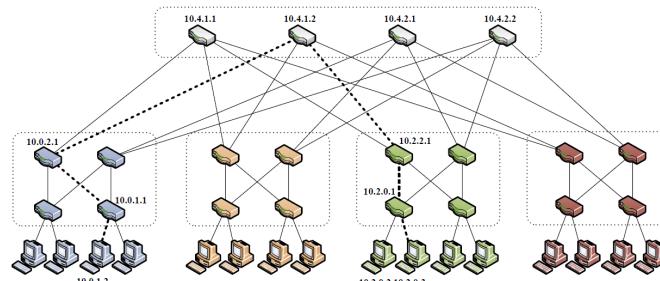
# RDMA is hot topic in HPC and DCN

- Lower latency, Higher bandwidth, Lower CPU utilization



# RDMA needs to deal with OoO packets

- Abundant parallel paths in DCN

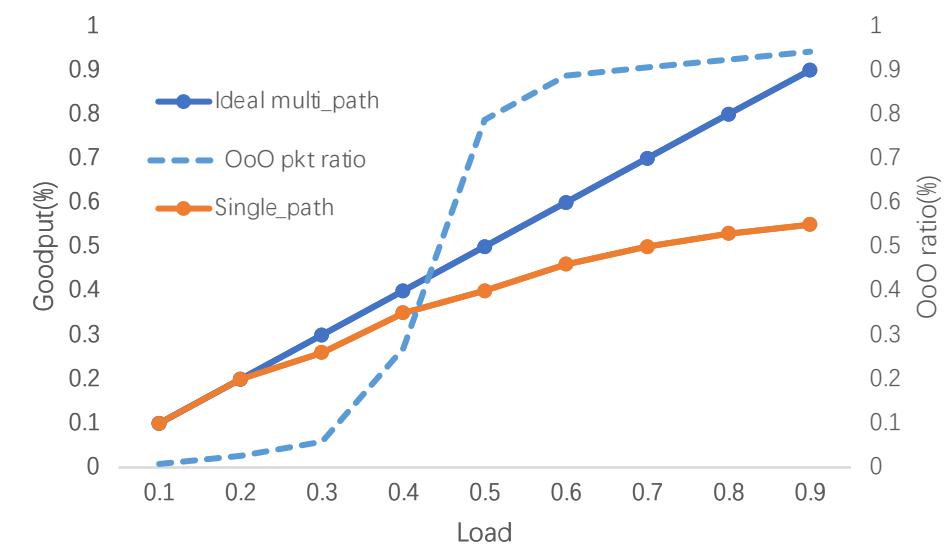
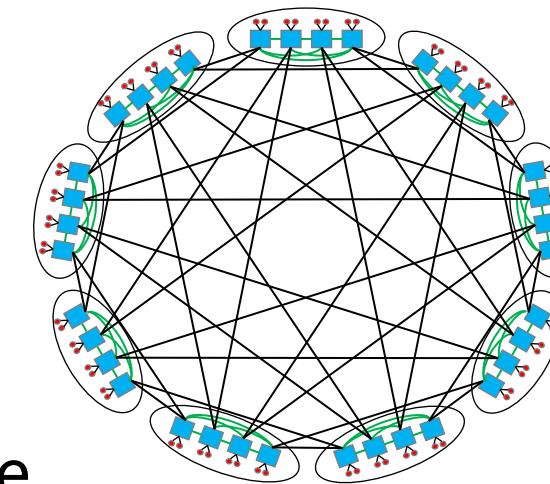


- Single-path transmission cannot fully utilize network bandwidth

- Uneven load
- Hash conflict

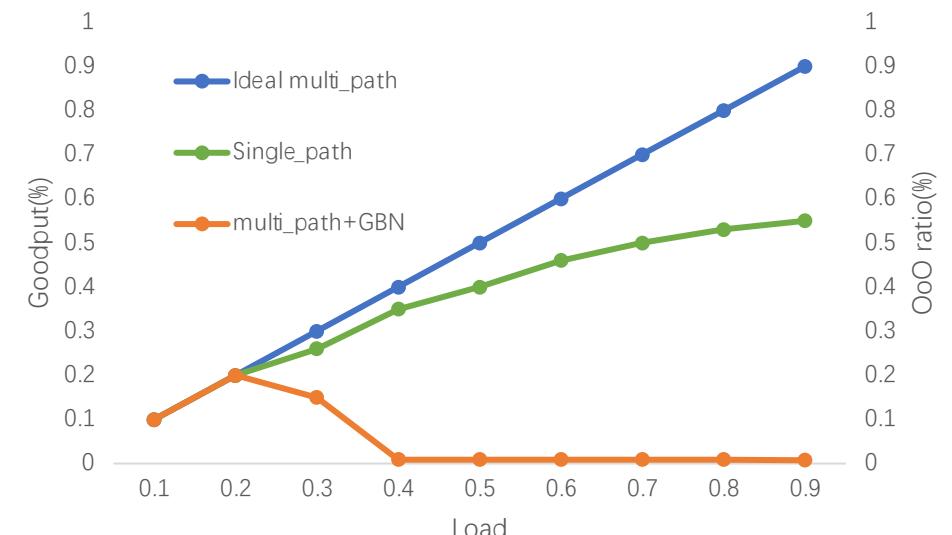
- Multipath brings massive OoO packets

- Path delays, switch queue lengths vary
- OoO packets ratio related to the load



# RDMA needs to deal with OoO packets

- Current RDMA behaves very poor in deal with OoO  
**Go-Back-N (GBN) loss recovery has very low efficiency**
  - Mis-treat OoO as loss, abundant unnecessary retransmission



# RDMA needs to deal with OoO packets

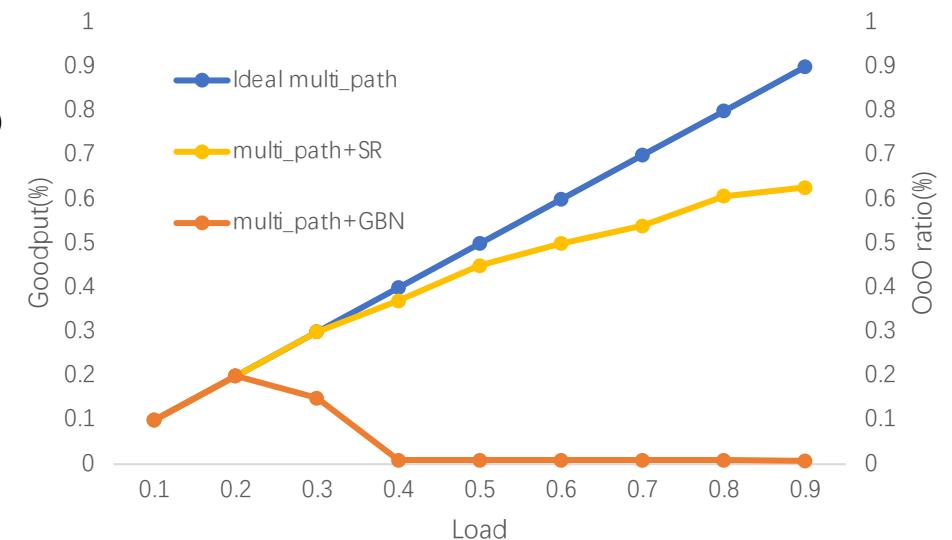
- Current RDMA behaves very poor in deal with OoO

## **Go-Back-N (GBN) loss recovery has very low efficiency**

- Mis-treat OoO as loss, abundant unnecessary retransmission

## **Selective retransmission (SR) is still inefficient to deal with OoO**

- Still Mis-treat OoO as loss
    - Sending rate decreased due to CC miss-acting to OoO, and unnecessary retransmission
  - Bitmap record OoO packets, which consumes too much RNIC memory
    - 5BDP bitmap required per connection
    - BDP:  $200\text{G} * 32\text{us}$  (1KB MTU)  $\sim 100\text{B}$
    - 2.5MB required for 5k connections



# Related work

- To reduce OoO packets, limit load-balancing granularity and flexibility thus less pressure on RNIC reordering
  - Letflow<sup>[1]</sup>, Conweave<sup>[2]</sup>, MPRDMA<sup>[3]</sup>

**Inferior network utilization due to sub-optimal load-balancing**

- Improve RNIC SR efficiency
  - IRN<sup>[4]</sup>, SRNIC<sup>[5]</sup>, MELO<sup>[6]</sup>

**Mismatch to the OoO scenario, and burden on RNIC**

- SR consumes much memory
- SR is slow in processing with massive OoO packets

**Fast and memory-efficient reorder scheme on RNIC to support fine-grained per-packet multi-path load balancing!**

[1] Let It Flow: Resilient Asymmetric Load Balancing with Flowlet Switching, NSDI'2017

[2] Network Load Balancing with In-network Reordering Support for RDMA, SIGCOMM'2023

[3] Multi-Path Transport for RDMA in Datacenters, NSDI'2018

[4] Revisiting Network Support for RDMA. SIGCOMM'2018

[5] SRNIC: A Scalable Architecture for RDMA NICs. USENIX2023

[6] Memory Efficient Loss Recovery for Hardware-based Transport in Datacenter. APNET'2017

# Challenges

- Key point
  - **Recording OoO, bitmaps must be lightweight and fast**
  - CC should differentiate normal OoO from packet loss
- Challenge 1: Massive Bitmaps in Small On-chip Memory
  - Bitmap Pooling for massive storage efficiency
  - Bitmap Caching for frequent random access efficiency
- Challenge 2: High Disorder Processing within Small Time Constraints
  - Packet gather to reduce average bitmap latency
  - Dynamic scheduling with fast and slow paths

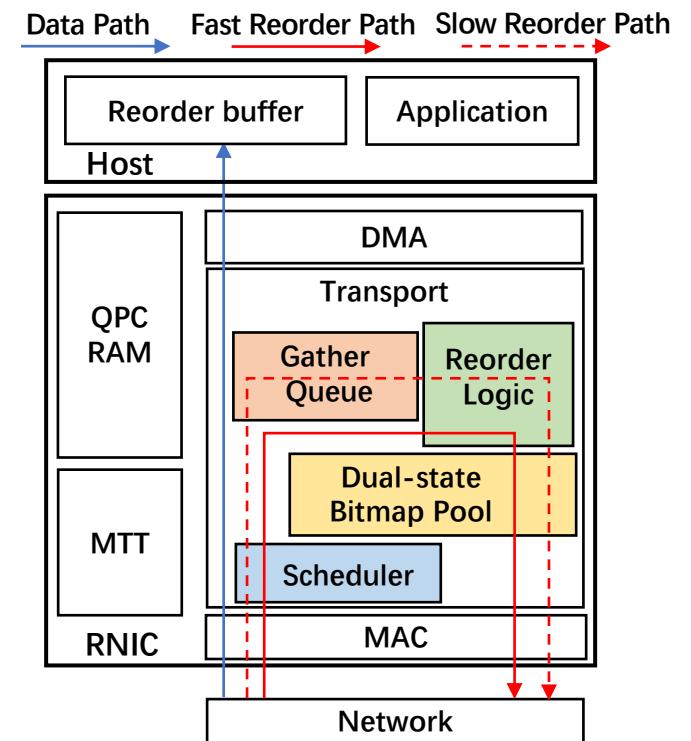
# Overview of LEFT

## ■ Data Path

- Reorder buffer is allocated to each connection for OoO data
- Submit to the application after OoO recovery

## ■ Control Path

- Reorder: OoO packets no longer trigger retransmission and CC
- Dual-state bitmap pool: Small, fast bitmap shared by all connections
- Gather queue: Aggregating packets from different connections
- Scheduler: Selects the path for incoming packets based on the degree of packet reordering



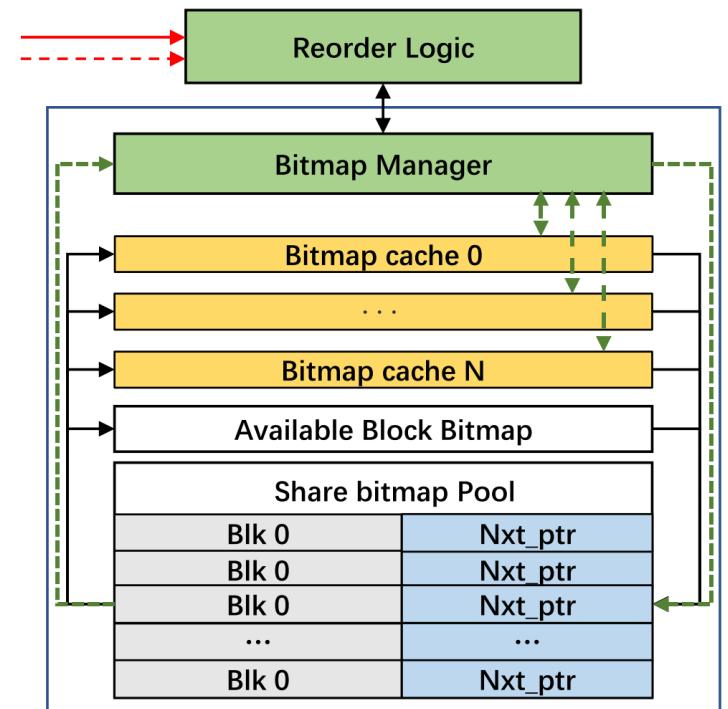
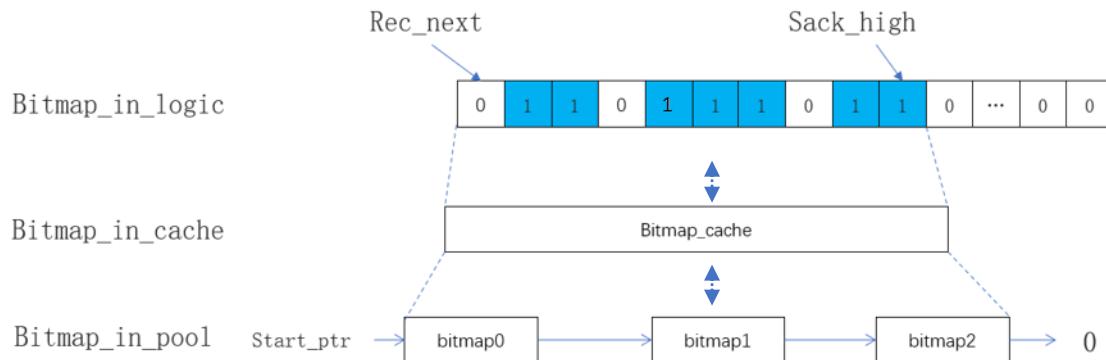
# Dual-state bitmap pool

## ■ Structure

- Available block bitmap marks the available bitmap blocks in pool
- Bitmap block as the smallest bitmap sharing unit
- Bitmap blocks of the same connection are linked by Nxt\_ptr
- Bitmap cache consists of loop registers

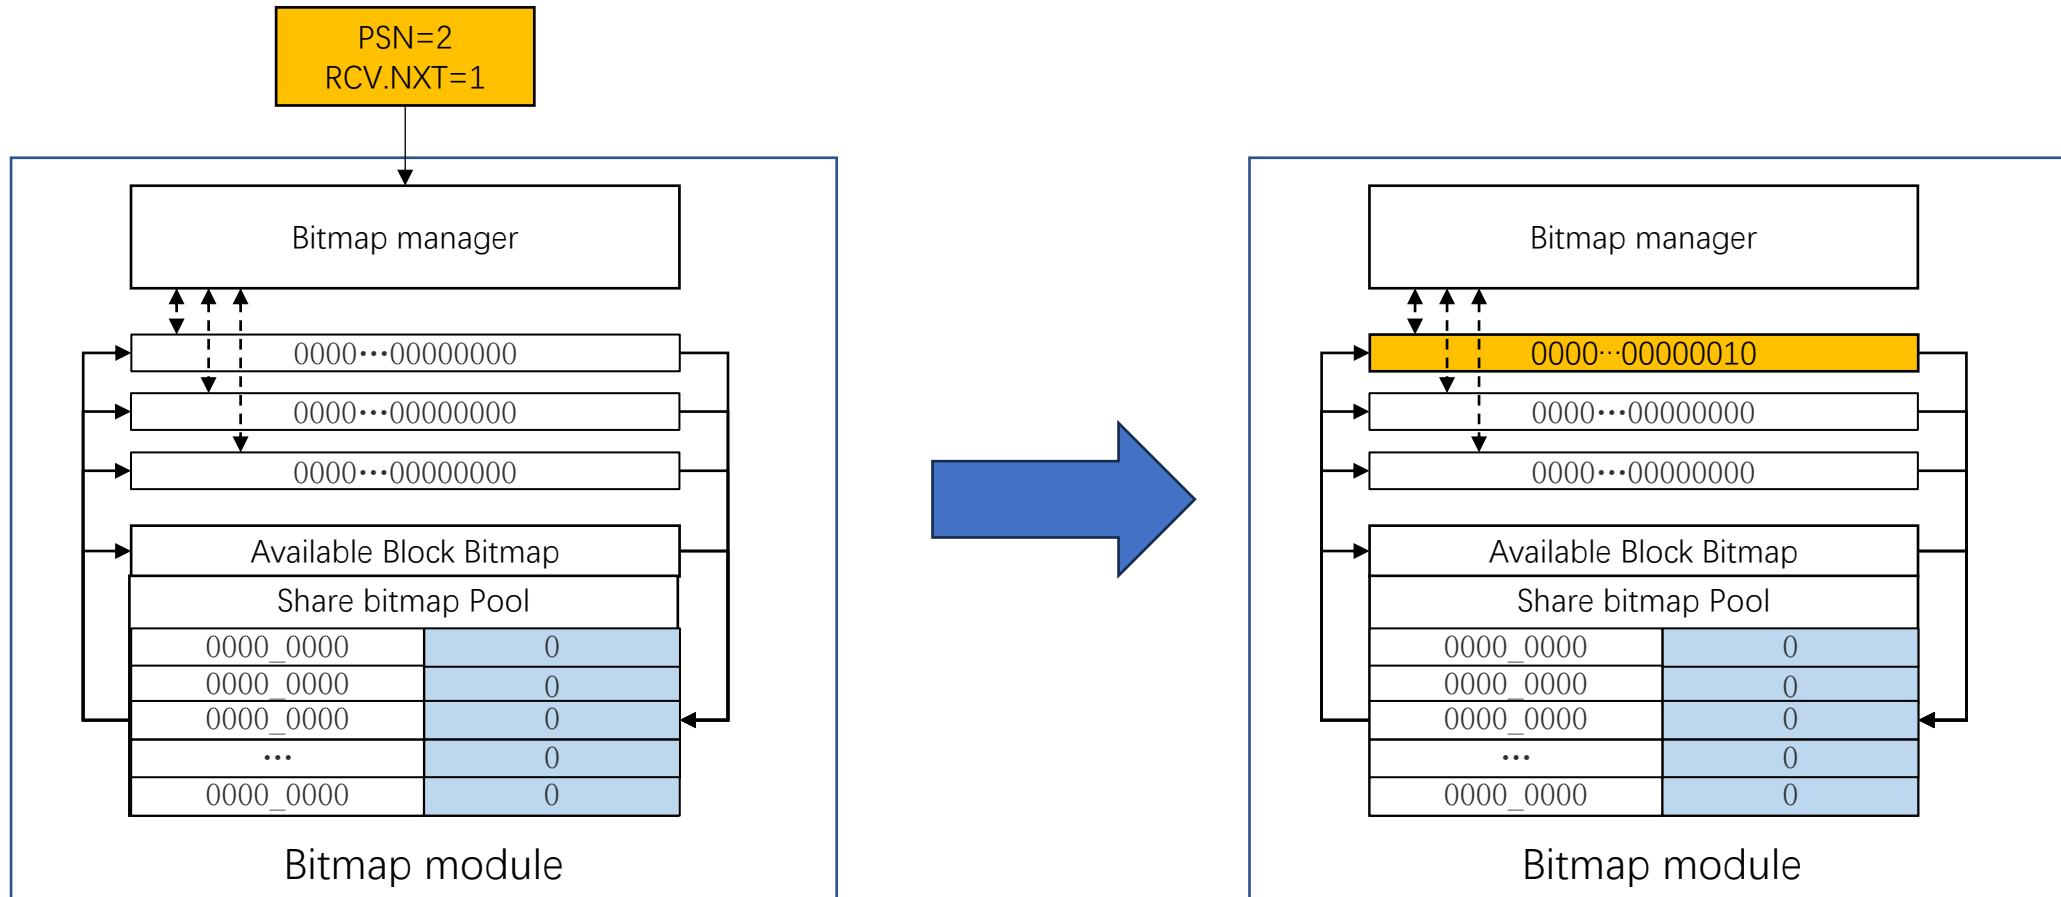
## ■ Connection bitmap will be stored in the cache or pool

- Simple bitmap in cache, fast but memory inefficient
- Link list bitmap in pool, slow but memory efficient [6]



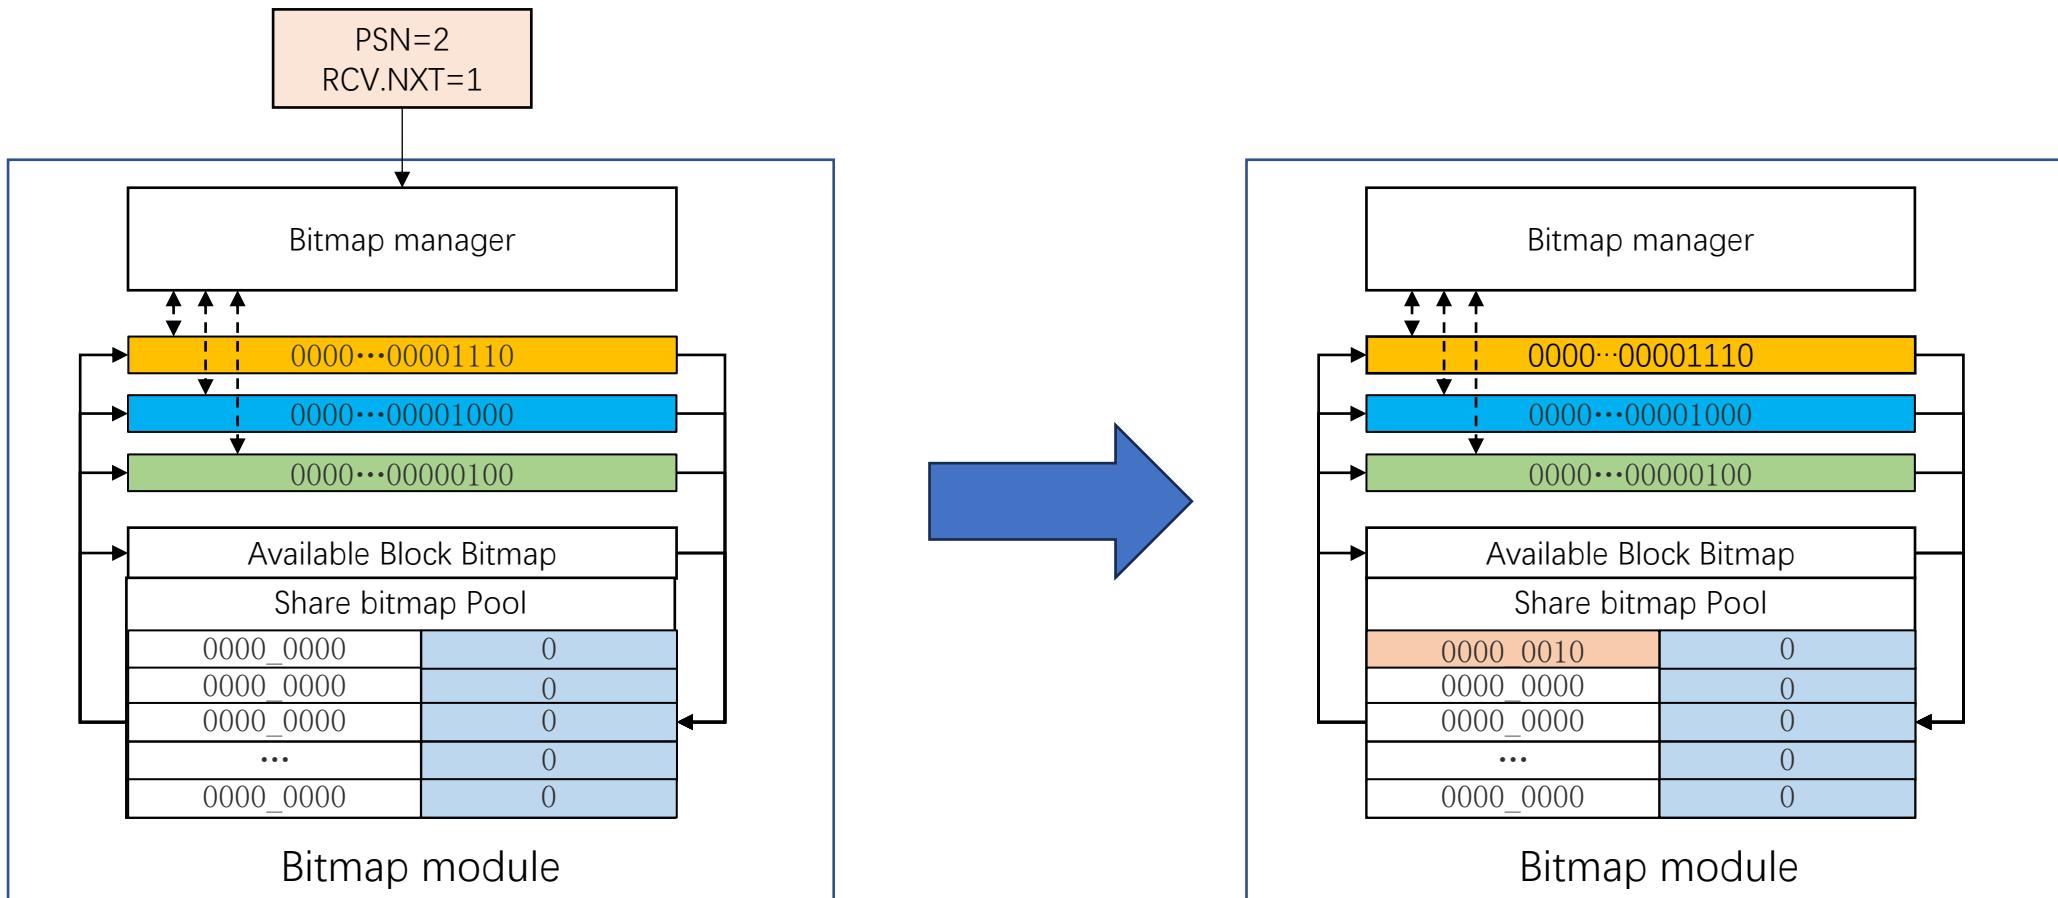
# Dual-state bitmap pool

- Cache hits, bitmap updated in cache



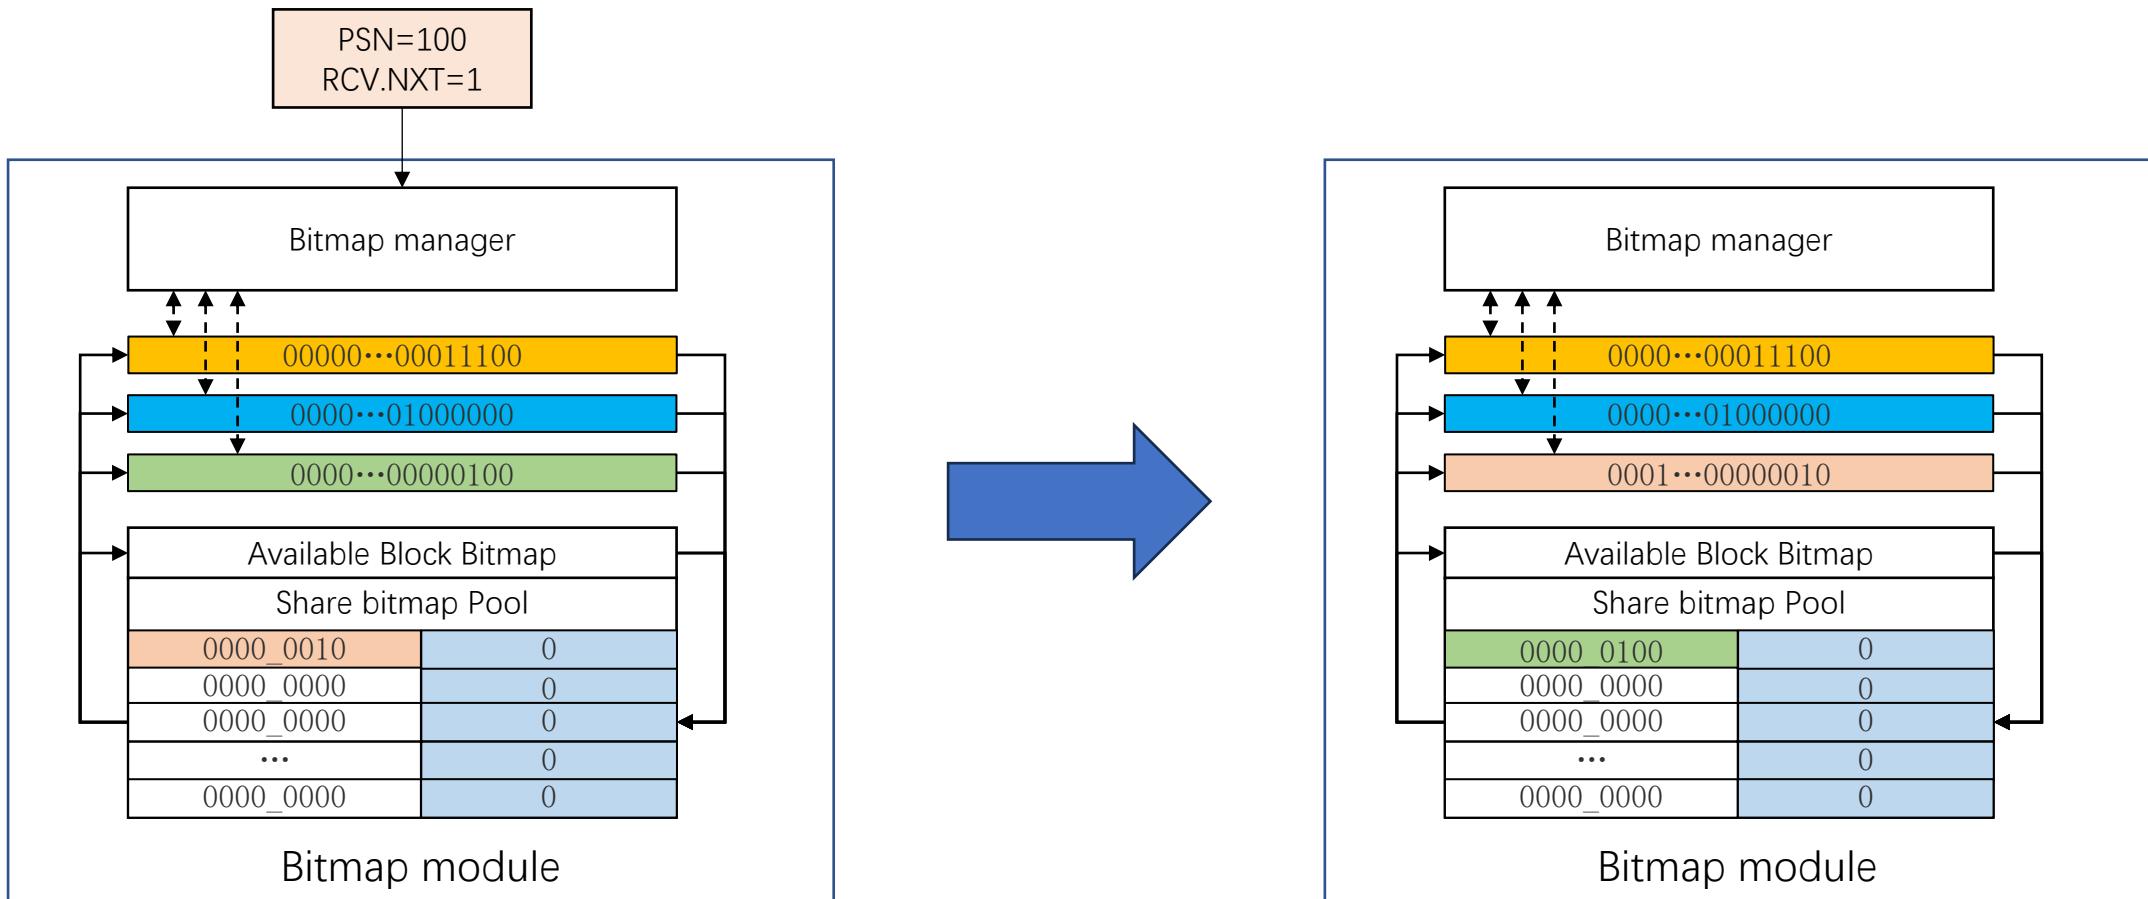
# Dual-state bitmap pool

- Cache miss, bitmap updated in bitmap pool



# Dual-state bitmap pool

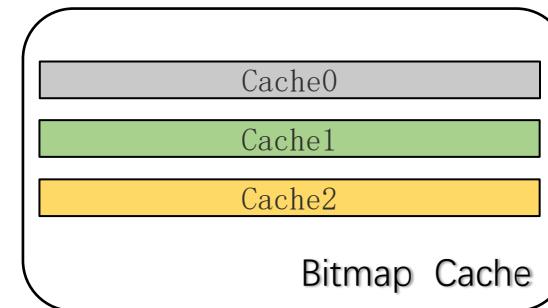
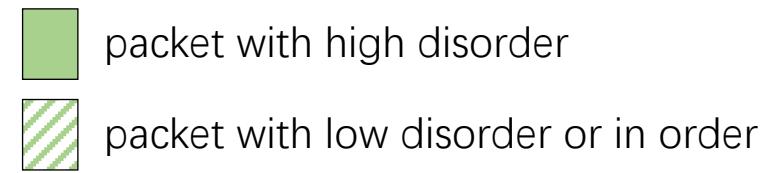
- Cache miss, bitmap updated in bitmap cache



# Bitmap Acceleration

## ■ Cache misses cause high latency

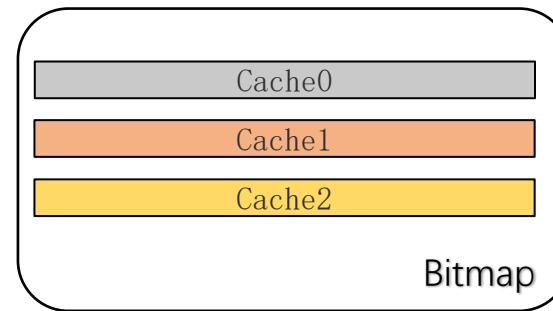
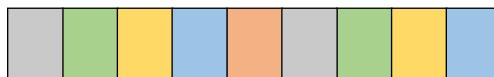
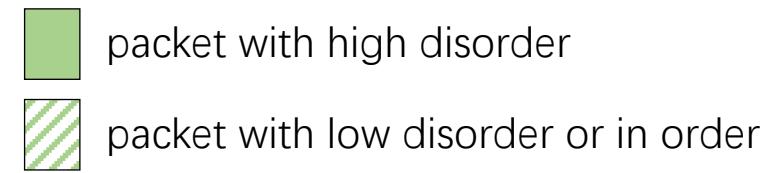
- Cache-pool swapping delay between proportional to the bitmap size
- Switching a bitmap block takes 5ns



# Bitmap Acceleration

## Cache misses cause high latency

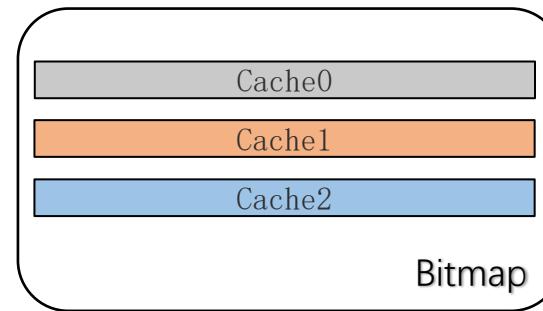
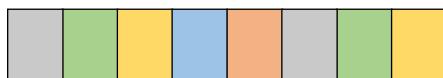
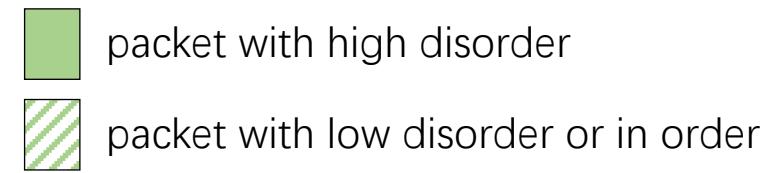
- Cache-pool swapping delay between proportional to the bitmap size
- Switching a bitmap block takes 5ns



# Bitmap Acceleration

## Cache misses cause high latency

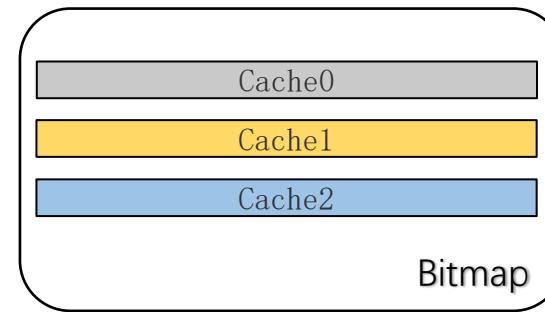
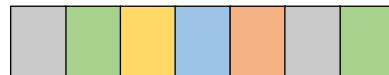
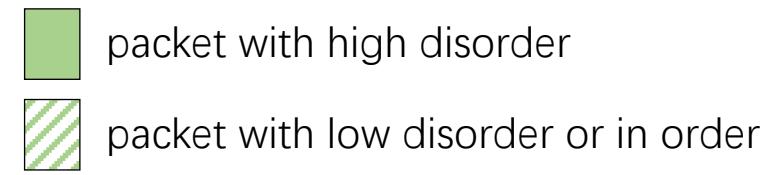
- Cache-pool swapping delay between proportional to the bitmap size
- Switching a bitmap block takes 5ns



# Bitmap Acceleration

## Cache misses cause high latency

- Cache-pool swapping delay between proportional to the bitmap size
- Switching a bitmap block takes 5ns

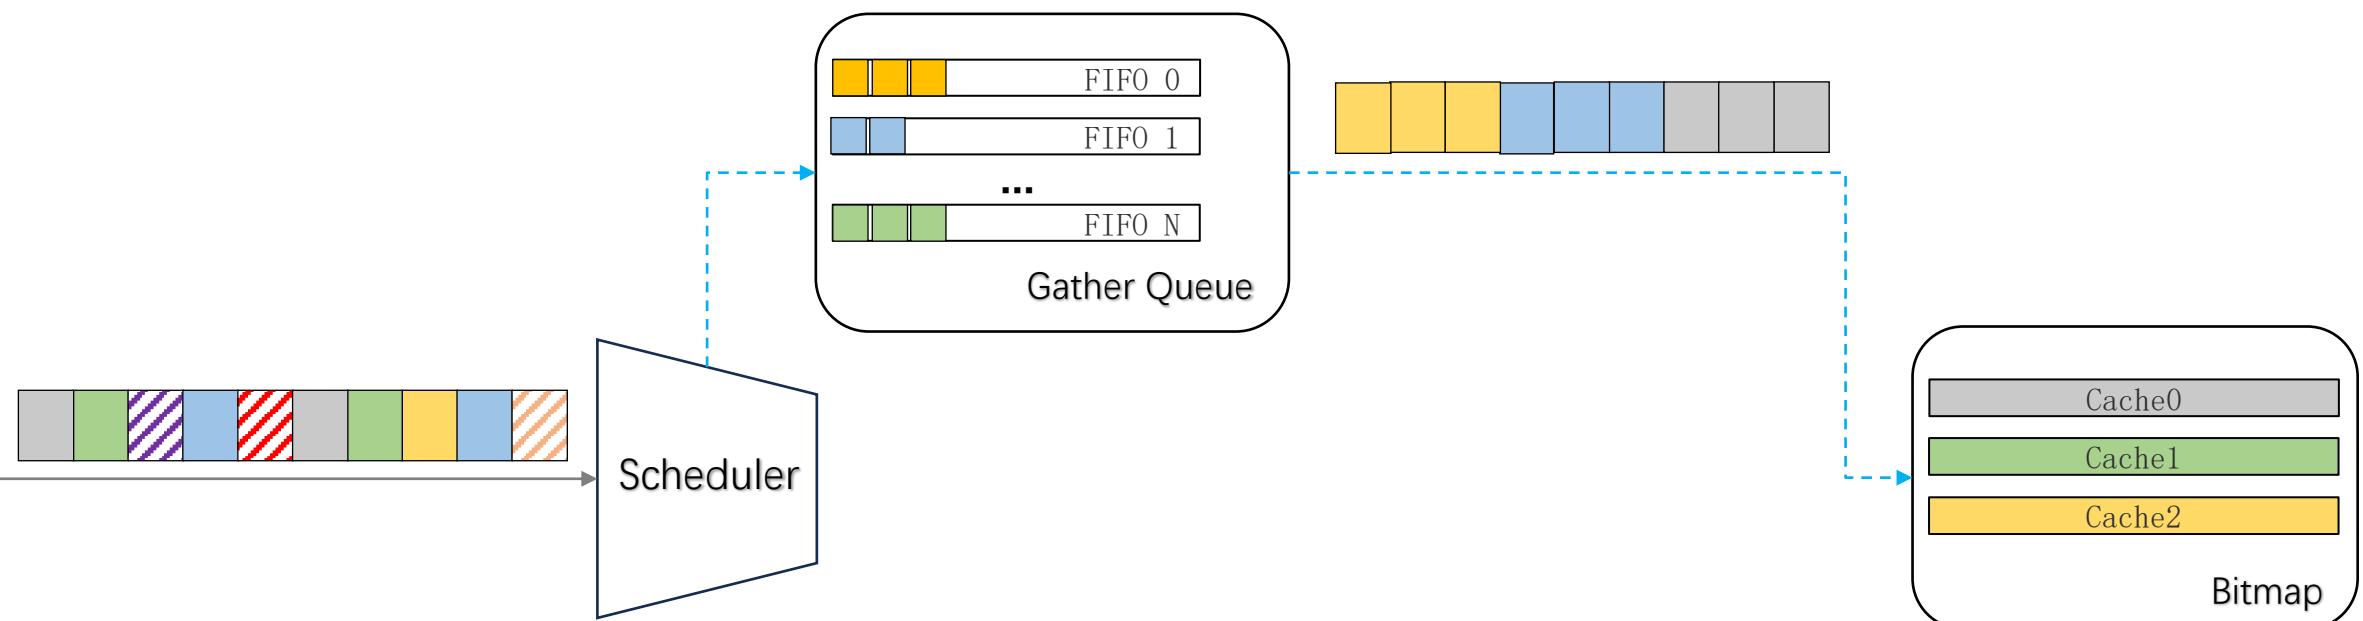
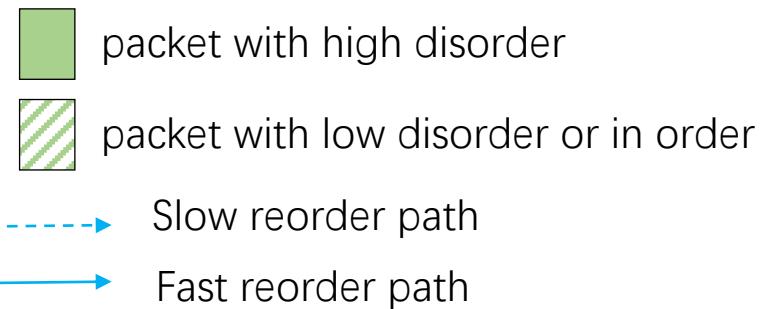


Frequent cache misses cause high bitmap latency and slow RNIC speeds!

# Bitmap Acceleration

## Packet gathering

- Trade-off between delay and cache hit rate
- 8 FIFOs with packet count and timeout
- Gather connection packets into batches for bitmap recording



# Bitmap Acceleration

## Packet Scheduler

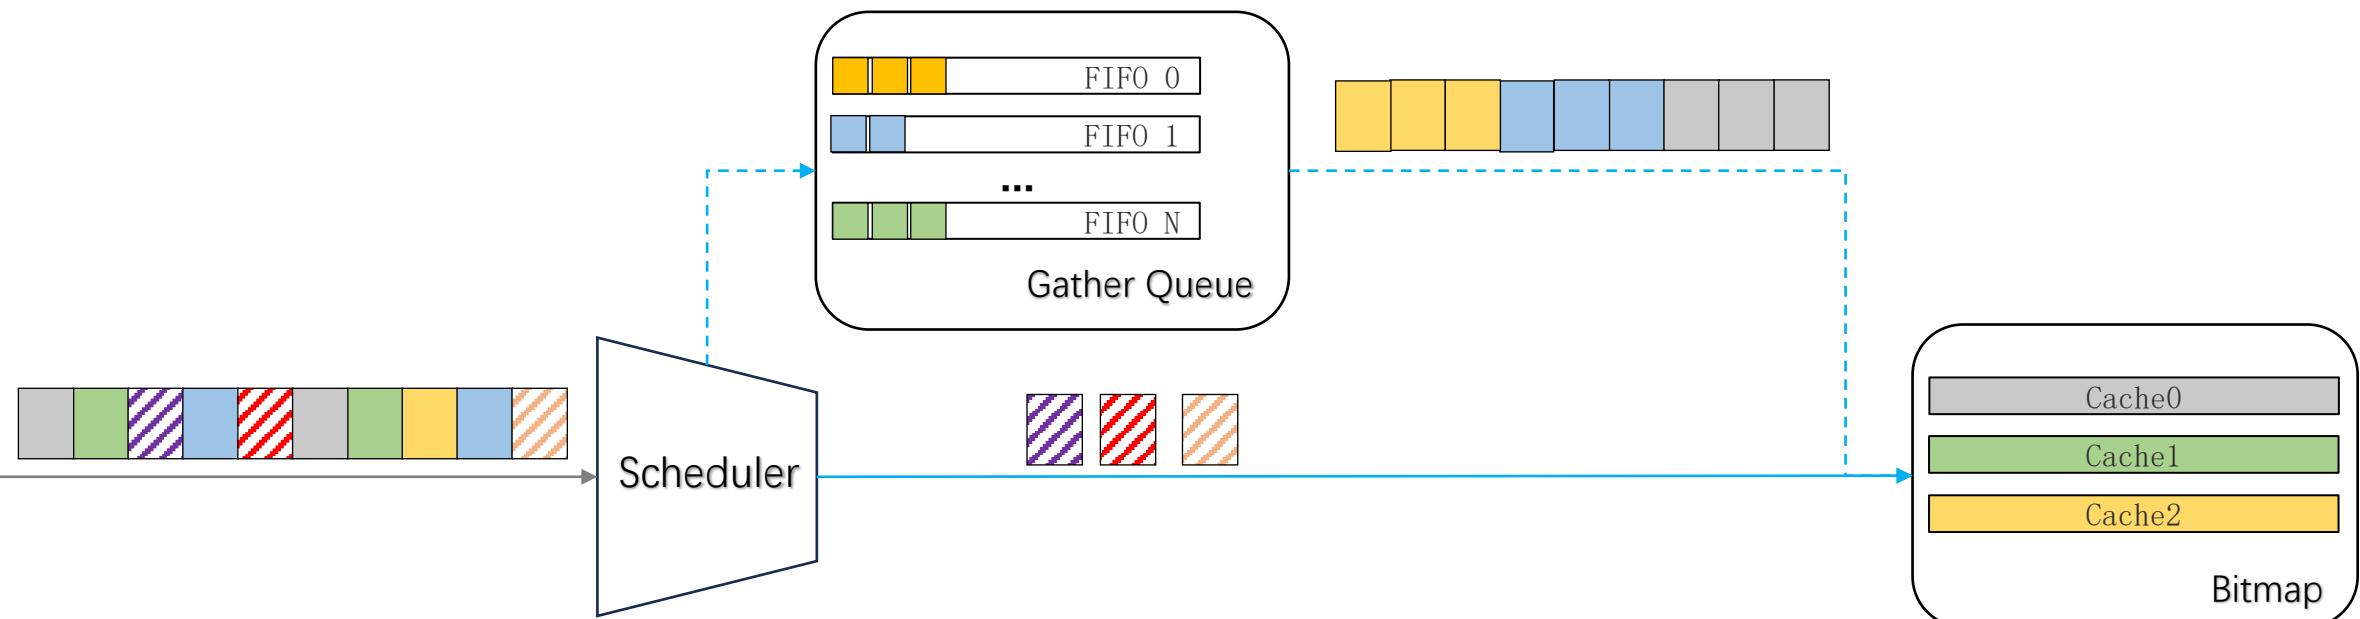
- No disorder, low disorder connection packets do not need to be gathered
- In large disorder connection, packets with low disorder do not need to be gathered

packet with high disorder

packet with low disorder

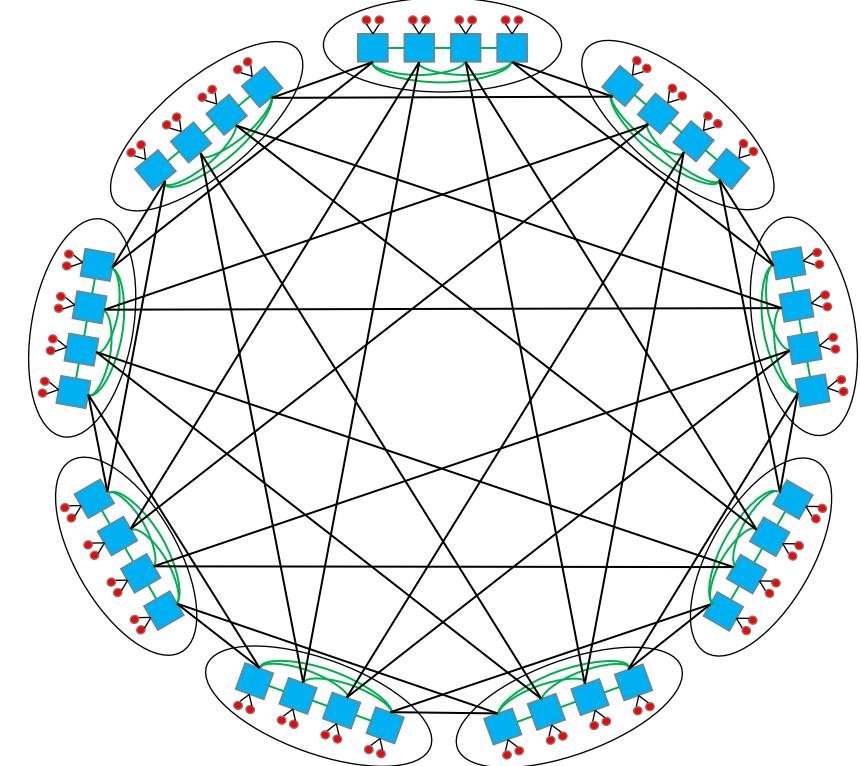
Slow reorder path

Fast reorder path



# NS3 Evaluation

- **Setup**
  - NS3: integrate LEFT into HPCC
- **Topology**
  - Dragonfly, 400Gbps direct between groups
  - Bandwidth: 200Gbps, link latency: 1us
  - MTU: 1KB
  - Bits Pool Size: 2Kb
  - Bitmap caches: 3, aggregate queues: 8
- **Comparison solutions**
  - Perflow, IRN, LEFT-(without schedule), MELO+-(with cache)



# Evaluation-Benefit FCT

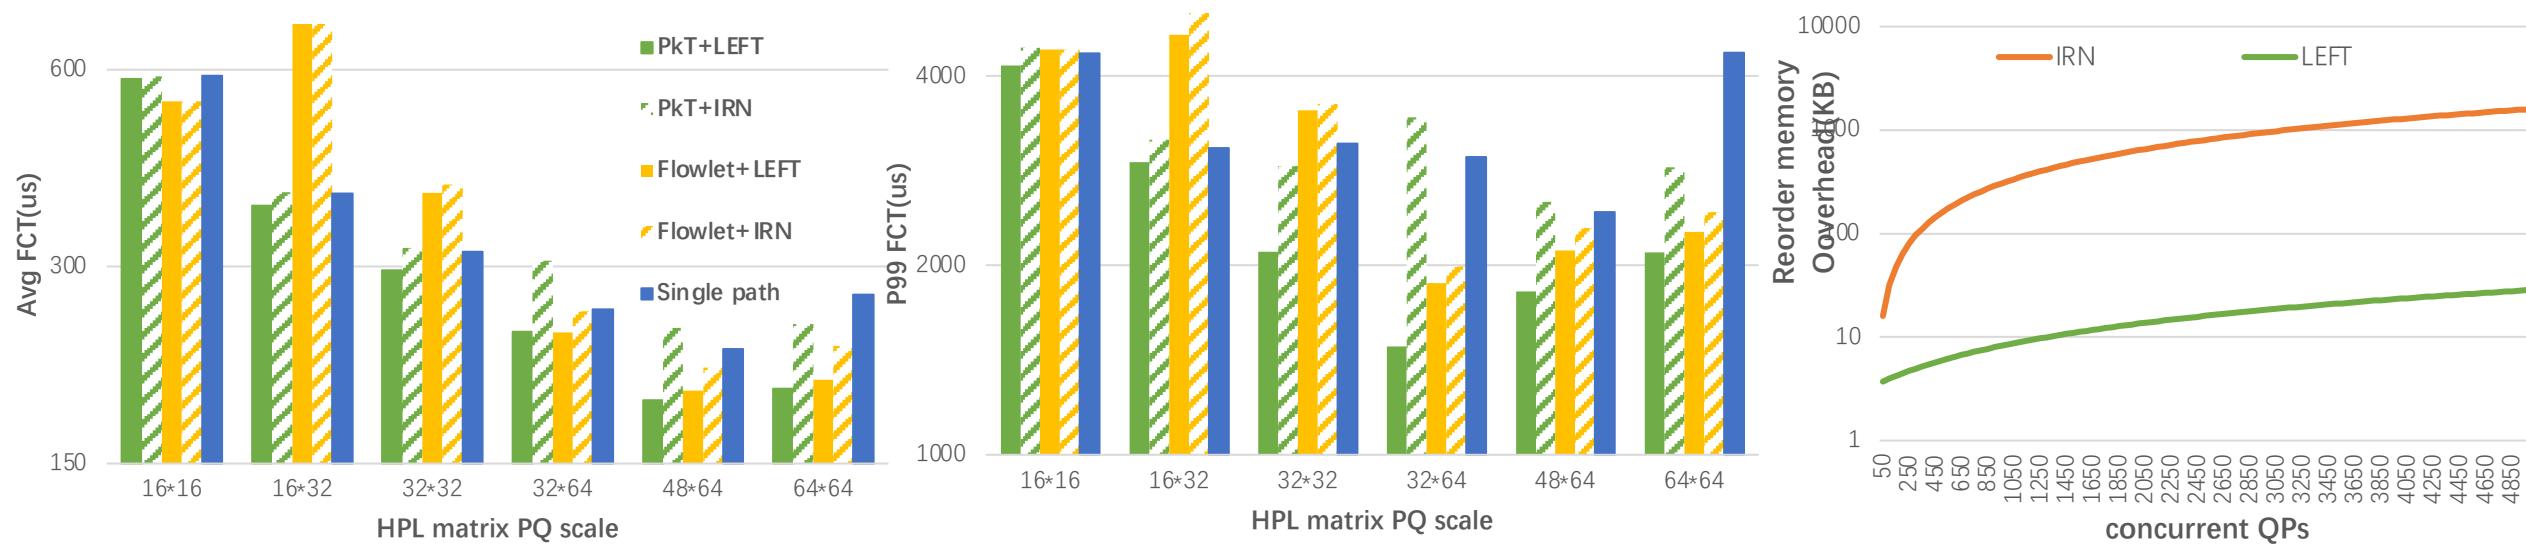
## ■ Average FCT:

- LEFT improve 28.1% compared with per-flow, 20% and 12% reduction compared with per-packet and per-flowlet without LEFT equipped specifically

## ■ 99% tail FCT:

- 48.2% reduction compared with per-flow, 26.7% and 8.1% reduction compared with per-packet and per-flowlet without LEFT equipped specifically

## ■ LEFT adds 27KB memory overhead, only 1.7% of the baseline (IRN)

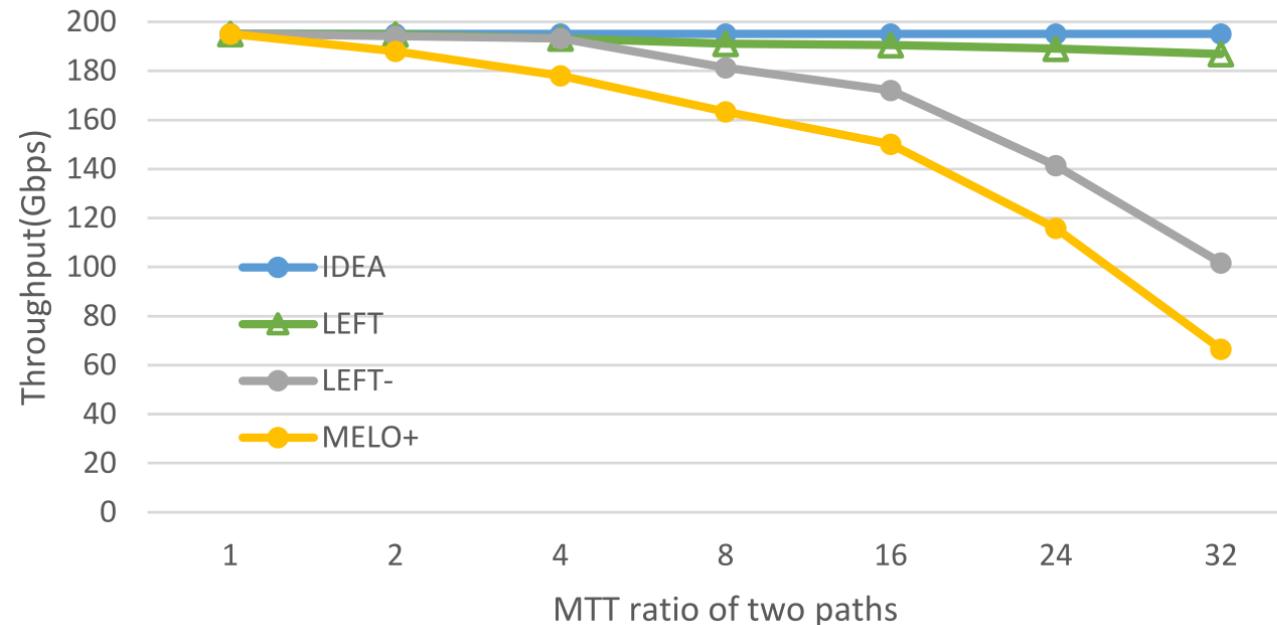


# Evaluation-Benefit Throughput

## ■ Setup:

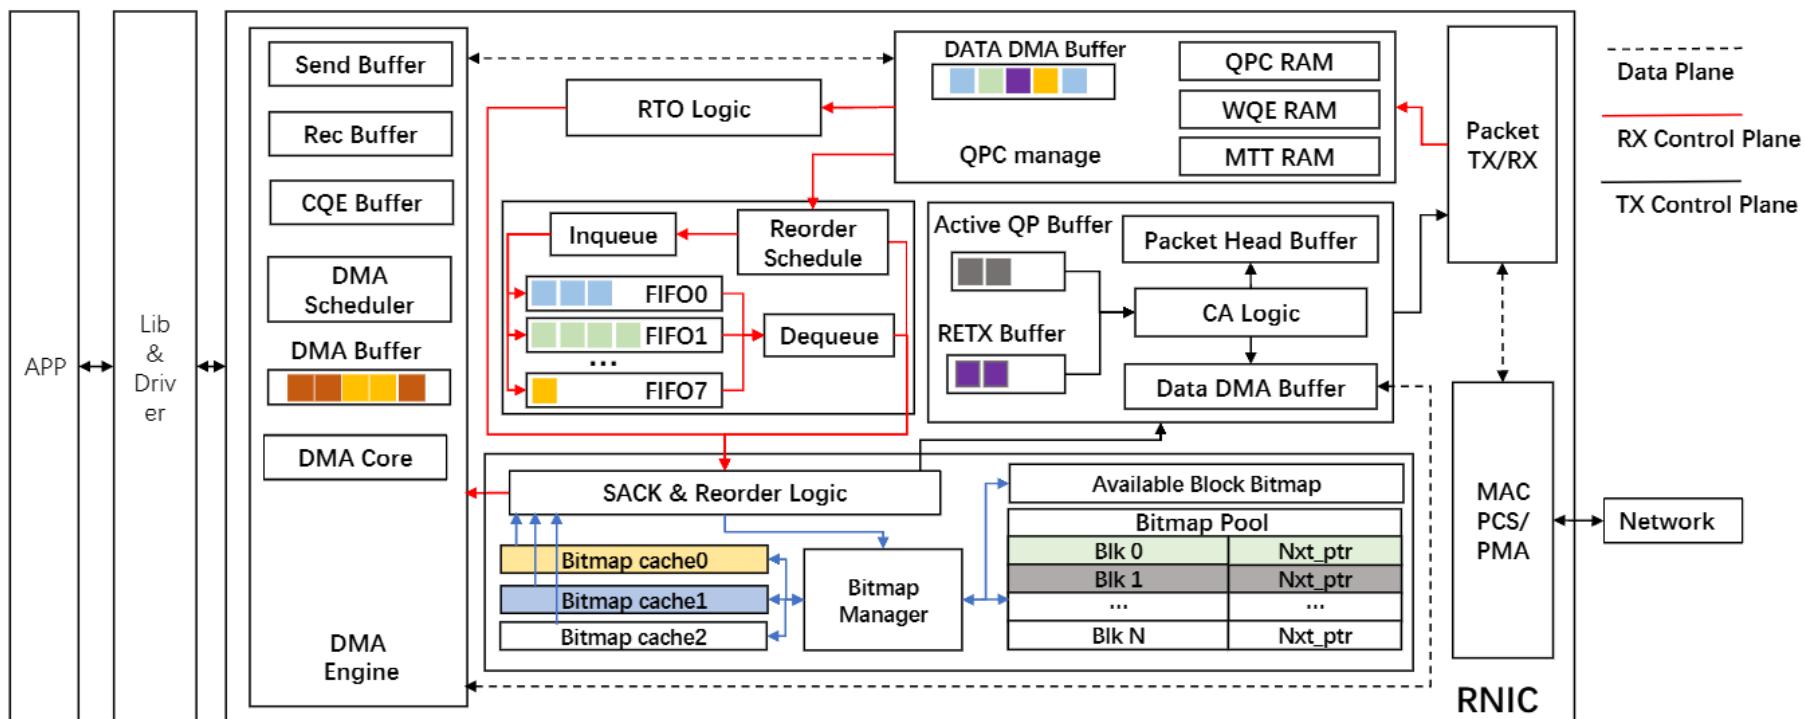
- Two servers connected by two links, and one of the link have different delay(1us,2us...32us)
- evenly distribute all packets on two path

■ 94% throughput under path RTT differs by 32 times. 83% and 180% higher than LEFT- and MELO+, respectively.



# LEFT FPGA implementation

- LEFT is implemented in XCKU040-FFVA1156-2I Xilinx FPGA.
- The board has four 10Gbps ports and one 25Gbps port.
- The RTL code utilizes VIVADO 2022.1 for logic simulation, synthesis, layout, and routing



# FPGA Result

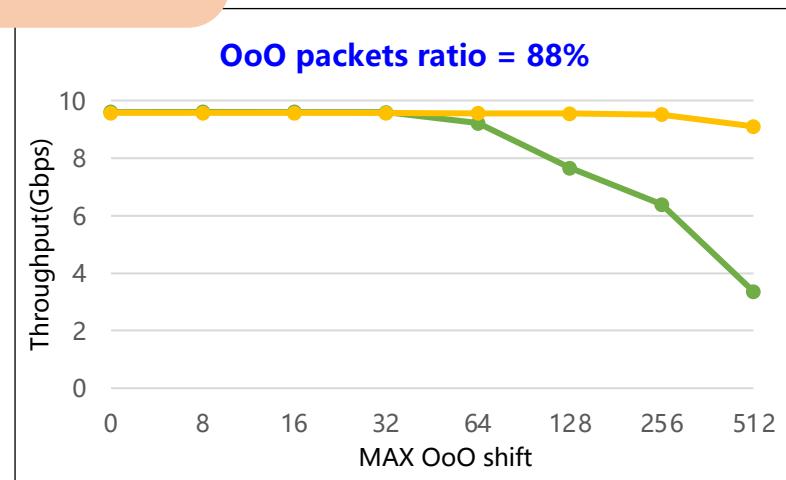
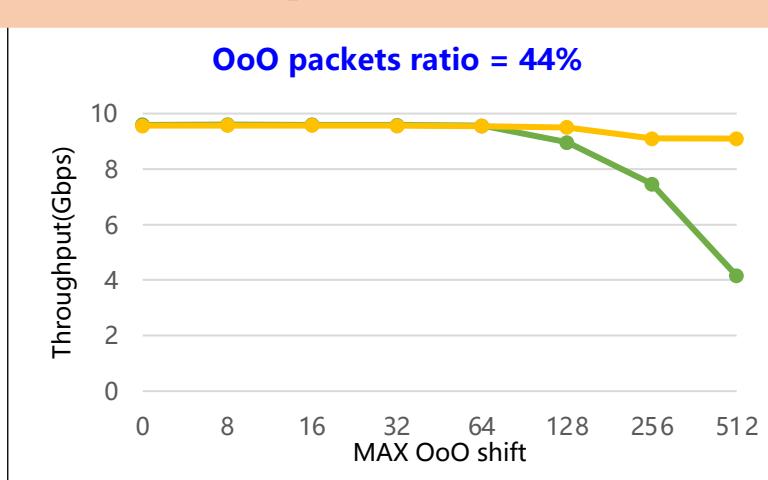
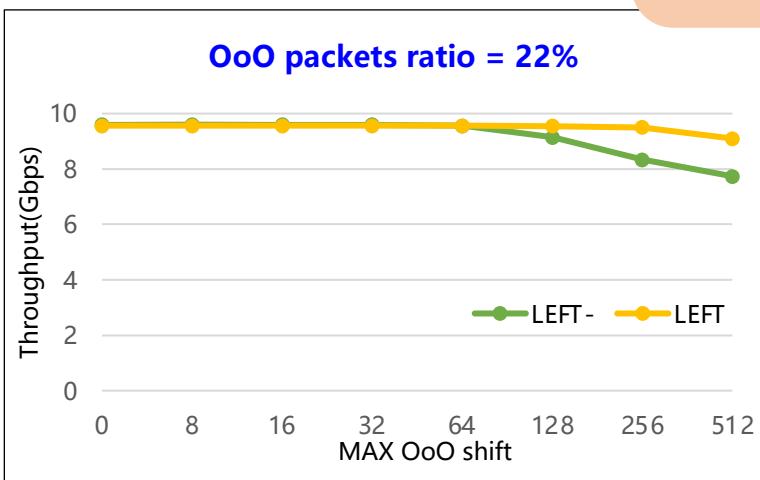
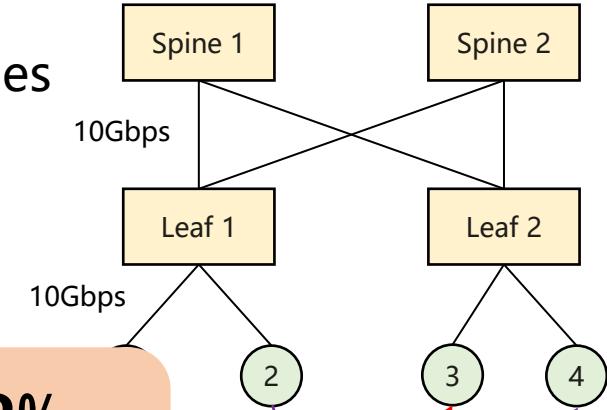
## Setup

- ① sends ③, ② sends ④, 5 connections between each pair of nodes
- Bandwidth: 10Gbps, link latency: 1us
- MTU: 75B
- Bits Pool Size: 2Kb, Bitmask cache: 2 aggregate queues, 8

## Comparison solution

- LEFT, LEFT-

92% throughput under 88% OoO packets ratio



# Conclusion & future work

- **LEFT supporting flexible multi-path transmission without affecting RNIC performance!**
  - By dual-state bitmap pool and In addition, LEFT achieves lightweight memory overhead, with an average of only 7B additional storage overhead per connection
  - By fast and slow reordering path, LEFT achieves exceptional speed, maintaining 94% throughput even when the path RTT differs by 32 times
- **There are some problems worth future study.**
  - Modeling and analysis of the number and size of bitmap blocks, caches, and gather queues
  - Large-scale testbed implementation
  - More efficient judgment of packet loss and reduced overall bitmap overhead



中南大學  
CENTRAL SOUTH UNIVERSITY



湖南大學  
HUNAN UNIVERSITY

中南大學  
CENTRAL SOUTH UNIVERSITY

# Thanks for listening

## Q & A

Contact: 204701003@csu.edu.cn